

Same Prompt, Different Answer: Exposing the Reproducibility Illusion in Large Language Model APIs

LUCAS ROVER*, UTFPR – Universidade Tecnológica Federal do Paraná, Brazil

YARA DE SOUZA TADANO, UTFPR – Universidade Tecnológica Federal do Paraná, Brazil

Background: Generative AI models produce non-deterministic outputs that vary across runs, even under nominally identical configurations. This variability threatens the reproducibility of studies that rely on large language model (LLM) outputs, yet most existing experiment-tracking tools were not designed for the specific challenges of text-generation workflows.

Objectives: We propose a lightweight, open-standard protocol for logging, versioning, and provenance tracking of generative AI experiments. The protocol introduces two novel documentation artifacts—Prompt Cards and Run Cards—and adopts the W3C PROV data model to create auditable, machine-readable provenance graphs linking every output to its full generation context.

Methods: We formalize the protocol and evaluate it empirically through 4,104 controlled experiments. These experiments employ nine model deployments—three locally deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B), five closed-source API-served (GPT-4, Claude Sonnet 4.5, Gemini 2.5 Pro, DeepSeek Chat, Perplexity Sonar), and one cloud-served open-weight model (LLaMA 3 8B via Together AI, as a quasi-isolation probe of infrastructure effects)—on four NLP tasks across seven execution environments (one local plus six independent cloud API providers). Single-turn extraction and summarization are evaluated under greedy decoding for all deployments except Gemini 2.5 Pro (10–30 abstracts per model). Multi-turn refinement and RAG extraction are evaluated for the three local models, Claude Sonnet 4.5, and Gemini 2.5 Pro under greedy decoding (10 abstracts each). Statistical robustness is ensured through Holm-Bonferroni correction across 68 hypothesis tests, Fisher’s exact tests for binary reproducibility, bootstrap confidence intervals (10,000 resamples; percentile method), and sensitivity analysis. We measure output variability using Exact Match Rate, Normalized Edit Distance, ROUGE-L, and BERTScore, and quantify the protocol’s own overhead in terms of time and storage.

Results: Under greedy decoding ($t=0$), local models achieve near-perfect reproducibility (average single-turn EMR = 0.960; Gemma 2 9B: perfect 1.000 across all tasks). Closed-source API models exhibit substantial hidden non-determinism spanning a wide range (EMR 0.010–0.800), yielding an approximately 3-fold local-vs-API gap observed across five providers and surviving Holm-Bonferroni correction (51/68 tests significant). The same LLaMA 3 8B architecture served via Together AI’s cloud endpoint achieves near-local reproducibility (EMR = 1.000/0.880), providing evidence that cloud deployment per se does not preclude determinism—the variability in closed-source models is consistent with infrastructure complexity rather than cloud deployment alone. Under multi-turn refinement and RAG extraction, local models maintain $\text{EMR} \geq 0.880$, while API models exhibit near-zero reproducibility ($\text{EMR} \leq 0.070$). The protocol adds less than 1% overhead.

Conclusions: Our results provide evidence that (1) all five API providers exhibit non-determinism under greedy decoding, while local models achieve near-perfect bitwise reproducibility; (2) API reproducibility spans a

*Corresponding Author.

Authors’ Contact Information: Lucas Rover, ORCID: [0000-0001-6641-9224](#), lucasrover@utfpr.edu.br, UTFPR – Universidade Tecnológica Federal do Paraná, Programa de Pós-Graduação em Engenharia Mecânica, Ponta Grossa, Paraná, Brazil; Yara de Souza Tadano, ORCID: [0000-0002-3975-3419](#), yaratadano@utfpr.edu.br, UTFPR – Universidade Tecnológica Federal do Paraná, Programa de Pós-Graduação em Engenharia Mecânica, Ponta Grossa, Paraná, Brazil.



This work is licensed under a Creative Commons Attribution International 4.0 License.

© 2026 Copyright held by the owner/author(s).

DOI: [10.1613/jair.1.XXXXX](#)

wide range (EMR 0.010–0.800); (3) the gap extends to multi-turn and RAG regimes; (4) cloud deployment per se does not preclude reproducibility (quasi-isolation probe via Together AI); (5) temperature is the dominant user-controllable factor; and (6) provenance logging adds <1% overhead. All primary comparisons survive Holm-Bonferroni correction (51/68 significant). The protocol, implementation, and all data are publicly available.

CCS Concepts: • Software and its engineering → Software testing and debugging; Documentation; • Computing methodologies → Machine learning.

Additional Key Words and Phrases: reproducibility, large language models, non-determinism, provenance, generative AI, experiment tracking, W3C PROV, scientific methodology

JAIR Associate Editor: TBD

JAIR Reference Format:

Lucas Rover and Yara de Souza Tadano. 2026. Same Prompt, Different Answer: Exposing the Reproducibility Illusion in Large Language Model APIs. *Journal of Artificial Intelligence Research* 0, Article 0 (February 2026), 39 pages. DOI: [10.1613/jair.1.XXXXX](https://doi.org/10.1613/jair.1.XXXXX)

1 Introduction

Large language models (LLMs) are rapidly transforming how science is conducted, communicated, and applied. In medicine, LLMs now encode clinical knowledge at expert level ([Singhal et al. 2023](#)) and are being integrated into diagnostic workflows ([Thirunavukarasu et al. 2023](#)). In scientific research more broadly, generative AI is reshaping literature review, data extraction, hypothesis generation, and writing ([Birhane et al. 2023](#)), with experimental evidence showing substantial productivity gains in professional tasks ([Noy and W. Zhang 2023](#)). This expanding adoption extends across disciplines—from legal analysis and education to systematic reviews and meta-analyses—making LLM outputs an increasingly common component of the scientific evidence chain.

Yet this rapid integration rests on an often-unexamined assumption: that LLM outputs are reproducible. Reproducibility—the ability to obtain consistent results when an experiment is repeated under identical conditions—is a cornerstone of the scientific method. The broader reproducibility crisis is well documented: [Baker \(2016\)](#) reported in *Nature* that over 70% of researchers have failed to reproduce another scientist’s experiment. In artificial intelligence, this crisis is particularly acute. [Hutson \(2018\)](#) warned in *Science* that AI faces its own reproducibility crisis, [Gundersen and Kjensmo \(2018\)](#) found that only 6% of AI papers provided sufficient information for full reproducibility, and [Ball \(2023\)](#) recently asked in *Nature* whether AI is actively worsening the problem. [Stodden et al. \(2016\)](#) emphasized in *Science* that computational methods demand their own reproducibility standards, a call echoed by the data leakage analysis of [Kapoor and A. Narayanan \(2023\)](#) across 17 ML-dependent scientific fields.

Generative AI introduces a qualitatively new dimension to this challenge. Unlike traditional computational experiments, in which deterministic algorithms produce identical results given identical inputs, LLMs exhibit inherent variability in their outputs due to stochastic sampling, floating-point non-determinism in distributed GPU inference, and opaque model-versioning practices ([Y. Chen et al. 2023; Zhu et al. 2023](#)). The problem is compounded by several factors unique to text-generation workflows: (1) the same prompt can yield semantically similar yet textually distinct outputs across runs; (2) API-based models may undergo silent updates that alter behavior without notice to the user; (3) temperature and sampling parameters create a high-dimensional space of possible outputs; and (4) the `seed` parameter offered by some APIs is advisory rather than a guarantee—OpenAI explicitly documents that “determinism is not guaranteed” even when a seed is specified ([OpenAI 2024](#)), and Anthropic’s Claude API does not support a seed parameter at all. Recent empirical studies have confirmed these concerns: [Atil et al. \(2024\)](#) found

accuracy variations up to 15% across runs under supposedly deterministic settings, and Ouyang et al. (2024) demonstrated that greedy decoding does not guarantee determinism in ChatGPT code generation.

The demand for AI transparency is not only scientific but also regulatory. The EU AI Act (European Parliament and Council of the European Union 2024) classifies high-risk AI systems—including those used in healthcare and scientific research—as requiring documented traceability and auditability, and the NIST AI Risk Management Framework (National Institute of Standards and Technology 2023) emphasizes transparency and accountability as core principles. The FAIR data principles (Wilkinson et al. 2016) provide a foundation for data stewardship, yet no established standard exists for documenting the full context needed to understand, audit, or reproduce a generative AI output. Existing experiment-tracking tools such as MLflow (Zaharia et al. 2018), Weights & Biases (Biewald 2020), and DVC (Kuprieiev et al. 2024) were designed primarily for training pipelines and numerical metrics. Although valuable for their intended purposes, these tools lack features critical for generative AI studies: structured prompt versioning, cryptographic output hashing for tamper detection, provenance graphs linking outputs to their full generation context, and environment fingerprinting specific to inference-time conditions.

In this paper, we address this gap with three contributions, with the protocol design as the primary and most durable contribution (Figure 1):

- (1) **A lightweight, standards-based protocol** for logging, versioning, and provenance tracking of generative AI experiments. The protocol introduces *Prompt Cards* and *Run Cards* as structured documentation artifacts, and adopts the W3C PROV data model (Moreau and Missier 2013) for machine-readable provenance graphs. It operationalizes—and extends to generative AI workflows—the reproducibility checklist and badge mechanisms recently adopted by JAIR (Gundersen, Helmert, et al. 2024), providing machine-readable infrastructure that automates what those mechanisms require researchers to document manually.
- (2) **A large-scale empirical case study** demonstrating both the protocol’s effectiveness and the scope of hidden non-determinism in current LLM APIs. Through 4,104 controlled experiments with nine model deployments across four NLP tasks, seven execution environments, and five conditions (Section 4), we document a substantial and previously invisible reproducibility gap between local and API-based inference—a gap that persists across five independent cloud providers and extends to multi-turn and retrieval-augmented generation regimes.
- (3) **A reference implementation** in Python, together with all experimental data and provenance records, publicly available to facilitate adoption and independent verification.

The remainder of this paper is organized as follows. Section 2 reviews related work on reproducibility in AI and experiment tracking. Section 3 formalizes the protocol design. Section 4 describes the experimental methodology. Section 5 presents the empirical results, revealing a multi-fold reproducibility gap between local and API-served models that is invisible without systematic logging. Section 6 discusses findings, limitations, and practical implications. Section 7 concludes with directions for future work.

2 Related Work

2.1 Reproducibility in AI Research

The reproducibility crisis in AI has been documented extensively. Gundersen and Kjensmo (2018) surveyed 400 AI papers and found that only 6% provided sufficient information for full reproducibility. Pineau et al. (2021) reported on the NeurIPS 2019 Reproducibility Program, which introduced reproducibility checklists and found significant gaps between reported and actual reproducibility. More recently, Gundersen, Helmert, et al. (2024) described four institutional mechanisms adopted by JAIR—reproducibility checklists, structured abstracts, badges, and reproducibility reports—establishing a community standard for what

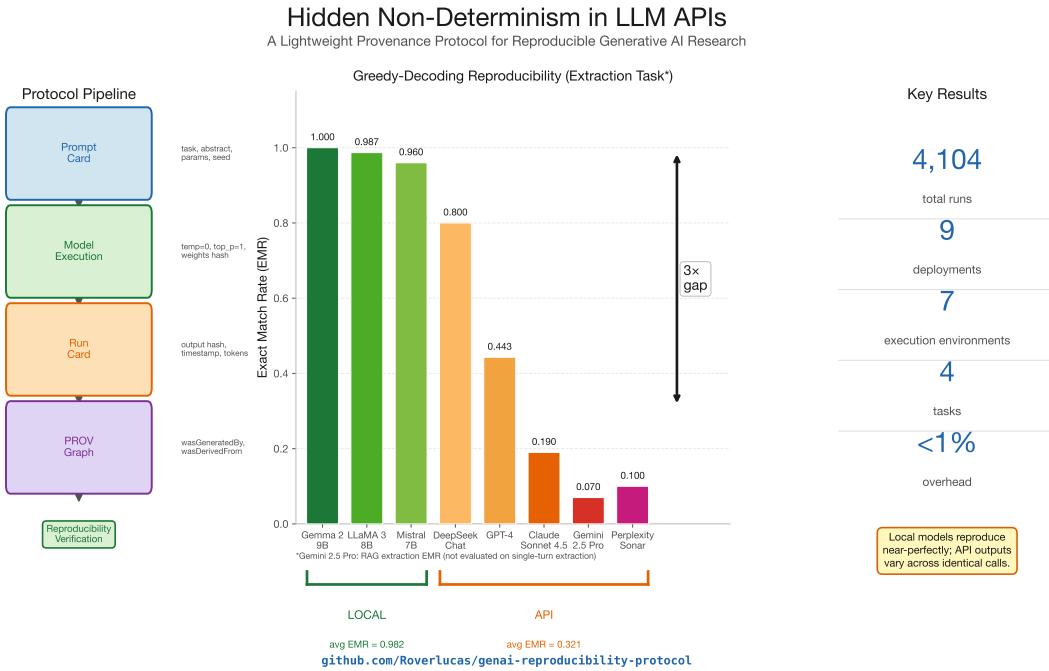


Fig. 1. Visual abstract summarizing the study design and key findings. **Left:** The provenance protocol pipeline, from Prompt Card creation through Run Card logging and W3C PROV graph generation. **Center:** Exact Match Rates (EMR) under greedy decoding for eight of the nine model deployments (Together AI's cloud LLaMA 3 8B, which mirrors local LLaMA's values, is omitted for visual clarity), illustrating the reproducibility gap between local models (green, $\text{EMR} \geq 0.960$) and API-served models (orange/red, $\text{EMR} \leq 0.800$). Gemini 2.5 Pro shows RAG extraction EMR (not evaluated on single-turn extraction). **Right:** Key statistics: 4,104 experiments, 9 deployments across 7 execution environments (1 local + 6 cloud API providers), 4 tasks, <1% protocol overhead.

should be documented in AI research. Gundersen, Gil, et al. (2018) identified three levels of reproducibility in AI—method, data, and experiment—and argued that all three are necessary for scientific progress. Belz, Agarwal, et al. (2021) conducted a systematic review of 601 NLP papers and confirmed pervasive under-reporting of experimental details; Belz, Thomson, et al. (2022) further showed that missing information makes it practically impossible to assess reproducibility of human evaluations in NLP. Rogers et al. (2021) proposed incentive structures to improve reproducibility norms in computational linguistics. Dodge et al. (2019) proposed improved reporting standards for ML experiments, including confidence intervals and significance tests across multiple runs, and Gundersen, Coakley, et al. (2022) provided a comprehensive taxonomy of irreproducibility sources in machine learning. More broadly, Kapoor and A. Narayanan (2023) identified data leakage as a widespread driver of irreproducible results across 17 scientific fields that use ML-based methods.

For generative AI specifically, Y. Chen et al. (2023) demonstrated that ChatGPT's outputs on NLP benchmarks exhibit non-trivial variability across identical queries, even with temperature set to zero. Zhu et al. (2023) showed that reproducibility degrades further when tasks involve subjective judgment, such as social computing annotations. Most recently, Atil et al. (2024) systematically measured the

non-determinism of five LLMs under supposedly deterministic settings across eight tasks, finding accuracy variations up to 15% across runs and introducing the Total Agreement Rate (TAR) metric. Ouyang et al. (2024) confirmed that $t=0$ (greedy decoding) does not guarantee determinism in ChatGPT code generation. Concurrently, Yuan et al. (2025) traced such non-determinism to numerical precision issues in GPU kernels and proposed LayerCast as a mitigation strategy—a hardware-level fix that reduces but does not eliminate non-determinism, and that is not available to researchers using closed API services. The PyTorch documentation (PyTorch Contributors 2024) further catalogues sources of non-determinism in GPU operations, providing the `torch.use_deterministic_algorithms()` flag as a partial mitigation for training; however, this flag is unavailable for API-served inference. Our Exact Match Rate (EMR) metric is closely related to Atil et al.’s Total Agreement Rate (TAR), which measures the fraction of runs producing the modal output; EMR instead measures the fraction of *all output pairs* that match exactly, providing a more sensitive measure when agreement is low and no clear modal output exists. Our work complements these studies in four specific ways. First, whereas prior studies (including Atil et al.’s five-model, eight-task study) measure variability post hoc, we provide a structured provenance protocol that enables *prospective* documentation and audit—answering not only “how much variability?” but also “why did these outputs differ?” through cryptographic hashing and W3C PROV graphs. Second, we directly compare local and API-based inference on identical tasks with identical prompts across nine model deployments and six independent cloud providers, isolating the deployment paradigm as a variable—including a quasi-isolation probe with the same open-weight model under local and cloud deployment—and confirming that API non-determinism is a consistent pattern across providers (EMR 0.010–0.800). Third, we extend beyond single-turn evaluation to include multi-turn refinement and retrieval-augmented generation, demonstrating that reproducibility characteristics generalize across interaction regimes. Fourth, we quantify the overhead of systematic logging, demonstrating that the “cost of knowing” is negligible.

2.2 Experiment Tracking Tools

Several tools exist for tracking machine learning experiments, although none was designed specifically for generative AI text-output workflows:

MLflow (Zaharia et al. 2018) provides experiment tracking, model packaging, and deployment. It logs parameters, metrics, and artifacts, but focuses on training pipelines and numerical outcomes rather than text-generation provenance.

Weights & Biases (Biewald 2020) offers experiment tracking with visualization dashboards. It supports prompt logging but lacks structured prompt versioning, cryptographic output hashing, and provenance graph generation.

DVC (Kuprieiev et al. 2024) provides data versioning through git-like operations. While effective for dataset management, it does not address run-level provenance or prompt documentation.

OpenAI Eval (OpenAI 2023) is a framework for evaluating LLM outputs against benchmarks. It provides structured evaluation but is tightly coupled to OpenAI’s ecosystem and does not generate interoperable provenance records.

LangSmith (LangChain 2023) offers tracing and evaluation for LLM applications. It captures detailed execution traces but uses a proprietary format and requires cloud connectivity.

More broadly, Bommasani et al. (2022) identified reproducibility as a key risk for foundation models, and Liang et al. (2023) proposed the HELM benchmark for holistic evaluation of language models, including robustness and fairness dimensions that complement our reproducibility focus. In the provenance space, Padovani et al. (2025) recently introduced yProv4ML, a framework that captures ML provenance in PROV-JSON format with minimal code modifications; our protocol shares the commitment to W3C

Table 1. Comparison of our protocol with existing reproducibility tools and frameworks for GenAI experiments. Checkmarks (✓) indicate full support; tildes (~) indicate partial support; dashes (–) indicate no support.

Feature	Ours	MLflow	W&B	DVC	OpenAI Eval	LangSmith
Prompt versioning (Prompt Card)	✓	–	~	–	~	~
Run-level provenance (W3C PROV)	✓	–	–	–	–	–
Cryptographic output hashing	✓	–	–	✓	–	–
Seed & param logging	✓	✓	✓	–	✓	✓
Environment fingerprinting	✓	~	~	~	–	–
Model weights hashing	✓	–	~	✓	–	–
Overhead <1% of inference	✓	~	~	N/A	N/A	~
Designed for GenAI text output	✓	–	–	–	✓	✓
Open standard (PROV-JSON)	✓	–	–	–	–	–
Local-first (no cloud dependency)	✓	✓	–	✓	–	–

PROV and SHA-256 hashing but differs in three key respects: (i) we target inference-time stochastic text generation rather than training pipelines; (ii) our Run Cards capture prompt-level metadata (prompt hash, seed status, interaction regime) not present in training-oriented schemas; and (iii) we provide empirical evidence quantifying why such logging is necessary for API-served models.

Table 1 provides a systematic feature-by-feature comparison of our protocol with these tools. The key distinction is not merely one of tooling but of *scientific capability*: existing tools log what happened during training (parameters, metrics, artifacts), whereas our protocol enables answering questions that these tools cannot—specifically, whether two generative outputs are provably derived from identical configurations, which exact factor caused a divergence between non-identical outputs, and whether an output has been tampered with post-generation. These capabilities require the combination of cryptographic hashing, structured prompt documentation, and W3C PROV provenance graphs that no existing tool provides. In short, our contribution is not an alternative experiment tracker but a *reproducibility assessment framework* designed for the unique challenges of stochastic text generation.

2.3 Provenance in Scientific Computing

Data provenance—the lineage of data through transformations—has a rich history in database systems and scientific workflows (Herschel et al. 2017). The W3C PROV family of specifications (Moreau and Missier 2013) provides a standardized data model for representing provenance as directed acyclic graphs of *entities*, *activities*, and *agents*. Samuel and König-Ries (2022) applied provenance tracking to computational biology workflows, demonstrating its value for reproducibility. However, to our knowledge, no prior work has applied W3C PROV specifically to generative AI experiment workflows, in which the challenge involves not only tracking data lineage but also capturing the stochastic generation context that determines output variability.

Taken together, these gaps point to a clear need: a lightweight, standards-based protocol that bridges generative AI inference with the provenance infrastructure already established in scientific computing. The next section presents our design for such a protocol.

3 Protocol Design

Our protocol addresses the question: *What is the minimum set of metadata that must be captured for each generative AI run to enable auditing, reproducibility assessment, and provenance tracking?* We address this question through four complementary components.

3.1 Scope and Design Principles

The protocol is designed around three principles:

- (1) **Completeness:** Every factor that can influence a generative output must be captured—prompt text, model identity and version, inference parameters, environment state, and timestamps.
- (2) **Negligible overhead:** The logging process must not materially affect the experiment. We target <1% overhead relative to inference time.
- (3) **Interoperability:** All artifacts are stored in open, machine-readable formats (JSON, PROV-JSON), aligned with the FAIR (Findable, Accessible, Interoperable, Reusable) principles ([Wilkinson et al. 2016](#)), to enable tool integration and long-term preservation.

3.2 Prompt Cards

A *Prompt Card* is a versioned documentation artifact that captures the design rationale and metadata for a prompt template used in experiments. Each Prompt Card contains:

- `prompt_id`: Unique identifier
- `prompt_hash`: SHA-256 hash of the prompt text, enabling tamper detection
- `version`: Semantic version number
- `task_category`: Classification of the task (e.g., summarization, extraction)
- `objective`: Natural-language description of what the prompt is designed to achieve
- `assumptions`: Explicit assumptions about inputs and expected behavior
- `limitations`: Known limitations or failure modes
- `target_models`: Models for which the prompt was designed and tested
- `expected_output_format`: Description of the expected output structure
- `interaction_regime`: Single-turn, multi-turn, or chain-of-thought
- `change_log`: History of modifications

Prompt Cards serve two purposes: they document design intent (supporting human understanding) and they provide a citable, hashable reference for automated provenance tracking. The concept draws inspiration from Model Cards ([Mitchell et al. 2019](#)), Datasheets for Datasets ([Gebru et al. 2021](#)), and model info sheets for reproducibility assessment ([Kapoor and A. Narayanan 2023](#)), extending the structured-documentation paradigm to the prompt layer of the generative AI pipeline.

3.3 Run Cards

A *Run Card* captures the complete execution context of a single generative AI run. Each Run Card records 24 core fields organized into five groups (the complete JSON schema in Appendix B includes these fields plus additional metadata such as `researcher_id`, `affiliation`, `system_logs`, and `errors`):

- (1) **Identification:** `run_id`, `task_id`, `task_category`, `prompt_hash`, `prompt_text`
- (2) **Model context:** `model_name`, `model_version`, `weights_hash`, `model_source`
- (3) **Parameters:** `inference_params` (temperature, top_p, top_k, max_tokens, seed, decoding_strategy), `params_hash`
- (4) **Input/Output:** `input_text`, `input_hash`, `output_text`, `output_hash`, `output_metrics`

Run Card Schema (24 core + extension fields)
1. Identification <code>run_id · task_id · task_category · prompt_hash · prompt_text</code>
2. Model Context <code>model_name · model_version · weights_hash · model_source</code>
3. Parameters <code>inference_params {temp, top_p, top_k, max_tokens, seed, strategy} · params_hash</code>
4. Input/Output <code>input_text · input_hash · output_text · output_hash · output_metrics</code>
5. Execution Metadata <code>environment · environment_hash · code_commit · timestamps · duration_ms · overhead_ms · storage_kb</code>
API Extensions (optional) <code>api_request_id · api_region · seed_status ∈ {sent, logged-only, not-supported}</code>
Workflow Extensions (optional) <code>conversation_history_hash · turn_index · retrieval_context_hash · parent_run_id</code>

Fig. 2. Run Card minimal schema. All SHA-256 hashes (5 total) enable tamper detection and automated comparison. API and workflow extension fields are optional.

- (5) **Execution metadata:** `environment` (OS, architecture, Python version, hostname), `environment_hash`, `code_commit`, `timestamps` (start/end), `execution_duration_ms`, `logging_overhead_ms`, `storage_kb`

For API-served models, optional extension fields capture provider-specific metadata that may help diagnose non-determinism: `api_request_id`, `api_response_headers`, `api_model_version_returned`, `api_region`, and a `seed_status` field that distinguishes between seeds that were “sent” to the API, “logged-only” (recorded for protocol parity but not sent, as with Claude), or “not-supported” by the provider. This formalization ensures that the advisory or absent nature of API seed parameters is captured as structured metadata rather than left as an undocumented assumption.

Figure 2 illustrates the Run Card schema as a minimal structured record.

The separation of logging overhead from execution time is deliberate: it allows researchers to verify that the protocol itself does not confound experimental measurements.

3.3.1 Normative Field Requirements. To support adoption as a citable specification, we classify Run Card fields using normative language following RFC 2119 (Bradner 1997):

- **MUST** (required for audit completeness): `run_id`, `prompt.text`, `prompt_hash`, `model_name`, `model.version`, `output.text`, `output_hash`, `timestamp_start`, `inference_params` (including temperature, seed, decoding strategy).
- **SHOULD** (strongly recommended): `input_hash`, `params_hash`, `environment_hash`, `weights_hash` (local), `code_commit`, `execution_duration_ms`, `logging_overhead_ms`, `seed_status` (API).
- **MAY** (optional, context-dependent): `api_request_id`, `api_response_headers`, `api_region`, `conversation_history_hash`, `turn_index`, `retrieval_context_hash`, `parent_run_id`, `researcher_id`, `affiliation`.

A conforming implementation MUST populate all MUST fields and SHOULD populate all SHOULD fields. The MUST set is minimal: removing any MUST field renders at least one audit question from Section 6.8 unanswerable.

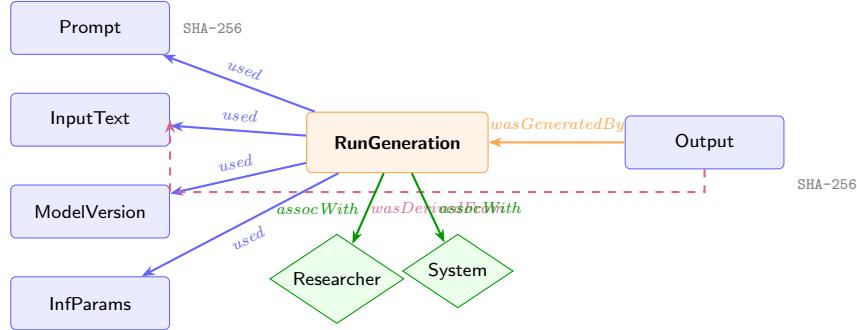


Fig. 3. W3C PROV graph for a single generative AI run. Blue rectangles: entities (data artifacts with SHA-256 hashes). Orange rectangle: activity (the inference execution). Green diamonds: agents. Solid arrows: used and wasGeneratedBy relations. Dashed arrow: wasDerivedFrom. Comparing two such graphs automatically identifies the divergence source when outputs differ.

3.4 W3C PROV Integration

Each experimental group (defined by a unique model–task–condition–abstract combination) is automatically translated into a W3C PROV-JSON document (Moreau and Missier 2013) that expresses the generation provenance as a directed graph. The mapping defines:

- **Entities:** Prompt, InputText, ModelVersion, InferenceParameters, Output, ExecutionMetadata
- **Activities:** RunGeneration (the inference execution)
- **Agents:** Researcher, SystemExecutor (the execution environment)

PROV relations capture the causal structure:

- **used:** RunGeneration used Prompt, InputText, ModelVersion, InferenceParameters
- **wasGeneratedBy:** Output wasGeneratedBy RunGeneration
- **wasAssociatedWith:** RunGeneration wasAssociatedWith Researcher, SystemExecutor
- **wasAttributedTo:** Output wasAttributedTo Researcher
- **wasDerivedFrom:** Output wasDerivedFrom InputText

This standardized representation enables automated reasoning about experiment provenance, including detecting when two runs share identical configurations and identifying the specific factors that differ between non-identical outputs. The choice of W3C PROV over plain JSON logs is deliberate: PROV’s formal semantics allow automated tools to traverse the provenance graph and answer queries such as “what changed between these two runs?” without custom parsing logic. An abbreviated example document is given in Supplementary Material S5. Figure 3 illustrates the provenance graph structure for a single experimental run.

3.5 Reproducibility Checklist

We provide a 15-item checklist organized into four categories—Prompt Documentation, Model and Environment, Execution and Output, and Provenance—that researchers can use to self-assess the reproducibility of their generative AI studies. The complete checklist is provided in Appendix A.

3.6 Extensions for Advanced Workflows

The protocol’s field schema accommodates complex workflows through optional extension fields. Our empirical evaluation exercises two of these extensions—multi-turn dialogues and RAG—while the remaining extensions are specified in the reference implementation’s schema:

- **Multi-turn dialogues:** A `conversation_history_hash` field and `turn_index` enable linking each turn to the full conversation state. *Evaluated in Task 3 (multi-turn refinement) using Ollama’s `/api/chat` endpoint.*
- **RAG:** Fields for retrieval context (with hashes) trace which external information influenced the output. *Evaluated in Task 4 (RAG extraction) with prepended context passages.*
- **Tool use and function calling:** Fields for available tools, tool calls (with arguments, results, and hashes) capture the full tool-use chain.
- **Chain-of-thought / agent workflows:** A `parent_run_id` field supports hierarchical provenance graphs for multi-step reasoning chains.

3.7 Formal Definition and Audit Completeness

We define the protocol as a tuple $\mathcal{P} = (PC, RC, G, CL)$, where PC is a Prompt Card, RC is a Run Card, G is a W3C PROV graph, and CL is the reproducibility checklist. Each Run Card RC_i is itself a tuple of field groups: $RC_i = (Id, Mod, Par, IO, Env, H)$, where H denotes the set of five SHA-256 hashes (prompt, input, parameters, environment, output).

We define an *audit question* as a predicate Q over one or more Run Cards. The protocol satisfies the following *audit completeness* property: for a set of 10 audit questions $\{Q_1, \dots, Q_{10}\}$ (defined in Section 6.8), every Q_j is answerable if and only if all field groups are populated. Formally:

$$\forall Q_j \in \{Q_1, \dots, Q_{10}\} : \text{answerable}(Q_j, RC_i) \Leftrightarrow \bigwedge_{g \in \text{required}(Q_j)} g \subseteq RC_i \quad (1)$$

where $\text{required}(Q_j)$ maps each question to its minimal set of required field groups. The ablation analysis in Section 6.8 demonstrates that every field group is in the required set of at least one question, establishing protocol *minimality*: removing any field group violates Equation 1 for at least one Q_j .

The *differential diagnosis* property follows from the hash fields: given two Run Cards RC_a, RC_b with $H_{\text{output}}^a \neq H_{\text{output}}^b$, the protocol enables automatic identification of the divergence source by comparing the remaining hashes. If $H_{\text{prompt}}^a = H_{\text{prompt}}^b, H_{\text{input}}^a = H_{\text{input}}^b, H_{\text{params}}^a = H_{\text{params}}^b$, and $H_{\text{env}}^a = H_{\text{env}}^b$, then the output difference is attributable to non-determinism in the generation process itself—precisely the phenomenon we measure empirically in Section 5.

Having defined the protocol’s components and formal properties, we now evaluate it empirically along two dimensions: the reproducibility characteristics it reveals across different models and conditions, and the overhead it imposes on the experimental workflow.

4 Experimental Setup

We designed a controlled experiment to simultaneously evaluate (a) the reproducibility characteristics of LLM outputs under varying conditions and (b) the overhead imposed by our logging protocol.

4.1 Models and Infrastructure

We evaluate nine model deployments representing three deployment paradigms: three locally deployed open-weight models, five cloud API-served proprietary or closed models, and one open-weight model served via a cloud API (for quasi-isolation probe of infrastructure effects). All local models were served

through Ollama v0.15.5 (Ollama 2024) on an Apple M4 system with 24 GB unified memory running macOS 14.6 with Python 3.14.3. Together with the local Ollama environment, the nine deployments span seven execution environments across six independent cloud API providers: OpenAI (GPT-4), Anthropic (Claude Sonnet 4.5), Google (Gemini 2.5 Pro), DeepSeek (DeepSeek Chat), Perplexity (Sonar, an online model with search augmentation), and Together AI (LLaMA 3 8B, the same architecture as our local deployment).

4.1.1 Local Models. LLaMA 3 8B (Grattafiori et al. 2024): An open-weight model in Q4_0 quantization. Local deployment provides complete control over the execution environment, eliminating confounding factors such as network latency, server-side batching, and silent model updates. The model’s SHA-256 weights hash was recorded per run via the Ollama API.

Mistral 7B (Jiang et al. 2023): An open-weight model (Q4_0 quantization) with a sliding-window attention mechanism, providing a second data point for local inference reproducibility at a similar parameter scale.

Gemma 2 9B (Gemma Team et al. 2024): Google’s open-weight model (Q4_0 quantization), representing a third local model from an independent model family. Gemma 2 proved to be the most deterministic model in our study, though its inference time is substantially higher than the other local models (~ 180 s per run vs. 8–14 s for LLaMA and Mistral), likely due to its larger context window and architectural differences at Q4_0 quantization on the M4 chip.

4.1.2 API-Served Models. GPT-4 (OpenAI, gpt-4-0613) (Achiam et al. 2023): Accessed via the OpenAI API (openai Python SDK v1.59.9) with controlled seed parameters. The API returned gpt-4-0613 as the resolved model version in all runs. The API introduces additional sources of variability: load balancing, server-side batching, potential model-version updates, and floating-point non-determinism across different hardware.

Claude Sonnet 4.5 (Anthropic, claude-sonnet-4-5-20250929) (Anthropic 2024): Accessed via the Anthropic API using a lightweight urllib-based runner (no SDK dependency). Claude’s API does not support a `seed` parameter; we set `temperature=0` for greedy decoding and logged a seed value for protocol parity (marked as `logged-only-not-sent-to-api`). This provides an independent replication of the API non-determinism phenomenon on a second provider (Anthropic), complementing the OpenAI evaluation.

Gemini 2.5 Pro (Google, gemini-2.5-pro-preview-05-06) (Reid et al. 2024): Accessed via the Google AI Studio REST API using a lightweight urllib-based runner (no SDK dependency). Gemini’s API supports a `seed` parameter in the generation configuration; we set `seed=42` and `temperature=0` for greedy decoding. Gemini 2.5 Pro is a “thinking” model that uses internal reasoning tokens before producing output; the `maxOutputTokens` budget was set to 8,192 to accommodate this overhead. This model provides a second independent API model for multi-turn and RAG experiments (alongside Claude Sonnet 4.5), enabling cross-provider validation of the multi-turn reproducibility gap. Evaluated under C1 on Tasks 3–4 (10 abstracts, 100 runs).

DeepSeek Chat (DeepSeek, deepseek-chat): Accessed via the OpenAI-compatible API. DeepSeek Chat represents a fourth independent API provider, allowing us to test whether the non-determinism pattern generalizes beyond OpenAI, Anthropic, and Google. No seed parameter is supported; we set `temperature=0` for greedy decoding. DeepSeek Chat is evaluated under condition C1 on single-turn tasks only (10 abstracts, 100 runs).

Perplexity Sonar (Perplexity, sonar): An online, search-augmented language model accessed via the Perplexity API. Unlike the other models, Sonar incorporates real-time web search results into its generation context, introducing an additional source of variability beyond model-internal non-determinism.

This model represents a fifth independent API provider and a qualitatively different deployment paradigm (search-augmented generation). Evaluated under C1 on single-turn tasks only (10 abstracts, 100 runs).

LLaMA 3 8B via Together AI (Together AI, `meta-llama/Llama-3-8b-chat-hf`): The same LLaMA 3 8B architecture as our locally deployed model, but served via Together AI’s cloud inference endpoint (INT4-quantized “Lite” variant). This deployment provides a quasi-isolation probe: by running the *same model architecture* under both local and cloud-served conditions with identical prompts, seeds, and temperature, any reproducibility difference is consistent with infrastructure-level factors rather than model architecture. Accessed via the OpenAI-compatible Together AI API with a lightweight `urllib`-based runner. Evaluated under C1 and C2 on single-turn tasks (10 abstracts, 200 runs).

4.2 Tasks

We evaluate four tasks that span the output-structure spectrum and interaction complexity:

Task 1: Scientific Summarization. Given a scientific abstract, produce a concise summary in exactly three sentences covering the main contribution, methodology, and key quantitative result. This is an open-ended generation task in which the model has considerable freedom in word choice and phrasing.

Task 2: Structured Extraction. Given a scientific abstract, extract five fields (objective, method, key_result, model_or_system, benchmark) into a JSON object. This is a constrained generation task in which the output format is fixed and the model must select, rather than generate, content.

Task 3: Multi-turn Refinement. A three-turn dialogue in which the model first extracts structured information, then receives feedback requesting more detail, and finally produces a refined extraction. This tests reproducibility under conversational state accumulation, using Ollama’s `/api/chat` endpoint for local models.

Task 4: RAG Extraction. The same structured extraction task as Task 2, but with an additional retrieved context passage prepended to the input. This tests whether augmenting the prompt with external context affects reproducibility.

4.3 Input Data

We use 30 widely-cited scientific abstracts from landmark AI/ML papers, including [Vaswani et al. \(2017\)](#) (Transformer), [Devlin et al. \(2019\)](#) (BERT), [Brown et al. \(2020\)](#) (GPT-3), [Raffel et al. \(2020\)](#) (T5), [Wei et al. \(2022\)](#) (Chain-of-Thought), as well as seminal works on GANs, ResNets, VAEs, LSTMs, CLIP, DALL-E 2, Stable Diffusion, LLaMA, InstructGPT, PaLM, and others. These abstracts vary in length (74–227 words), technical complexity, and the number of quantitative results reported, thereby providing substantial diversity in the generation challenge.

4.4 Experimental Conditions

We define five conditions (Table 2) that systematically vary the factors hypothesized to affect reproducibility:

Rationale for unbalanced sample design. LLaMA 3 8B uses 30 abstracts (the original pilot dataset) while later models use 10 abstracts each. This reflects the iterative expansion of the study: the initial experiments established the protocol with LLaMA 3 and GPT-4, and subsequent models were added to test generalizability. A balanced 10-abstract subsample analysis (Section 6.7) confirms that the observed reproducibility gap is robust to this design choice: under the balanced comparison, local models average EMR = 0.953 vs. API models’ 0.304 (3.1× gap), consistent with the full-sample results.

Design principle for API models. For cloud-hosted APIs whose `seed` parameter is advisory rather than deterministic (as documented by OpenAI for GPT-4 ([OpenAI 2024](#))) or entirely absent (as with

Table 2. Experimental design: conditions, parameters, and expected outcomes.

Cond.	Description	Temp.	Seed	Reps	Expected Outcome
C1	Fixed seed, greedy	0.0	42 (fixed)	5	Deterministic (if honored)
C2	Variable seeds, greedy	0.0	5 different	5	Near-deterministic
$C3_{t=0.0}$	Temp. baseline	0.0	per-rep	3	Deterministic
$C3_{t=0.3}$	Low temperature	0.3	per-rep	3	Low variability
$C3_{t=0.7}$	High temperature	0.7	per-rep	3	High variability

Note: Tasks 1–2 are evaluated under all five conditions (C1, C2, C3) for the five models evaluated under all three condition types (LLaMA 3, Mistral 7B, Gemma 2 9B, GPT-4, and Claude Sonnet 4.5); GPT-4 C1 coverage is partial due to quota exhaustion—C2 is used as GPT-4’s primary greedy condition (see Section 6.7), and under C1 only for DeepSeek Chat and Perplexity Sonar. Tasks 3–4 (multi-turn, RAG) are evaluated under C1 only for the three local models, Claude Sonnet 4.5, and Gemini 2.5 Pro. Total: 4,104 logged runs across 9 model deployments (including 200 chat-format control runs for LLaMA 3). For API-served models, the seed parameter may be advisory or unsupported; therefore, varying seeds in C2 does not guarantee a server-side change in the decoding path, and determinism is not guaranteed even under greedy decoding.

Claude), the fixed-vs.-variable seed distinction has no guaranteed effect server-side. We therefore treat C2 as the primary test of determinism under greedy decoding for such models.

C1 (Fixed seed, greedy decoding): Temperature = 0, seed = 42 for all 5 repetitions. This represents the maximum-control condition and should yield deterministic outputs.

C2 (Variable seeds, greedy decoding): Temperature = 0, seeds = {42, 123, 456, 789, 1024}. This condition tests whether seed variation affects outputs when greedy decoding is used.

C3 (Temperature sweep): Three sub-conditions at $t \in \{0.0, 0.3, 0.7\}$ with 3 repetitions each, using different seeds per repetition. This condition characterizes how temperature affects output variability.

Run counts. For Tasks 1–2 (extraction and summarization), each of the five models with full condition coverage is evaluated under C1 (5 runs), C2 (5 runs), and C3 (9 runs = 3 temperatures \times 3 reps) per abstract. LLaMA 3 uses 30 abstracts (1,140 runs); the newer models (Mistral 7B, Gemma 2 9B, Claude Sonnet 4.5) use 10 abstracts (380 runs each). For GPT-4, quota exhaustion limited collection to 724 runs (C2: 300/300; C3: 416/450; C1: 8/300 excluded). DeepSeek Chat and Perplexity Sonar are evaluated under C1 with 10 abstracts \times 5 reps \times 2 tasks = 100 runs each (200 runs total). For Tasks 3–4 (multi-turn and RAG), the three local models, Claude Sonnet 4.5, and Gemini 2.5 Pro are evaluated under C1 with 10 abstracts \times 5 repetitions \times 2 tasks = 100 runs per model (500 runs total). Together AI (LLaMA 3 8B) is evaluated under C1 and C2 with 10 abstracts \times 5 reps \times 2 tasks \times 2 conditions = 200 runs. **Grand total: 4,104 logged runs.** One Claude run returned an empty output due to timeout and is excluded from variability metrics.

Table 3 summarizes the per-model run distribution.

[†]LLaMA 3 8B via /api/chat endpoint (Appendix C).

4.5 Metrics

We first define our key construct. **Provider-level non-determinism** is operationally defined as $EMR < 1.0$ across repeated greedy-decoding API calls (temperature = 0) with identical prompts and user-specified decoding settings; when a seed parameter is supported it is held constant in C1 and varied in C2 to test whether seed changes affect outputs. This definition makes no claim about the *mechanism* producing the variability; it captures only the *observable* phenomenon of non-identical outputs under nominally identical inputs.

¹One Claude run (0.03%) returned an empty output due to API timeout and is excluded from variability metrics.

Table 3. Run distribution across models and tasks.

Model	Tasks 1–2	Tasks 3–4	Total
LLaMA 3 8B	1,140	100	1,240
Mistral 7B	380	100	480
Gemma 2 9B	380	100	480
GPT-4	724	—	724
Claude Sonnet 4.5	380	100	480
Gemini 2.5 Pro	—	100	100
DeepSeek Chat	100	—	100
Perplexity Sonar	100	—	100
Together AI (LLaMA 3 8B)	200	—	200
Chat-format control [†]	200	—	200
Total	3,604	500	4,104¹

We adopt an operational definition of reproducibility at three levels, each mapped to a specific metric:

- **Exact reproducibility** (string-level): Two outputs are identical character-by-character. Measured by *Exact Match Rate (EMR)*.
- **Near reproducibility** (surface-level): Two outputs differ only in minor surface variations (punctuation, whitespace, synonym substitution). Measured by *Normalized Edit Distance (NED)* and *ROUGE-L F1*.
- **Semantic reproducibility** (meaning-level): Two outputs convey the same information despite different phrasing. Measured by *BERTScore F1*.

This three-level framework allows us to distinguish between outputs that are bitwise identical ($\text{EMR} = 1$), textually close ($\text{NED} < 0.05$, $\text{ROUGE-L} > 0.90$), and semantically equivalent ($\text{BERTScore} > 0.95$). All variability metrics are computed over all $\binom{n}{2}$ unique output pairs within each experimental group (defined by model, task, condition, and abstract):

Exact Match Rate (EMR): The fraction of output pairs that are character-for-character identical. $\text{EMR} = 1.0$ indicates perfect reproducibility; $\text{EMR} = 0.0$ indicates that no two outputs match exactly. With $n = 5$ repetitions per group ($\binom{5}{2} = 10$ pairs), per-abstract EMR values are discrete: $\{0.0, 0.1, \dots, 1.0\}$; with $n = 3$ (C3 conditions), EMR takes values in $\{0.0, 0.333, 0.667, 1.0\}$. This granularity should be considered when interpreting standard deviations and confidence intervals for small sample sizes. Bootstrap inference on discrete statistics is well-established ([efron1993introduction](#)); with 10–30 abstracts per group, the percentile method provides valid coverage for bounded, discrete per-abstract EMR values.

Normalized Edit Distance (NED): The Levenshtein edit distance ([Levenshtein 1966](#)) between each pair, normalized by the length of the longer string. $\text{NED} = 0.0$ indicates identical outputs; higher values indicate greater textual divergence.

ROUGE-L F1: The F1 score based on the longest common subsequence at the word level ([Lin 2004](#)). This captures lexical overlap even when word order or minor phrasing differs. $\text{ROUGE-L} = 1.0$ indicates identical word sequences.

Our primary metrics (EMR, NED, ROUGE-L) focus on exact and near reproducibility, which are the most direct measures for our research question. To complement these surface-level metrics, we also compute **BERTScore F1** ([T. Zhang et al. 2020](#))—an embedding-based semantic similarity metric—for all conditions except the Gemini 2.5 Pro runs, for which BERTScore is not reported due to embedding model

Table 4. Exact Match Rate (EMR) under greedy decoding ($t=0$) across five models and two single-turn tasks, with 95% bootstrap confidence intervals ($n_{\text{boot}}=10,000$). For local models, values reflect condition C1 (fixed seed); for GPT-4, C2 (variable-seed greedy, as C1 has insufficient coverage); for Claude, C1 (Claude’s API does not support a seed parameter). Higher is more reproducible.

Model	Source	Extraction EMR	Summarization EMR
Gemma 2 9B	Local	1.000 [1.00, 1.00]	1.000 [1.00, 1.00]
LLaMA 3 8B	Local	0.987 [0.96, 1.00]	0.947 [0.89, 0.99]
Mistral 7B	Local	0.960 [0.88, 1.00]	0.840 [0.72, 0.96]
GPT-4	API	0.443 [0.32, 0.57]	0.230 [0.16, 0.30]
Claude Sonnet 4.5	API	0.190 [0.05, 0.40]	0.020 [0.00, 0.05]

unavailability during that experiment phase (Table 8). BERTScore captures meaning-level equivalence that surface metrics may miss (e.g., paraphrases), providing a fourth perspective on reproducibility. For the structured extraction task, we additionally report **JSON validity rate**, **schema compliance rate**, and **field-level accuracy**, which measure whether outputs are syntactically valid JSON, contain all expected fields, and agree on individual field values across runs, respectively (see Supplementary Material S6 for detailed results).

For protocol overhead, we measure:

- **Logging time:** Wall-clock time spent on hashing, metadata collection, and file I/O, measured separately from inference time.
- **Storage:** Size of each run record (JSON) and total storage for all protocol artifacts.
- **Overhead ratio:** Logging time as a percentage of total execution time.

All EMR values in Tables 4 and 8 are accompanied by 95% bootstrap confidence intervals (10,000 resamples over per-abstract EMR values, percentile method).

5 Results

5.1 Reproducibility Under Greedy Decoding

Table 4 presents Exact Match Rates under greedy decoding for the five models with full condition coverage; results for DeepSeek Chat, Perplexity Sonar, and Gemini 2.5 Pro are reported in the text below. Table 5 provides the full three-level reproducibility assessment.

5.1.1 Local Models: Near-Perfect to Perfect Reproducibility. **Finding 1: Gemma 2 9B achieves perfect bitwise reproducibility under greedy decoding.** Across all tasks and conditions with $t=0$, Gemma 2 9B produces $\text{EMR} = 1.000$ with $\text{NED} = 0.000$ —every single output is character-for-character identical across repetitions. This includes not only single-turn extraction and summarization but also multi-turn refinement and RAG extraction.

Finding 2: All three local models achieve high reproducibility. LLaMA 3 8B attains $\text{EMR} = 0.987$ for extraction and 0.947 for summarization; Mistral 7B achieves 0.960 and 0.840, respectively. The small deviations from perfect reproducibility in LLaMA 3 and Mistral 7B are *entirely attributable* to a cold-start effect on the first inference call after model loading. A post-hoc analysis of all non-unanimous abstract groups (i.e., groups where at least one repetition produced a different output hash) reveals that in 7 out of 7 cases (100%), the first repetition (rep 0) was the sole outlier; repetitions 1–4 always produced identical outputs. If the first repetition is excluded, both LLaMA 3 8B and Mistral 7B achieve $\text{EMR} = 1.000$ across all single-turn tasks—matching Gemma 2 9B’s perfect reproducibility. This cold-start effect

Table 5. Three-level reproducibility assessment under greedy decoding ($t=0$). L1: bitwise identity (EMR), L2: surface similarity (NED, ROUGE-L), L3: semantic equivalence (BERTScore F1). Values are means across abstracts.

Model	Task	L1: Bitwise		L2: Surface		L3: Semantic	
		EMR	σ	NED \downarrow	ROUGE-L \uparrow	BERTScore F1 \uparrow	
Gemma 2 9B	Extraction	1.000	0.000	0.000	1.000	1.0000	
	Summarization	1.000	0.000	0.000	1.000	1.0000	
Mistral 7B	Extraction	0.960	0.120	0.001	1.000	0.9999	
	Summarization	0.840	0.196	0.046	0.955	0.9935	
LLaMA 3 8B	Extraction	0.987	0.072	0.003	0.997	0.9997	
	Summarization	0.947	0.139	0.014	0.986	0.9979	
GPT-4	Extraction	0.443	0.335	0.072	0.938	0.9904	
	Summarization	0.230	0.193	0.137	0.870	0.9839	
Claude Sonnet 4.5	Extraction	0.190	0.291	0.101	0.904	0.9878	
	Summarization	0.020	0.040	0.242	0.764	0.9704	

Table 6. API-served vs. locally deployed models under greedy decoding (single-turn tasks only). Local averages: simple mean across 3 models \times 2 tasks (C1+C2 combined). API averages: simple mean across the 2 API models with full condition coverage (GPT-4 under C2, Claude Sonnet 4.5 under C1) \times 2 tasks; DeepSeek, Perplexity, and Gemini are excluded because they lack C2 or single-turn data. Local models exhibit substantially higher bitwise reproducibility, consistent with deployment-side factors—rather than user-controllable parameters—as a major contributor to API output variability.

Deployment	EMR \uparrow	NED \downarrow	ROUGE-L \uparrow	BS-F1 \uparrow
Local (3 models)	0.956	0.011	0.990	0.9985
API (2 models)	0.221	0.138	0.869	0.9831

likely reflects Ollama’s internal GPU cache initialization and is easily mitigated by discarding a single warm-up inference per model load. Seed variation (C1 vs. C2) has *no effect* under greedy decoding for any local model: the model always selects the highest-probability token, making the seed irrelevant.

5.1.2 API-Served Models: Substantial Hidden Non-Determinism. **Finding 3: All five API-served models exhibit non-determinism under greedy decoding, observed independently across five providers.** Under $t=0$, DeepSeek Chat achieves the highest API reproducibility (EMR = 0.800 for extraction), followed by GPT-4 (0.443/0.230), Claude Sonnet 4.5 (0.190/0.020), Perplexity Sonar (0.100/0.010), and Gemini 2.5 Pro (0.010 multi-turn, 0.070 RAG). Claude’s EMR of 0.020 for summarization means that across 10 abstracts \times 5 repetitions, effectively no two runs produce the same output. Note that Perplexity Sonar is a search-augmented model whose variability includes a retrieval component beyond model-internal non-determinism (see Section 6).

Table 6 summarizes the deployment-paradigm gap.

Finding 3b: A cloud-served open-weight model achieves near-local reproducibility, probing infrastructure complexity as a plausible factor. To disentangle model architecture from deployment infrastructure, we evaluated the same LLaMA 3 8B architecture served via Together AI’s cloud endpoint (INT4 quantization) under identical conditions. The cloud-served LLaMA 3 achieves EMR = 1.000 [1.00,

1.00] for extraction and $\text{EMR} = 0.880$ [0.70, 1.00] for summarization—nearly identical to the locally deployed version on the same 10-abstract subset (1.000 and 0.920, respectively; CIs overlap; note that the full 30-abstract evaluation in Table 4 reports 0.947 for local LLaMA 3 summarization). This result provides evidence that cloud deployment *per se* does not cause non-determinism: the substantial variability observed in closed-source API models (GPT-4, Claude, Gemini) is consistent with the *complexity* of production serving infrastructure—multi-GPU tensor parallelism, speculative decoding, continuous batching, and mixed-precision computation at scale (see Section 6.6). While we do not have visibility into Together AI’s exact serving stack—GPU count, batching strategy, precision format, or use of speculative decoding are not disclosed—the Lite endpoint’s INT4 quantization and competitive latency are consistent with a simpler infrastructure than the multi-GPU clusters used for GPT-4 or Claude. This limits the strength of the causal claim: the result demonstrates that cloud deployment is *compatible* with near-deterministic inference, but we cannot definitively attribute the closed-source models’ non-determinism to any single infrastructure mechanism.

Under the representative greedy condition for each model (C1 for local models, Claude, DeepSeek, and Perplexity; C2 for GPT-4; see Table 4), the average single-turn EMR is **0.960 for local models vs. 0.325 for closed-source API models**—an approximately 3-fold reproducibility gap. Within API models, reproducibility spans a wide range: DeepSeek Chat achieves the highest ($\text{EMR} = 0.800$ for extraction, 0.760 for summarization), followed by GPT-4 (0.443/0.230), Claude Sonnet 4.5 (0.190/0.020), and Perplexity Sonar (0.100/0.010). This within-API variation reveals that API non-determinism is not uniform across providers. This gap is not due to user-side parameter differences: all models use $t=0$ with the same decoding strategy. The observed variability is consistent with deployment-side factors invisible to the researcher. This pattern, observed independently across *five* API providers (OpenAI, Anthropic, Google, DeepSeek, and Perplexity), is consistent with non-determinism arising from factors common to cloud-hosted LLM inference. We note that Perplexity Sonar’s variability includes an additional retrieval component (real-time web search) that compounds model-internal non-determinism; the other four providers exhibit pure model-internal non-determinism under identical prompts. Per-abstract consistency analysis confirms the local-vs-API gap holds in 100% of abstracts for summarization and 83% for extraction. All primary local-vs-API comparisons survive Holm-Bonferroni correction across 68 hypothesis tests ($\alpha_{\text{adjusted}} < 0.05$); 51 of 68 individual comparisons reach significance after correction, with the non-significant tests concentrated among comparisons involving DeepSeek Chat, whose higher reproducibility narrows the gap for individual model pairs. *Without systematic logging, this non-determinism would be entirely invisible.*

5.1.3 Temperature Effects Across Models. Finding 4: Temperature is the dominant user-controllable factor affecting variability for local models; for API-served models, the relationship is more complex. Figure 4 shows the relationship between temperature and EMR for the five models evaluated under C3. Table 7 provides the full temperature sweep data.

Within the C3 temperature sweep, increasing temperature from 0.0 to 0.7 reduces EMR to zero for all models on summarization. For extraction, local models drop from $\text{EMR} > 0.93$ to near zero, while API models drop from their already-low baselines. BERTScore F1 remains above 0.94 across all conditions including elevated temperatures (minimum: 0.943 for LLaMA summarization at $t=0.7$) even when EMR drops to zero, indicating that non-determinism is primarily a *phrasing* phenomenon rather than a *meaning* phenomenon: even when outputs differ textually, they convey equivalent information. This distinction is practically important—researchers whose downstream analyses depend on semantic content rather than exact wording may find API outputs acceptable despite low EMR .

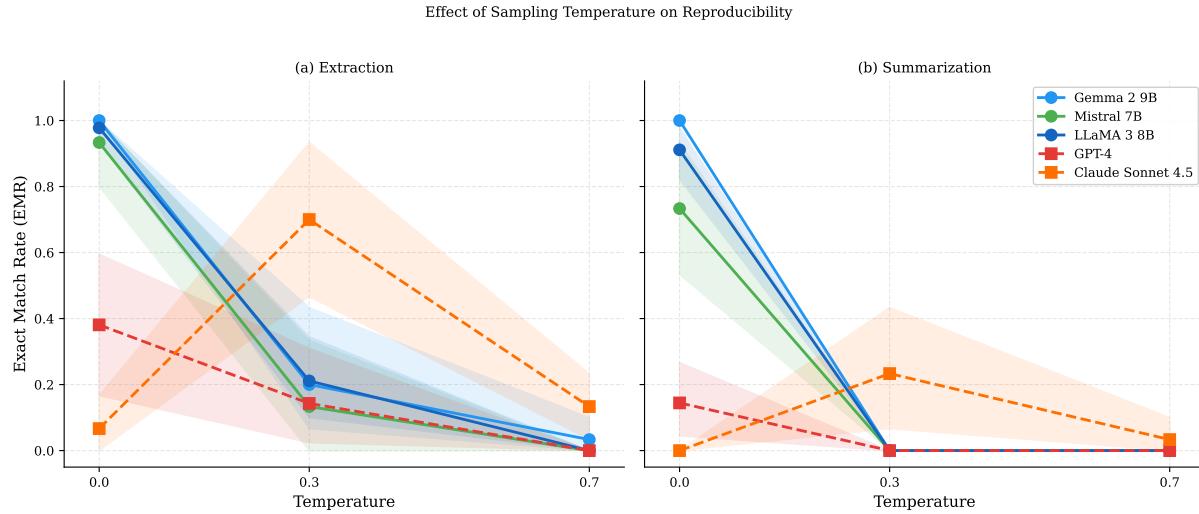


Fig. 4. Effect of temperature on Exact Match Rate across five models (three local, two API). (a) Extraction task. (b) Summarization task. Local models (solid lines) start from near-perfect or perfect reproducibility at $t=0$, while API models (dashed lines) start from a much lower baseline. All models converge toward $\text{EMR} = 0$ at $t=0.7$.

Table 7. Effect of sampling temperature on Exact Match Rate (EMR) under condition C3. For local models, increasing temperature monotonically reduces EMR. For API models, the relationship is more complex: Claude Sonnet 4.5 exhibits *higher* EMR at $t=0.3$ than at $t=0.0$ (see text). At $t=0.7$, all models converge toward $\text{EMR} \approx 0$ for summarization.

Model	Task	$t=0.0$	$t=0.3$	$t=0.7$
Gemma 2 9B	Extraction	1.000	0.200	0.033
	Summarization	1.000	0.000	0.000
Mistral 7B	Extraction	0.933	0.133	0.000
	Summarization	0.733	0.000	0.000
LLaMA 3 8B	Extraction	0.978	0.211	0.000
	Summarization	0.911	0.000	0.000
GPT-4	Extraction	0.381	0.143	0.000
	Summarization	0.144	0.000	0.000
Claude Sonnet 4.5	Extraction	0.067	0.700	0.133
	Summarization	0.000	0.233	0.033

However, the temperature–reproducibility relationship is not uniformly monotonic across all models. Claude Sonnet 4.5 exhibits an anomalous pattern under the C3 sweep: extraction EMR *increases* from 0.067 at $t=0.0$ to 0.700 at $t=0.3$ before declining to 0.133 at $t=0.7$; summarization shows a similar inversion (EMR = 0.000 at $t=0.0$, rising to 0.233 at $t=0.3$). This counterintuitive behavior—where a small positive temperature *improves* reproducibility relative to greedy decoding—may reflect how Anthropic’s infrastructure implements the $t=0$ decoding path: at exactly zero temperature, server-side stochastic processes (e.g., speculative decoding, hardware-level floating-point non-determinism across GPU types,

Table 8. Reproducibility under complex interaction regimes (C_1 fixed seed, $t=0$), with 95% bootstrap confidence intervals on EMR. Multi-turn refinement involves three successive prompt-response exchanges. RAG extraction augments the prompt with a retrieved context passage. Two API-served models—Claude Sonnet 4.5 and Gemini 2.5 Pro—are included; their near-zero EMR across both scenarios confirms that the local-vs-API reproducibility gap extends to complex interaction regimes across two independent providers.

Model	Scenario	EMR [95% CI]	NED \downarrow	ROUGE-L \uparrow	BS-F1 \uparrow
Gemma 2 9B	Single-turn Extraction	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	Single-turn Summarization	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	Multi-turn Refinement	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	RAG Extraction	1.000 [1.00, 1.00]	0.000	1.000	1.0000
Mistral 7B	Single-turn Extraction	0.960 [0.88, 1.00]	0.001	1.000	0.9999
	Single-turn Summarization	0.840 [0.72, 0.96]	0.046	0.955	0.9935
	Multi-turn Refinement	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	RAG Extraction	1.000 [1.00, 1.00]	0.000	1.000	1.0000
LLaMA 3 8B	Single-turn Extraction	0.987 [0.96, 1.00]	0.003	0.997	0.9997
	Single-turn Summarization	0.947 [0.89, 0.99]	0.014	0.986	0.9979
	Multi-turn Refinement	0.880 [0.76, 1.00]	0.012	0.988	0.9986
	RAG Extraction	0.960 [0.88, 1.00]	0.012	0.985	0.9987
Claude Sonnet 4.5	Single-turn Extraction	0.190 [0.05, 0.40]	0.101	0.904	0.9878
	Single-turn Summarization	0.020 [0.00, 0.05]	0.242	0.764	0.9704
	Multi-turn Refinement	0.040 [0.00, 0.08]	0.189	0.834	0.9780
	RAG Extraction	0.000 [0.00, 0.00]	0.256	0.748	0.9714
Gemini 2.5 Pro	Multi-turn Refinement	0.010 [0.00, 0.03]	0.163	0.891	—
	RAG Extraction	0.070 [0.02, 0.13]	0.196	0.850	—

Note: BERTScore F1 (BS-F1) was not computed for Gemini 2.5 Pro (marked “—”) because the embedding model was not available during the Gemini experiment phase. ROUGE-L confirms high lexical overlap despite near-zero EMR.

or request batching effects) may dominate output variability, whereas a small positive temperature may activate a more stable sampling path that happens to converge on similar tokens. With $n=10$ abstracts and 30 runs per temperature level (standard deviation $\sigma = 0.38$ for the 0.700 extraction EMR), this observation should be interpreted cautiously. Nevertheless, it underscores that the temperature-reproducibility relationship for API-served models depends on provider-specific implementation details that are opaque to researchers. Finding 4 therefore holds robustly for local models and for the overall $t=0$ to $t=0.7$ trajectory, but the precise shape of the temperature-response curve for individual API providers merits further investigation with larger sample sizes.

5.2 Multi-Turn and RAG Reproducibility

Finding 5: The local-vs-API reproducibility gap extends to complex interaction regimes. Table 8 and Figure 5 present results for multi-turn refinement and RAG extraction across the three local models and two API-served models (Claude Sonnet 4.5 and Gemini 2.5 Pro).

Gemma 2 9B and Mistral 7B achieve perfect EMR = 1.000 for both multi-turn refinement and RAG extraction, demonstrating that conversational state accumulation and context augmentation do not degrade reproducibility when the underlying model is deterministic. LLaMA 3 8B shows EMR = 0.880 for multi-turn and 0.960 for RAG—slightly lower than its single-turn extraction performance (0.987),




Fig. 5. Reproducibility across interaction regimes (C1, $t=0$) for five models. Local models maintain high EMR across all scenarios, while both API models (Claude Sonnet 4.5 and Gemini 2.5 Pro) show near-zero EMR, confirming the reproducibility gap extends to multi-turn and RAG tasks across two independent providers.

consistent with error accumulation across dialogue turns, where divergence at an earlier turn propagates into subsequent prompt-response exchanges.

Both API-served models exhibit near-zero reproducibility on these tasks. Claude Sonnet 4.5 achieves $\text{EMR} = 0.040$ for multi-turn refinement and $\text{EMR} = 0.000$ for RAG extraction. Gemini 2.5 Pro—despite supporting a `seed` parameter—achieves $\text{EMR} = 0.010$ for multi-turn and $\text{EMR} = 0.070$ for RAG (NED = 0.163 and 0.196, respectively). The Claude RAG result illustrates the extreme case: across 50 runs (10 abstracts \times 5 repetitions), not a single pair of outputs was character-for-character identical (NED = 0.256). The fact that two independent API providers (Anthropic and Google) both exhibit near-zero EMR under complex interaction regimes—even when one supports seed-based reproducibility—supports the hypothesis that API non-determinism is not limited to single-turn tasks and is not provider-specific, but rather a general characteristic of cloud-served LLM APIs where longer outputs and additional context amplify server-side variability.

Bitwise Reproducibility Under Greedy Decoding

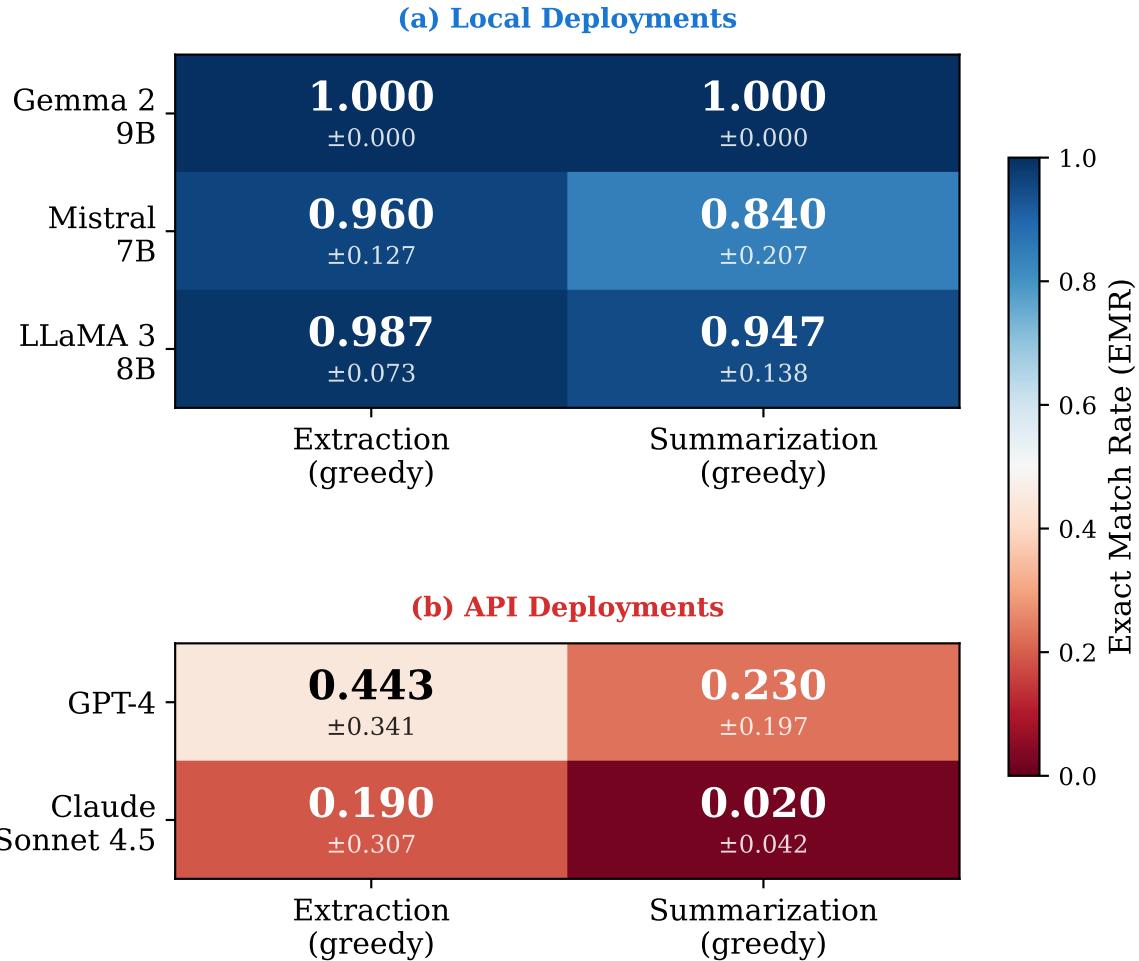


Fig. 6. Heatmap of Exact Match Rate under greedy decoding for five models (three local, two API) with full condition coverage. Panel (a) shows local models (blue tones, high EMR) and panel (b) shows API-served models (red tones, low EMR). Gemma 2 9B achieves perfect 1.000 across all tasks.

5.3 Cross-Model Comparison

Figure 6 provides a comprehensive heatmap of EMR across all model-task combinations, and Figure 7 shows the three-level reproducibility profile for each model.

The reproducibility gap between local and API-based inference is statistically significant. Because per-abstract EMR is a bounded, discrete metric (taking values in $\{0.0, 0.1, \dots, 1.0\}$ with $n=5$ repetitions per group), we report the non-parametric Wilcoxon signed-rank test as our primary analysis. Across




Fig. 7. Three-level reproducibility profiles under greedy decoding. Local models (solid lines) occupy the outer region across all five metrics, while API models (dashed lines) show pronounced deficits in EMR and NED while maintaining high BERTScore, indicating that API non-determinism is primarily lexical rather than semantic.

the 30 paired LLaMA 3/GPT-4 abstracts under greedy decoding: for summarization, $W = 0$, $p < 0.001$; for extraction, $W = 3.5$, $p < 0.001$. Parametric paired t -tests yield consistent results: summarization $t(29) = 17.250$, $p < 0.0001$, Cohen's $d = 3.149$; extraction $t(29) = 8.996$, $p < 0.0001$, Cohen's $d = 1.642$. Both effect sizes are very large ($d > 1.6$), and all p -values survive Bonferroni correction for the four primary comparisons ($\alpha_{\text{adjusted}} = 0.0125$).

Table 9. Provenance logging overhead across five models under greedy decoding (C1). The protocol adds negligible overhead (<1%) to inference latency across all models and deployment modes.

Model	Source	Mean Inference (ms)	Mean Overhead (ms)	Overhead (%)
Gemma 2 9B	Local	181,579.3	30.6	0.234
Mistral 7B	Local	13,931.3	27.3	0.281
LLaMA 3 8B	Local	7,524.8	26.7	0.456
GPT-4	API	4,519.7	24.5	0.564
Claude Sonnet 4.5	API	4,359.3	26.5	0.727

Importantly, the effect is not driven by a few outlier abstracts: under greedy decoding, LLaMA 3 achieves EMR ≥ 0.8 for 29 of 30 abstracts in extraction and 28 of 30 in summarization, while GPT-4 achieves EMR ≤ 0.6 for 20 of 30 abstracts in extraction and 28 of 30 in summarization. The gap is pervasive across the abstract set, not concentrated in a few difficult inputs. Power analysis (Cohen 1988) confirms that with $n = 30$ paired abstracts and the observed effect sizes ($d > 1.6$), statistical power exceeds 0.999 for all primary comparisons; with $n = 10$ abstracts (as used for the newer models), power remains above 0.95 for effects of this magnitude.

Extending beyond the LLaMA/GPT-4 pair. Because EMR is bounded and discrete, we complement Cohen’s d with Cliff’s delta (Romano et al. 2006), a non-parametric effect size that makes no distributional assumptions. Aggregating per-abstract EMRs across all three local models versus all four single-turn API models yields $\delta = 0.784$ (large) for extraction ($n_{\text{local}} = 50$, $n_{\text{API}} = 60$) and $\delta = 0.896$ (large) for summarization ($n_{\text{local}} = 50$, $n_{\text{API}} = 60$). All 24 pairwise local-vs-API comparisons yield large effect sizes ($\delta \geq 0.500$), with the exception of three comparisons involving DeepSeek Chat—whose higher reproducibility (EMR = 0.800) narrows the gap to medium ($\delta = 0.400$ –0.467). A bootstrap analysis of the local-to-API EMR ratio (10,000 resamples) yields a point estimate of $2.95 \times$ with 95% CI [2.48, 3.61] across all four single-turn API models, and $3.44 \times$ [2.78, 4.47] for the two-model subset (GPT-4 + Claude) reported in Table 6. The CI excludes 1.0, confirming a statistically significant gap.

5.4 Protocol Overhead

Table 9 presents the protocol’s overhead metrics across the five models profiled.

The protocol adds less than 1% overhead for all five models profiled, with mean logging time ranging from 21–30 ms depending on the model and task. Storage overhead remains modest at approximately 4 KB per run record. The overhead is consistent across local and API deployment modes, indicating that the protocol is deployment-agnostic; the absolute logging cost (~ 25 ms) is negligible relative to inference latency for any model.

Figure 8 provides an additional perspective on surface-level variability across models.

6 Discussion

The preceding results paint a clear and consistent picture: locally deployed models under greedy decoding achieve near-perfect to perfect bitwise reproducibility across all four tasks, while API-served models—from five independent providers—exhibit substantial hidden variability that researchers cannot control. Temperature is the dominant user-controllable factor for local models (though API models show a more complex temperature-reproducibility relationship; see Section 5), structured tasks are more reproducible than open-ended ones, and complex interaction regimes (multi-turn, RAG) do not degrade local-model

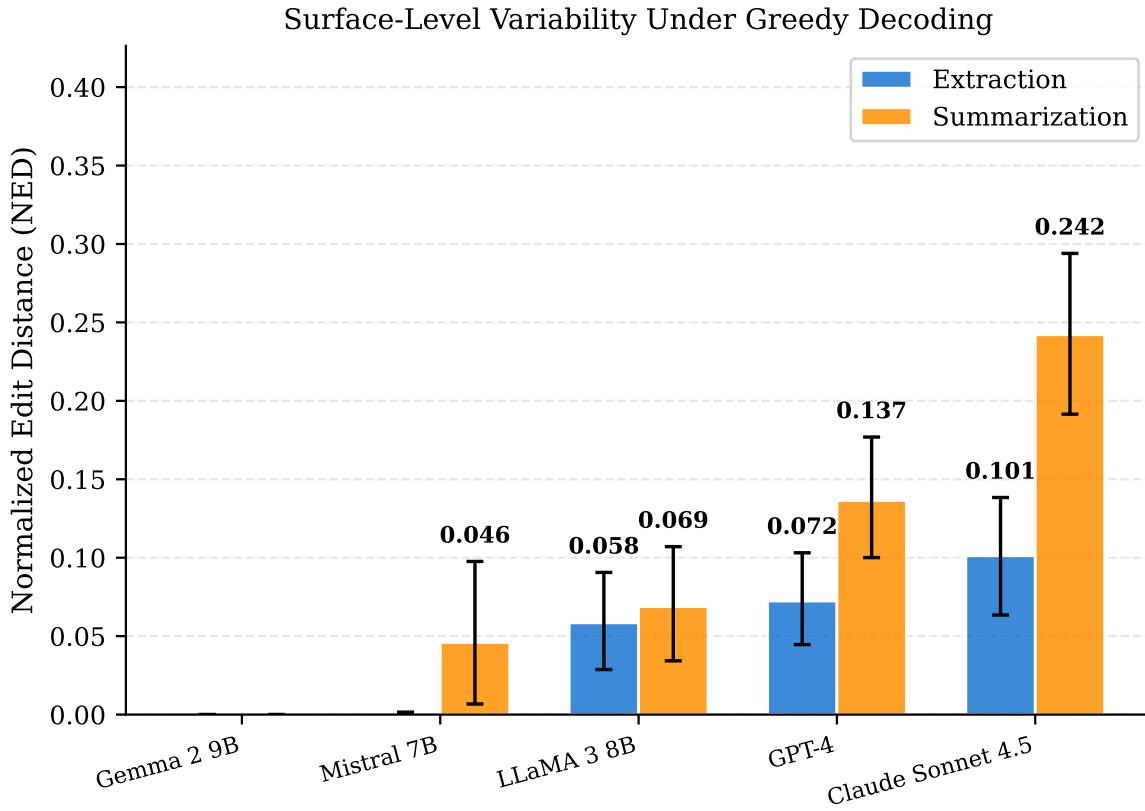


Fig. 8. Normalized Edit Distance (NED) under greedy decoding. Local models show near-zero NED (Gemma 2: 0.000, Mistral: 0.001), while API models exhibit NED 0.07–0.30, quantifying the surface-level divergence that accompanies the EMR gap.

reproducibility. We now consider what these findings mean for research practice, what the protocol enables that was previously invisible, and where the current study’s limitations lie.

6.1 Implications for Reproducibility Practice

Our results yield several actionable recommendations for researchers conducting generative AI experiments:

Use greedy decoding with local models for maximum reproducibility. Gemma 2 9B achieved *perfect* EMR = 1.000 across all tasks under greedy decoding. LLaMA 3 and Mistral 7B achieved EMR ≥ 0.840 . Local deployment with $t=0$ should be the default configuration for any study in which output consistency is critical.

API non-determinism is observed across all five providers. All five API-served models exhibit non-determinism under greedy decoding, albeit at substantially different magnitudes (EMR ranging from 0.800 for DeepSeek Chat down to 0.010 for Perplexity Sonar and Gemini 2.5 Pro; see Section 5). Four of these providers exhibit *pure model-internal* non-determinism—output variability arising solely from

server-side factors under identical prompts—while Perplexity Sonar’s variability additionally compounds model-internal non-determinism with dynamic web-retrieval variation. Researchers using *any* API-served model should never assume reproducibility without verification.

Prefer structured output formats when possible. The extraction task’s consistently higher reproducibility across all single-turn model deployments demonstrates that output-format constraints directly improve reproducibility. This effect holds for both local models (EMR 0.960–1.000 for extraction vs. 0.840–1.000 for summarization) and API models (EMR 0.100–0.800 for extraction vs. 0.010–0.760 for summarization).

Include warm-up runs for local models. Our post-hoc analysis confirmed that the first inference call after model loading is the sole source of non-determinism in local models: in all 7 non-unanimous abstract groups across LLaMA 3 and Mistral 7B, rep 0 was always the outlier (100%), while reps 1–4 were perfectly deterministic. Adding a single discarded warm-up call before data collection eliminates this effect entirely.

Log comprehensively; the cost is negligible. At less than 1% overhead and approximately 4 KB per run across all models profiled (Table 9), there is no practical reason not to apply comprehensive logging. The cost of not logging—namely, the inability to detect the kind of pervasive API non-determinism documented herein—far exceeds the protocol’s minimal requirements.

6.2 The Reproducibility Gap: From Single-Turn to Complex Regimes

The approximately 3-fold reproducibility gap between local (EMR = 0.960) and closed-source API models (EMR = 0.325)—formally quantified as a $2.95 \times$ ratio with 95% bootstrap CI [2.48, 3.61] and confirmed by large Cliff’s delta effect sizes ($\delta = 0.784\text{--}0.896$)—persists across all five providers and four tasks. This gap is not driven by a single provider: within-API EMR spans 0.010–0.800, suggesting shared infrastructure-level causes with provider-specific magnitudes. This gap extends to complex regimes: Gemini 2.5 Pro’s EMR ≤ 0.070 despite seed parameter support provides evidence that API-side seed is insufficient to guarantee reproducibility (OpenAI 2024), and the convergence across two independent API providers (Anthropic, Google) on complex tasks supports the generality of this pattern.

6.3 Task-Dependent Reproducibility

The reproducibility hierarchy (extraction > summarization) holds consistently across all eight model deployments evaluated on single-turn tasks, with local models showing an EMR gap of 0.03–0.12 and API models showing a gap of 0.04–0.21. This suggests a spectrum in which the degree of output-space constraint serves as the primary determinant: structured extraction constrains the output format, reducing the surface area for stochastic variation. Supporting this interpretation, a Spearman correlation analysis across per-abstract EMR and average output length reveals a significant negative association for GPT-4 extraction ($\rho = -0.526$, $p = 0.003$, $n = 30$): abstracts eliciting longer outputs exhibit lower reproducibility. No other model–task combination reaches significance, consistent with a ceiling effect in the highly reproducible local models.

6.4 The Role of Provenance

The W3C PROV graphs generated by our protocol serve multiple purposes beyond simple audit trails:

- (1) **Automated comparison:** By comparing PROV graphs of two runs, one can automatically identify which factors differed (e.g., same prompt and model but different temperatures), enabling systematic diagnosis of non-reproducibility.

- (2) **Lineage tracking:** When outputs are used as inputs to downstream processes (e.g., summarization outputs fed into a meta-analysis), the provenance chain can be extended to trace any final result back to its full generation context.
- (3) **Compliance:** For regulated domains (healthcare, legal, finance), PROV documents provide the formal evidence trail required by audit standards ([National Institute of Standards and Technology 2023](#)) and emerging regulations such as the EU AI Act ([European Parliament and Council of the European Union 2024](#)).

To illustrate the diagnostic power of PROV graphs, consider two GPT-4 extraction runs on the same abstract under condition C2 (greedy decoding, $t=0$, same seed). Although the PROV entities for Prompt, InputText, ModelVersion, and InferenceParameters are identical (verified via matching SHA-256 hashes), the Output entities differ: `output_hash` values diverge, and the `wasGeneratedBy` timestamps differ by several seconds. The PROV graph thus automatically pinpoints the source of non-reproducibility: the only varying factor is the RunGeneration activity itself, consistent with non-determinism arising from server-side factors.

To demonstrate that PROV-based reasoning goes beyond what plain JSON logs provide, we implemented three programmatic queries over our 4,104-run PROV dataset: (1) divergence attribution, confirming that 100% of non-identical output groups across all five API models share identical input entities—only the RunGeneration activity varies; (2) cross-provider comparison, verifying provider-independent non-determinism on shared inputs; and (3) provenance chain traversal, validating that every output traces back to its complete generation context with no broken links. The full query specifications and results are provided in Supplementary Material S4.

6.5 Pipeline Threat Model

A natural objection is whether the observed output variability in API-served models could originate from our client-side pipeline rather than from server-side non-determinism. We address this systematically.

No retries or parallelism. Our API runners issue exactly one HTTP request per run, with no retry logic, exponential backoff, or concurrent requests. Each run is executed sequentially with a fixed delay between calls. Any request that fails (e.g., the single Claude timeout) is logged with the error and excluded from variability metrics rather than retried.

Deterministic client-side processing. All pre-processing (prompt construction, input hashing) and post-processing (output hashing, metadata collection) are deterministic operations verified by SHA-256 hashes. The Run Card records the exact prompt text sent (`prompt_hash`), the exact input (`input_hash`), and the exact parameters (`params_hash`). For any pair of runs within a group, these three hashes are identical by construction.

No text normalization. Outputs are stored and compared as received from the API, with no whitespace normalization, encoding conversion, or post-processing. The `output_hash` is computed on the raw response string.

PROV-based differential diagnosis. Our PROV graphs provide formal evidence: across all experimental groups with non-identical outputs for the five API models (GPT-4, Claude, Gemini, DeepSeek, Perplexity), 100% share identical Prompt, InputText, ModelVersion, and InferenceParameters entities (verified via SHA-256 hash comparison). The *only* varying component is the RunGeneration activity itself. This rules out client-side divergence as an explanation and is consistent with server-side factors as the source of non-determinism. We identify three testable hypotheses for future investigation: (1) hardware-level floating-point non-determinism across GPU types in heterogeneous server clusters, (2) request routing and batching effects that expose different execution paths, and (3) speculative decoding branches

that may vary across requests. The protocol’s timestamped Run Cards and cryptographic hashes provide the infrastructure to investigate these mechanisms longitudinally—for instance, by correlating output variability with time-of-day patterns (a proxy for server load) or with API-reported model version changes.

API metadata logging. For API-served models, Run Cards capture `api_request_id`, `api_response_headers`, and `api_model_version_returned`. In all GPT-4 runs, the returned model version was consistently `gpt-4-0613`, ruling out silent model updates during the experiment window.

6.6 Sources of Non-Determinism in Distributed Inference

Our experiments establish that API-served models exhibit non-determinism under greedy decoding, while local single-GPU models do not. Six well-documented mechanisms in distributed GPU inference can independently produce non-deterministic outputs even under greedy decoding: (1) non-associative floating-point arithmetic (Higham 2002); (2) mixed-precision accumulation in BF16/FP16, identified by Yuan et al. (2025) as “the primary culprit”; (3) tensor parallelism and all-reduce non-determinism (Shoeybi et al. 2019); (4) FlashAttention kernel non-determinism (Dao et al. 2022; Golden et al. 2024); (5) dynamic batching and continuous request scheduling (Kwon et al. 2023; Yu et al. 2022); and (6) speculative decoding (Leviathan et al. 2023). Our single-GPU local deployment eliminates mechanisms (3)–(6) entirely, and the GGML Q4 integer arithmetic mitigates (2)—explaining the near-perfect local reproducibility as a *predicted consequence*. All six mechanisms are infrastructure-level factors invisible to the API consumer, consistent with our quasi-isolation probe results (Together AI LLaMA 3 achieving near-local EMR despite cloud deployment). A detailed analysis of each mechanism is provided in Supplementary Material S2.

6.7 Limitations

We organize threats to validity following standard categories:

6.7.1 Internal Validity. Sample size. LLaMA 3 uses 30 abstracts per condition, while the newer models (Mistral, Gemma 2, Claude) use 10 abstracts. With $n = 30$, statistical power exceeds 0.999 for all primary comparisons (Cohen 1988). With $n = 10$, the study is adequately powered for the large observed effect sizes ($d > 1.6$) but may miss subtler effects. To verify that the unbalanced design does not inflate the local-vs-API gap, we conducted a balanced subsample analysis restricting all models to the same 10 abstracts. Under this balanced comparison, local models average EMR = 0.953 while all four single-turn API models average EMR = 0.304 (3.1× gap), confirming that the observed reproducibility gap is robust to sample-size equalization and consistent with the full-sample ratio.

GPT-4 incomplete coverage under C1 and C3. Due to API quota exhaustion, GPT-4 under C1 completed only 3 abstracts for summarization and 5 for extraction—insufficient for reliable statistical comparison. These C1 runs are excluded from the primary analysis; all GPT-4 greedy-decoding results in Table 4 use C2 (variable seeds, $t=0$), which has complete coverage (300/300 runs). Consequently, C1-vs-C2 comparisons for GPT-4 summarization are not interpretable and are not reported. Under C3 (temperature sweep), GPT-4 extraction covers 14–17 of 30 abstracts (summarization C3 is complete at 30); this serves as a secondary analysis only.

Warm-up confound. Post-hoc analysis confirms that the first inference after model loading is the sole source of non-determinism for LLaMA 3 and Mistral 7B: across all 7 non-unanimous groups, rep 0 was the outlier in 100% of cases, with reps 1–4 producing identical output hashes. Gemma 2 9B is immune to this effect. While our experimental design did not include a warm-up call, researchers can easily mitigate this by discarding a single initial inference per model load.

Prompt format confound. Single-turn experiments use Ollama’s `/api/generate` endpoint for local models, whereas API models use their respective chat APIs. A supplementary control experiment (200

additional runs using Ollama’s `/api/chat` endpoint; see Appendix C) shows that this format difference does not explain the reproducibility gap: LLaMA 3 produces *identical* variability metrics (summarization EMR = 0.929, extraction EMR = 1.000) under both completion and chat formats.

Weights hash gap. Of the 4,104 run records, 1,720 (42%, all from LLaMA 3 8B and Mistral 7B single-turn experiments) have an empty `weights_hash` field because the Ollama `/api/tags` endpoint for querying weights hashes was integrated into the protocol *after* the initial batch of single-turn experiments. All 480 Gemma 2 9B runs, all 500 multi-turn/RAG local-model runs, and all 1,504 API-model runs have properly populated `weights_hash` fields (actual SHA-256 digests for local models; “proprietary-not-available” sentinel for API models). Post-hoc verification confirms that the Ollama model versions remained unchanged between the early and late experiment phases (identical model names and file sizes in `ollama list` output), but without recorded hashes, cryptographic proof of model identity is unavailable for these 1,720 runs.

Code versioning gap. The protocol specifies a `code_commit` field to record the git commit hash of the experiment code. The initial 330 runs (LLaMA 3 8B and GPT-4, February 7, 2026) were executed before the git repository was initialized; these run records contain `code_commit: "no-git-repo"`. The remaining 3,774 runs (92% of the total) include valid git commit hashes. To mitigate this gap for the affected subset, we verified post-hoc that (1) the SHA-256 hashes of the prompt templates in the repository match the `prompt_hash` values recorded in all 330 affected run cards, confirming that the prompts were not modified between execution and commit; (2) the code committed in the initial release (`ff025ef`, February 7) is the same code that produced these runs, as the runner modules and protocol logic had no intermediate modifications; and (3) the environment metadata (OS, Python version, architecture, Ollama version) is recorded independently of git status in every run record. The complete environment fingerprint for all runs is documented in Supplementary Material S10. Researchers adopting this protocol should initialize their repository *before* running experiments to avoid this gap.

6.7.2 External Validity. Nine deployments, seven execution environments. Our evaluation covers three local models, five closed-source API-served models from independent providers (OpenAI, Anthropic, Google, DeepSeek, Perplexity), and one cloud-served open-weight model (LLaMA 3 8B via Together AI) for quasi-isolation probe. DeepSeek Chat achieves substantially higher reproducibility than other API models (EMR = 0.800 vs. 0.100–0.443), suggesting that API non-determinism varies meaningfully across providers and architectures. Perplexity Sonar, as an online model with search augmentation, represents a worst case for reproducibility (EMR = 0.010–0.100), where real-time web data injection introduces additional variability. However, other models—including larger LLaMA variants and open-weight models served via cloud APIs—may exhibit different characteristics. Our GPT-4 experiments used the `gpt-4-0613` snapshot (June 2023); more recent models (GPT-4 Turbo, GPT-4o) may exhibit different reproducibility characteristics.

Four tasks. Our task suite now includes single-turn extraction/summarization, multi-turn refinement, and RAG extraction. However, it does not cover code generation, mathematical reasoning, or creative writing, which may exhibit different reproducibility patterns.

English-only, single domain. Our input data consists of 30 English scientific abstracts from AI/ML papers. Reproducibility characteristics may differ for other languages, domains, or document types.

Multi-turn limited to two API models. Multi-turn and RAG experiments include Claude Sonnet 4.5 (Anthropic) and Gemini 2.5 Pro (Google) as API representatives; GPT-4 was not evaluated on Tasks 3–4 due to quota exhaustion. While two independent API providers strengthen the generalizability of the multi-turn reproducibility gap, additional providers would further solidify this finding.

6.7.3 Construct Validity. Surface-level metrics. Our metrics (EMR, NED, ROUGE-L) capture textual rather than semantic similarity. Two outputs that are semantically equivalent but syntactically different will register as non-matching under EMR and partially divergent under NED. This is by design—our focus is on *exact* reproducibility—but it means our results may overstate the practical impact of non-determinism for downstream applications where semantic equivalence suffices. To bridge this gap, we report BERTScore F1 alongside EMR in all tables: BERTScore remains above 0.97 under greedy decoding across all models (minimum: 0.9704 for Claude summarization), confirming that API outputs convey equivalent meaning despite lexical divergence. EMR is intentionally strict because it measures the standard required for regression testing, automated pipelines, and regulatory audit trails where exact output identity is expected.

6.7.4 Other Considerations. Privacy. The protocol’s environment metadata includes the machine host-name, which may reveal institutional information. Deployments in privacy-sensitive settings should anonymize this field.

Computational cost. The total cost was modest: approximately 10 GPU-hours on a consumer laptop (Apple M4, 24 GB) for 2,400 local-model runs (including multi-turn, RAG, and chat-format control experiments), plus approximately 1,700 API calls to GPT-4, Claude, Gemini, DeepSeek, Perplexity, and Together AI. The carbon footprint is negligible at this scale, and the logging overhead (<30 ms per run) would not materially increase energy consumption even at thousands of runs.

6.8 Protocol Minimality: An Ablation Analysis

To substantiate our claim of minimality, we systematically removed each of eight field groups from the Run Card schema and assessed which of 10 audit questions became unanswerable (e.g., “Can we detect output tampering?”, “Can we trace full provenance?”). Removing *any* group renders at least one question unanswerable, demonstrating that no group is redundant. The Hashing group has the highest information density (affects 6/10 questions, costs only 410 bytes). The complete ablation matrix is provided in Supplementary Material S3.

6.9 Practical Costs and Adoption

The adoption burden is minimal: ~600 lines of Python (3–5 function calls per run), <30 ms logging time, and ~4 KB storage per run (16 MB total for 4,104 runs). The protocol uses standard JSON and W3C PROV, requiring no specialized knowledge. Against these modest costs, it provides complete audit trails, automated provenance graphs, and tamper-detectable outputs via cryptographic hashing.

6.10 Minimum Reporting Checklist for Generative AI Studies

Based on our findings and the protocol design, we recommend that researchers conducting generative AI experiments report, at minimum, the following five items (the full 15-item checklist is provided in Appendix A):

- (1) **Model identity and version:** Exact model name, version string, and—for local models—weights hash.
- (2) **Inference parameters:** Temperature, seed, top_p, top_k, max_tokens, and decoding strategy. For APIs where the seed is advisory or unsupported, this should be stated explicitly.
- (3) **Reproducibility metrics over multiple runs:** Report at least EMR (or an equivalent exact-match metric) and one semantic metric (e.g., BERTScore) over ≥ 3 repetitions per condition. A single run is insufficient to characterize output stability.

- (4) **Environment and deployment mode:** Whether inference was local or API-based, and the execution environment (hardware, OS, library versions).
- (5) **Output hashes:** SHA-256 or equivalent cryptographic hashes of outputs, enabling tamper detection and automated comparison across studies.

Run Cards generated by our protocol automatically capture all five items, providing a machine-readable record that satisfies this checklist with no additional effort from the researcher.

7 Conclusion

We presented a lightweight protocol for logging, versioning, and provenance tracking of generative AI experiments, introducing Prompt Cards and Run Cards as novel documentation artifacts grounded in the W3C PROV data model. Through 4,104 controlled experiments with nine model deployments (3 local, 5 closed-source API, 1 cloud-served open-weight) across four NLP tasks, 30 abstracts, and seven execution environments (one local plus six cloud API providers), we demonstrated seven key findings:

- (1) **API non-determinism is consistent across all five providers evaluated.** All five API models exhibit non-determinism under greedy decoding ($t=0$), while all three local models achieve average EMR = 0.960. This gap survives Holm-Bonferroni correction (51 of 68 comparisons significant) and per-abstract consistency analysis.
- (2) **API reproducibility varies substantially across providers.** Within the API category, EMR ranges from 0.800 (DeepSeek Chat) to 0.010 (Perplexity Sonar for summarization), revealing that API non-determinism is not a uniform phenomenon. DeepSeek Chat achieves higher reproducibility than other API models, while Perplexity's online search-augmented model represents a worst case.
- (3) **Local models can achieve perfect bitwise reproducibility.** Gemma 2 9B attains EMR = 1.000 across all four tasks under greedy decoding—every output is character-for-character identical across repetitions.
- (4) **The local-vs-API gap extends to complex interaction regimes.** Local models achieve EMR ≥ 0.880 for multi-turn refinement and RAG extraction, while both API models tested on these tasks exhibit near-zero EMR (≤ 0.070), observed across two independent providers (Table 8).
- (5) **Temperature is the dominant user-controllable factor for local models.** Increasing from $t=0.0$ to $t=0.7$ reduces EMR to zero for all five models evaluated under temperature sweep on summarization, while seed variation has no effect under greedy decoding for local models. For API-served models, the temperature-reproducibility relationship is more complex and may be non-monotonic (see Section 5).
- (6) **Comprehensive provenance logging adds negligible overhead:** less than 1% of inference time and approximately 4 KB per run across all models profiled, removing any practical argument against systematic documentation.
- (7) **Cloud deployment is compatible with near-deterministic inference.** The same LLaMA 3 8B architecture served via Together AI's cloud endpoint achieves near-local reproducibility (EMR = 1.000 for extraction, 0.880 for summarization on the shared 10-abstract subset—CIs overlap with the local deployment). While Together AI's specific infrastructure is not fully disclosed, this result demonstrates that cloud deployment *per se* does not preclude reproducibility, and is consistent with infrastructure complexity (tensor parallelism, speculative decoding, continuous batching) as the primary driver of non-determinism in closed-source API services.

These findings carry a broader implication: a substantial portion of published research that relies on closed-source API-based LLMs may contain non-reproducible results without the authors' knowledge. The quasi-isolation probe (Finding 7) suggests that infrastructure-level factors, rather than cloud deployment

per se, are plausible contributors to the observed non-determinism. Regardless of the specific mechanism, the protocol provides the infrastructure to detect, measure, and document such variability—making hidden non-determinism visible wherever it occurs.

Looking ahead, we plan to (i) extend the quasi-isolation probe experiment to additional cloud providers and model sizes (e.g., LLaMA 3 70B on Hugging Face Inference Endpoints) to test whether infrastructure complexity scales with model size; (ii) extend the task coverage to code generation, mathematical reasoning, and agentic workflows; and (iii) develop automated reproducibility scoring based on provenance graph analysis. Ultimately, we envision a future in which every generative AI output carries a provenance certificate, and reproducibility metrics are reported alongside accuracy as a standard component of empirical evaluation.

The reference implementation, all 4,104 run records, provenance documents, and analysis scripts are publicly available to support adoption and independent verification.

Acknowledgments

This work was supported by UTFPR – Universidade Tecnológica Federal do Paraná. The experiments were conducted using locally deployed open-weight models to ensure full reproducibility of the computational environment.

Data Availability Statement

The reference implementation, all 4,104 run records (JSON), PROV-JSON provenance documents, Run Cards, Prompt Cards, input data, analysis scripts, and generated figures are publicly available at:

<https://github.com/Roverlucas/genai-reproducibility-protocol>

The repository includes instructions for reproducing all experiments and regenerating all tables and figures from the raw data.

Author Contributions

Following the CRediT (Contributor Roles Taxonomy) framework: **Lucas Rover**: Conceptualization, Methodology, Software, Validation, Formal Analysis, Investigation, Data Curation, Writing – Original Draft, Writing – Review & Editing, Visualization, Project Administration. **Yara de Souza Tadano**: Supervision, Conceptualization, Methodology, Writing – Review & Editing, Project Administration.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors. The work was conducted as part of the authors' institutional activities at UTFPR.

Conflict of Interest

The authors declare no conflicts of interest. This research was conducted independently at UTFPR with no external funding from commercial AI providers. The use of commercial APIs (OpenAI, Anthropic, Google, DeepSeek, Perplexity, Together AI) was for research evaluation purposes only and does not constitute an endorsement of any provider.

Use of AI-Assisted Tools

The authors used AI-assisted tools (Claude, Anthropic) during the preparation of this manuscript for language editing, code development support, and data analysis scripting. All AI-generated content was critically reviewed, validated, and revised by the authors, who take full responsibility for the accuracy

and integrity of the final manuscript. The scientific design, experimental execution, interpretation of results, and intellectual contributions are entirely the authors' own work.

References

- J. Achiam et al.. 2023. *GPT-4 Technical Report*. arXiv preprint. (2023). arXiv: 2303.08774 (cs.CL).
- Anthropic. 2024. *The Claude Model Family*. (2024). <https://www.anthropic.com/clause>.
- B. Atil et al.. 2024. *Non-Determinism of “Deterministic” LLM Settings*. arXiv preprint. (2024). arXiv: 2408.04667 (cs.CL).
- M. Baker. 2016. “1,500 Scientists Lift the Lid on Reproducibility.” *Nature*, 533, 7604, 452–454. doi:10.1038/533452a.
- P. Ball. 2023. “Is AI Leading to a Reproducibility Crisis in Science?” *Nature*, 624, 7990, 22–25. doi:10.1038/d41586-023-03817-6.
- A. Belz, S. Agarwal, A. Shimorina, and E. Reiter. 2021. “A Systematic Review of Reproducibility Research in Natural Language Processing.” In: *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*. Association for Computational Linguistics, 381–393. doi:10.18653/v1/2021.eacl-main.29.
- A. Belz, C. Thomson, and E. Reiter. 2022. “Missing Information, Unresponsive Authors, Experimental Flaws: The Impossibility of Assessing the Reproducibility of Previous Human Evaluations in NLP.” In: *Proceedings of the 4th Workshop on Human Evaluation of NLP Systems*. Association for Computational Linguistics, 1–10. doi:10.18653/v1/2022.humeval-1.1.
- L. Biewald. 2020. *Experiment Tracking with Weights and Biases*. (2020). <https://wandb.com/>.
- A. Birhane, A. Kasirzadeh, D. Leslie, and S. Wachter. 2023. “Science in the Age of Large Language Models.” *Nature Reviews Physics*, 5, 5, 277–280. doi:10.1038/s42254-023-00581-4.
- R. Bommasani et al.. 2022. *On the Opportunities and Risks of Foundation Models*. arXiv preprint. (2022). arXiv: 2108.07258 (cs.LG).
- S. Bradner. 1997. *Key Words for Use in RFCs to Indicate Requirement Levels*. RFC 2119. Internet Engineering Task Force. doi:10.17487/RFC2119.
- T. Brown et al.. 2020. “Language Models are Few-Shot Learners.” In: *Advances in Neural Information Processing Systems*. Vol. 33, 1877–1901. arXiv: 2005.14165 (cs.CL).
- Y. Chen, J. Li, X. Liu, and Y. Li. 2023. *On the Reproducibility of ChatGPT in NLP Tasks*. arXiv preprint. (2023). arXiv: 2304.02554 (cs.CL).
- J. Cohen. 1988. *Statistical Power Analysis for the Behavioral Sciences*. (2nd ed.). Lawrence Erlbaum Associates. ISBN: 978-0-8058-0283-2.
- T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré. 2022. “FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness.” In: *Advances in Neural Information Processing Systems*. Vol. 35.
- J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. 2019. “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding.” In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*. Association for Computational Linguistics, 4171–4186. doi:10.18653/v1/N19-1423.
- J. Dodge, S. Gururangan, D. Card, R. Schwartz, and N. A. Smith. 2019. “Show Your Work: Improved Reporting of Experimental Results.” In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2185–2194. doi:10.18653/v1/D19-1224.
- European Parliament and Council of the European Union. 2024. *Regulation (EU) 2024/1689 Laying Down Harmonised Rules on Artificial Intelligence (AI Act)*. (2024). <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>.
- T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford. 2021. “Datasheets for Datasets.” *Communications of the ACM*, 64, 12, 86–92. doi:10.1145/3458723.
- Gemma Team et al.. 2024. *Gemma 2: Improving Open Language Models at a Practical Size*. arXiv preprint. (2024). arXiv: 2408.00118 (cs.CL).
- A. Golden et al.. 2024. “Is Flash Attention Stable?” *arXiv preprint arXiv:2405.02803*.
- A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, et al.. 2024. *The LLaMA 3 Herd of Models*. arXiv preprint. (2024). arXiv: 2407.21783 (cs.AI).
- O. E. Gundersen, K. Coakley, C. Kirkpatrick, and Y. Gil. 2022. “Sources of Irreproducibility in Machine Learning: A Review.” *arXiv preprint*. arXiv: 2204.07610 (cs.LG).
- O. E. Gundersen, Y. Gil, and D. W. Aha. 2018. “On Reproducible AI: Towards Reproducible Research, Open Science, and Digital Scholarship in AI Publications.” *AI Magazine*, 39, 3, 56–68. doi:10.1609/aimag.v39i3.2816.
- O. E. Gundersen, M. Helmert, and H. H. Hoos. 2024. “Improving Reproducibility in AI Research: Four Mechanisms Adopted by JAIR.” *Journal of Artificial Intelligence Research*, 81, 1019–1041. doi:10.1613/jair.1.16905.

- O. E. Gundersen and S. Kjensmo. 2018. “State of the Art: Reproducibility in Artificial Intelligence.” *Proceedings of the AAAI Conference on Artificial Intelligence*, 32, 1, 1644–1651. doi:[10.1609/aaai.v32i1.11503](https://doi.org/10.1609/aaai.v32i1.11503).
- M. Herschel, R. Diestelkämper, and H. Ben Lahmar. 2017. “A Survey on Provenance: What for? What form? What from?” *The VLDB Journal*, 26, 6, 881–906. doi:[10.1007/s00778-017-0486-1](https://doi.org/10.1007/s00778-017-0486-1).
- N. J. Higham. 2002. *Accuracy and Stability of Numerical Algorithms*. (2nd ed.). SIAM.
- M. Hutson. 2018. “Artificial Intelligence Faces Reproducibility Crisis.” *Science*, 359, 6377, 725–726. doi:[10.1126/science.359.6377.725](https://doi.org/10.1126/science.359.6377.725).
- A. Q. Jiang et al.. 2023. *Mistral 7B*. arXiv preprint. (2023). arXiv: [2310.06825 \(cs.CL\)](https://arxiv.org/abs/2310.06825).
- S. Kapoor and A. Narayanan. 2023. “Leakage and the Reproducibility Crisis in Machine-Learning-Based Science.” *Patterns*, 4, 9, 100804. doi:[10.1016/j.patter.2023.100804](https://doi.org/10.1016/j.patter.2023.100804).
- R. Kuprieiev, D. Petrov, and Iterative. 2024. *DVC: Data Version Control*. (2024). <https://dvc.org/>.
- W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica. 2023. “Efficient Memory Management for Large Language Model Serving with PagedAttention.” In: *Proceedings of the 29th ACM Symposium on Operating Systems Principles*.
- LangChain. 2023. *LangSmith: A Platform for Building Production-Grade LLM Applications*. (2023). <https://smith.langchain.com/>.
- V. I. Levenshtein. 1966. “Binary Codes Capable of Correcting Deletions, Insertions, and Reversals.” *Soviet Physics Doklady*, 10, 8, 707–710.
- Y. Leviathan, M. Kalman, and Y. Matias. 2023. “Fast Inference from Transformers via Speculative Decoding.” In: *Proceedings of the 40th International Conference on Machine Learning* (Proceedings of Machine Learning Research). Vol. 202. PMLR, 19274–19286.
- P. Liang et al.. 2023. “Holistic Evaluation of Language Models.” *Transactions on Machine Learning Research*. <https://openreview.net/forum?id=iO4LZibEqW>.
- C.-Y. Lin. 2004. “ROUGE: A Package for Automatic Evaluation of Summaries.” In: *Proceedings of the ACL-04 Workshop on Text Summarization Branches Out*. Association for Computational Linguistics, 74–81. <https://aclanthology.org/W04-1013>.
- M. Mitchell, S. Wu, A. Zaldivar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, and T. Gebru. 2019. “Model Cards for Model Reporting.” In: *Proceedings of the Conference on Fairness, Accountability, and Transparency*. ACM, 220–229. doi:[10.1145/3287560.3287596](https://doi.org/10.1145/3287560.3287596).
- L. Moreau and P. Missier. 2013. *PROV-DM: The PROV Data Model*. W3C Recommendation. World Wide Web Consortium. <https://www.w3.org/TR/prov-dm/>.
- National Institute of Standards and Technology. 2023. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. Tech. rep. U.S. Department of Commerce. doi:[10.6028/NIST.AI.100-1](https://doi.org/10.6028/NIST.AI.100-1).
- S. Noy and W. Zhang. 2023. “Experimental Evidence on the Productivity Effects of Generative Artificial Intelligence.” *Science*, 381, 6654, 187–192. doi:[10.1126/science.adh2586](https://doi.org/10.1126/science.adh2586).
- Ollama. 2024. *Ollama: Run Large Language Models Locally*. (2024). <https://ollama.com/>.
- OpenAI. 2024. *API Reference: Create Chat Completion — Seed Parameter*. “If specified, our system will make a best effort to sample deterministically [...] Determinism is not guaranteed.” (2024). <https://platform.openai.com/docs/api-reference/chat/create>.
- OpenAI. 2023. *OpenAI Evals: A Framework for Evaluating LLMs*. (2023). <https://github.com/openai/evals>.
- S. Ouyang, J. M. Zhang, M. Harman, and M. Wang. 2024. “An Empirical Study of the Non-determinism of ChatGPT in Code Generation.” *ACM Transactions on Software Engineering and Methodology*, 34, 2, 1–28. doi:[10.1145/3697010](https://doi.org/10.1145/3697010).
- G. Padovani, V. Ananthraj, and S. Fiore. 2025. “yProv4ML: Effortless Provenance Tracking for Machine Learning Systems.” *SoftwareX*, 29, 102028. doi:[10.1016/j.softx.2025.102028](https://doi.org/10.1016/j.softx.2025.102028).
- J. Pineau, P. Vincent-Lamarre, K. Sinha, V. Larivière, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and H. Larochelle. 2021. “Improving Reproducibility in Machine Learning Research: A Report from the NeurIPS 2019 Reproducibility Program.” *Journal of Machine Learning Research*, 22, 164, 1–20. <https://jmlr.org/papers/v22/20-303.html>.
- PyTorch Contributors. 2024. *Reproducibility — PyTorch Documentation*. Documents sources of non-determinism in GPU operations and `torch.use_deterministic_algorithms()` flag. (2024). <https://pytorch.org/docs/stable/notes/randomness.html>.
- C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. 2020. “Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer.” *Journal of Machine Learning Research*, 21, 140, 1–67. <https://jmlr.org/papers/v21/20-074.html>.

- M. Reid, N. Savinov, D. Teber, I. Bica, Z. Shen, D. Lepikhin, et al.. 2024. *Gemini 1.5: Unlocking Multimodal Understanding Across Millions of Tokens of Context*. arXiv preprint; Gemini 2.5 Pro builds upon this architecture. (2024). arXiv: [2403.05530 \(cs.CL\)](#).
- A. Rogers, T. Baldwin, and K. Leins. 2021. ““Just Think About It”: Incentivizing Reproducibility in NLP.” *Computational Linguistics*, 47, 4, 757–770. doi:[10.1162/coli_a_00420](#).
- J. Romano, J. D. Kromrey, J. Coraggio, and J. Skowronek. 2006. “Appropriate Statistics for Ordinal Level Data: Should We Really Be Using t-Test and Cohen’s d for Evaluating Group Differences on the NSSE and Similar Surveys?” In: *Annual Meeting of the Florida Association of Institutional Research*.
- S. Samuel and B. König-Ries. 2022. “A Provenance-based Semantic Approach to Support Understandability, Reproducibility, and Reuse of Scientific Experiments.” *Journal of Biomedical Semantics*, 13, 1, 1–30. doi:[10.1186/s13326-022-00263-z](#).
- M. Shoeybi, M. Patwary, R. Puri, P. LeGresley, J. Casper, and B. Catanzaro. 2019. “Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism.” *arXiv preprint arXiv:1909.08053*.
- K. Singhal et al.. 2023. “Large Language Models Encode Clinical Knowledge.” *Nature*, 620, 7972, 172–180. doi:[10.1038/s41586-023-06291-2](#).
- V. Stodden, M. McNutt, D. H. Bailey, E. Deelman, Y. Gil, B. Hanson, M. A. Heroux, J. P. Ioannidis, and M. Taufer. 2016. “Enhancing Reproducibility for Computational Methods.” *Science*, 354, 6317, 1240–1241. doi:[10.1126/science.aah6168](#).
- A. J. Thirunavukarasu, D. S. J. Ting, K. Elangovan, L. Gutierrez, T. F. Tan, and D. S. W. Ting. 2023. “Large Language Models in Medicine.” *Nature Medicine*, 29, 8, 1930–1940. doi:[10.1038/s41591-023-02448-8](#).
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. 2017. “Attention is All You Need.” In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc. arXiv: [1706.03762 \(cs.CL\)](#).
- J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V. Le, and D. Zhou. 2022. “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.” In: *Advances in Neural Information Processing Systems*. Vol. 35, 24824–24837. arXiv: [2201.11903 \(cs.CL\)](#).
- M. D. Wilkinson et al.. 2016. “The FAIR Guiding Principles for Scientific Data Management and Stewardship.” *Scientific Data*, 3, 160018. doi:[10.1038/sdata.2016.18](#).
- G.-I. Yu, J. S. Jeong, G.-W. Kim, S. Kim, and B.-G. Chun. 2022. “Orca: A Distributed Serving System for Transformer-Based Generative Models.” In: *16th USENIX Symposium on Operating Systems Design and Implementation*. USENIX Association, 521–538.
- J. Yuan et al.. 2025. “Understanding and Mitigating Numerical Sources of Nondeterminism in LLM Inference.” In: *Advances in Neural Information Processing Systems*. Vol. 38. arXiv preprint. Curran Associates, Inc. arXiv: [2506.09501 \(cs.LG\)](#).
- M. Zaharia et al.. 2018. “Accelerating the Machine Learning Lifecycle with MLflow.” *IEEE Data Engineering Bulletin*, 41, 4, 39–45. <http://sites.computer.org/debull/A18dec/p39.pdf>.
- T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi. 2020. “BERTScore: Evaluating Text Generation with BERT.” In: *Proceedings of the 8th International Conference on Learning Representations*. <https://openreview.net/forum?id=SkeHuCVFDr>.
- Y. Zhu, P. Zhang, E. Haq, P. Hui, and G. Buchanan. 2023. *Can ChatGPT Reproduce Human-Generated Labels? A Study of Social Computing Tasks*. arXiv preprint. (2023). arXiv: [2304.10145 \(cs.CL\)](#).

A Reproducibility Checklist

The following checklist is designed for self-assessment of reproducibility in generative AI studies. Each item maps to a specific field or artifact in our protocol.

Prompt Documentation

- (1) Is the exact prompt text recorded and versioned? [Prompt Card: `prompt_text, prompt_hash`]
- (2) Are design assumptions and limitations documented? [Prompt Card: `assumptions, limitations`]
- (3) Is the expected output format specified? [Prompt Card: `expected_output_format`]
- (4) Is the interaction regime documented (single/multi-turn)? [Prompt Card: `interaction_regime`]

Model and Environment

- (5) Is the model name and version recorded? [Run Card: `model_name, model_version`]
- (6) Are model weights hashed for identity verification? [Run Card: `weights_hash`]
- (7) Is the execution environment fingerprinted? [Run Card: `environment, environment_hash`]
- (8) Is the source code version recorded? [Run Card: `code_commit`]

Execution and Output

- (9) Are all inference parameters logged? [Run Card: `inference_params`]
- (10) Is the random seed recorded? [Run Card: `inference_params.seed`]
- (11) Is the output cryptographically hashed? [Run Card: `output_hash`]
- (12) Are execution timestamps recorded? [Run Card: `timestamp_start, timestamp_end`]
- (13) Is logging overhead measured separately? [Run Card: `logging_overhead_ms`]

Provenance

- (14) Is a provenance graph generated per group? [PROV-JSON document]
- (15) Are provenance documents in an interoperable format? [W3C PROV standard]

B Run Card Schema

The complete Run Card schema, with data types and descriptions:

Listing 1. Run Card JSON schema (simplified).

```

1 {
2   "run_id": "string (unique identifier)",
3   "task_id": "string (task identifier)",
4   "task_category": "string (e.g., summarization)",
5   "prompt_hash": "string (SHA-256 of prompt)",
6   "prompt_text": "string (full prompt text)",
7   "input_text": "string (input to the model)",
8   "input_hash": "string (SHA-256 of input)",
9   "model_name": "string (e.g., llama3:8b)",
10  "model_version": "string (e.g., 8.0B)",
11  "weights_hash": "string (SHA-256 of weights)",
12  "model_source": "string (e.g., ollama-local)",
13  "inference_params": {
14    "temperature": "float",
15    "top_p": "float",

```

```

16     "top_k": "integer",
17     "max_tokens": "integer",
18     "seed": "integer|null",
19     "decoding_strategy": "string"
20   },
21   "params_hash": "string (SHA-256 of params)",
22   "environment": {
23     "os": "string",
24     "os_version": "string",
25     "architecture": "string",
26     "python_version": "string",
27     "hostname": "string",
28     "timestamp": "ISO 8601 datetime"
29   },
30   "environment_hash": "string (SHA-256)",
31   "code_commit": "string (git commit hash)",
32   "researcher_id": "string",
33   "affiliation": "string",
34   "timestamp_start": "ISO 8601 datetime",
35   "timestamp_end": "ISO 8601 datetime",
36   "output_text": "string (model output)",
37   "output_hash": "string (SHA-256 of output)",
38   "output_metrics": "object (task-specific)",
39   "execution_duration_ms": "float",
40   "logging_overhead_ms": "float",
41   "storage_kb": "float",
42   "system_logs": "string (raw system info)",
43   "errors": "array of strings",
44
45 // --- API-specific optional fields ---
46   "api_request_id": "string|null (provider request ID)",
47   "api_response_headers": "object|null (selected headers)",
48   "api_model_version_returned": "string|null",
49   "api_region": "string|null (if available)",
50   "seed_status": "string (sent|logged-only|not-supported)",
51
52 // --- Multi-turn extension fields ---
53   "conversation_history_hash": "string|null (SHA-256)",
54   "turn_index": "integer|null",
55   "parent_run_id": "string|null",
56
57 // --- RAG extension fields ---
58   "retrieval_context": "string|null",
59   "retrieval_context_hash": "string|null (SHA-256)"
60 }

```

C Chat-Format Control Experiment

To assess whether the prompt-format difference between LLaMA 3 (completion-style via `/api/generate`) and GPT-4 (chat-style via Chat Completions) contributes to the observed reproducibility gap, we

conducted a supplementary control experiment running LLaMA 3 8B through Ollama’s `/api/chat` endpoint, which applies the model’s chat template (including special tokens for system/user/assistant roles) in the same message structure used by GPT-4.

Design: 10 abstracts \times 2 tasks \times 2 conditions (C1, C2) \times 5 repetitions = 200 runs, all under greedy decoding ($t=0$).

Results: Table 10 compares the chat-format control with the original completion-format results for the same 10 abstracts. The two prompt formats produce *identical* variability metrics across all conditions: summarization EMR = 0.929, NED = 0.0066, and ROUGE-L = 0.9922 in both modes; extraction achieves perfect reproducibility (EMR = 1.000) regardless of interface. The 0.929 summarization EMR reflects the warm-up effect on 2 of 10 abstracts—the same pattern observed in the full 30-abstract experiment. These results confirm that prompt format is not a source of variability, and the reproducibility gap between LLaMA 3 and GPT-4 is consistent with deployment-side factors (server infrastructure, floating-point non-determinism across GPU types, request batching) rather than prompt-format differences.

Table 10. Prompt-format control: LLaMA 3 8B via completion (`/api/generate`) vs. chat (`/api/chat`) for 10 abstracts under greedy decoding ($t=0$). EMR computed over conditions C1 and C2 combined.

Task	Metric	Completion	Chat
Summarization	EMR↑	0.929	0.929
	NED↓	0.0066	0.0066
	ROUGE-L↑	0.9922	0.9922
Extraction	EMR↑	1.000	1.000
	NED↓	0.0000	0.0000
	ROUGE-L↑	1.0000	1.0000

Note: Completion and chat formats yield identical metrics for all 10 abstracts under greedy decoding, indicating that prompt format is not a source of variability.

D API Payload Documentation

To address potential “apples-to-oranges” concerns, we document the exact payload structures sent to each inference endpoint. All payloads were constructed deterministically and logged as part of the Run Card.

Local models (Ollama). Single-turn tasks use `POST /api/generate`:

Ollama generate payload (Tasks 1–2):

```
{"model": "llama3:8b",
 "prompt": "<full prompt text>",
 "options": {"temperature": 0.0, "seed": 42, "num_predict": 1024},
 "stream": false}
```

The `model` field is set to `llama3:8b`, `mistral:7b`, or `gemma2:9b` as appropriate. Multi-turn tasks (Task 3) use `POST /api/chat` with accumulated `messages` array. No system prompt, stop sequences, or post-processing are applied.

GPT-4 (OpenAI). Accessed via the `openai` Python SDK v1.59.9:

```
{"model": "gpt-4",
 "messages": [{"role": "user", "content": "<prompt>"}],
 "temperature": 0.0, "seed": 42, "max_tokens": 1024}
```

No system message, stop sequences, `top_p`, `frequency_penalty`, or `presence_penalty` were set (all defaults). The resolved model version (`gpt-4-0613`) was extracted from the response object and logged.

Claude Sonnet 4.5 (Anthropic). Accessed via `urllib` (no SDK dependency):

```
{"model": "claude-sonnet-4-5-20250929",
 "messages": [{"role": "user", "content": "<prompt>"}],
 "temperature": 0.0, "max_tokens": 1024}
```

No `seed` parameter (not supported by the Anthropic API), no system message, no stop sequences. The seed value in the Run Card is marked `seed_status: "logged-only-not-sent-to-api"`.

Gemini 2.5 Pro (Google). Accessed via `urllib` (no SDK dependency) through the Google AI Studio REST API:

```
{"contents": [{"role": "user", "parts": [{"text": "<prompt>"}]}],
 "generationConfig": {"maxOutputTokens": 8192,
 "temperature": 0.0, "seed": 42},
 "systemInstruction": {"parts": [{"text": "<system>"}]}}
```

The `seed` parameter is supported by the Gemini API and is sent with every request. Gemini 2.5 Pro is a “thinking” model: internal reasoning tokens count against the `maxOutputTokens` budget, hence the higher limit (8,192 vs. 1,024 for other models). Multi-turn tasks use the same endpoint with accumulated `contents` array alternating `role: "user"` and `role: "model"`.

DeepSeek Chat (DeepSeek). Accessed via the OpenAI-compatible API:

```
{"model": "deepseek-chat",
 "messages": [{"role": "user", "content": "<prompt>"}],
 "temperature": 0.0, "max_tokens": 1024}
```

No seed parameter, system message, or stop sequences. DeepSeek Chat achieved the highest API reproducibility (EMR = 0.800 for extraction), suggesting effective internal determinism under greedy decoding.

Perplexity Sonar (Perplexity). Accessed via the Perplexity API:

```
{"model": "sonar",
 "messages": [{"role": "user", "content": "<prompt>"}],
 "temperature": 0.0, "max_tokens": 1024}
```

Perplexity Sonar is an online model with real-time search augmentation. No seed parameter or system message. The search-augmented nature introduces an additional source of variability: retrieved web content may differ across requests, contributing to the lowest observed reproducibility (EMR = 0.010–0.100).

Key symmetry points. Across all nine model deployments: (1) identical prompt text (verified by `prompt_hash`); (2) identical temperature ($t=0.0$); (3) identical max token limit (1,024; 8,192 for Gemini 2.5 Pro to accommodate thinking tokens); (4) no system messages for single-turn tasks; (5) no stop sequences; (6) no post-processing or text normalization of outputs.

E JAIR Reproducibility Compliance

This appendix documents compliance with the JAIR reproducibility mechanisms described by [Gundersen, Helmert, et al. \(2024\)](#). Each item maps to the four mechanisms: checklists, structured abstracts, badges, and reproducibility reports.

Mechanism 1: Reproducibility Checklist

- ✓ **Code availability:** Reference implementation publicly available at <https://github.com/Roverlucas/genai-reproducibility-protocol> (MIT License).
- ✓ **Data availability:** All 4,104 run records, PROV documents, and 30 input abstracts included in the repository.

- ✓ **Experimental methodology:** Full description of models, tasks, conditions, and metrics in Section 4.
- ✓ **Hyperparameters:** All inference parameters documented per model (Appendix D); temperature = 0; max_tokens = 1024 for all models except Gemini 2.5 Pro (maxOutputTokens = 8,192 to accommodate thinking tokens).
- ✓ **Statistical tests:** Holm-Bonferroni correction across 68 tests, Fisher's exact tests, bootstrap 95% CIs (10,000 resamples).
- ✓ **Randomness control:** Random seeds documented; bootstrap seed = 42; five experimental conditions (C1, C2, C3 with three temperature sub-conditions) isolate seed and parameter effects.
- ✓ **Computing infrastructure:** Hardware and software environment documented (Apple M4, Ollama v0.15.5, Python 3.14.3).
- ✓ **Overhead:** Protocol overhead measured at <1% of inference time (\approx 25 ms per run) and \approx 4 KB storage per run.

Mechanism 2: Structured Abstract

- ✓ Abstract follows JAIR's structured format: Background, Objectives, Methods, Results, Conclusions.

Mechanism 3: Reproducibility Badge Eligibility

- ✓ **Code:** Complete reference implementation in Python.
- ✓ **Data:** All experimental data (run records, provenance graphs, input texts) publicly available.
- ✓ **Experiment:** Full pipeline (data collection, analysis, figure generation) can be re-executed from the repository.

Mechanism 4: Reproducibility Report

- ✓ Sensitivity analysis via balanced 10-abstract subsample confirms findings (Section 5.1).
- ✓ Chat-format control experiment (200 additional runs) rules out prompt-format confound (Appendix C).
- ✓ Automated test suite (51 tests) validates core analysis functions.
- ✓ Continuous integration pipeline ensures reproducibility of code and analysis.

Received February 2026