

Hidden Non-Determinism in Large Language Model APIs: A Lightweight Provenance Protocol for Reproducible Generative AI Research

LUCAS ROVER*, UTFPR – Universidade Tecnológica Federal do Paraná, Brazil

YARA DE SOUZA TADANO, UTFPR – Universidade Tecnológica Federal do Paraná, Brazil

Background: Generative AI models produce non-deterministic outputs that vary across runs, even under nominally identical configurations. This variability threatens the reproducibility of studies that rely on large language model (LLM) outputs, yet most existing experiment-tracking tools were not designed for the specific challenges of text-generation workflows.

Objectives: We propose a lightweight, open-standard protocol for logging, versioning, and provenance tracking of generative AI experiments. The protocol introduces two novel documentation artifacts—Prompt Cards and Run Cards—and adopts the W3C PROV data model to create auditable, machine-readable provenance graphs linking every output to its full generation context.

Methods: We formalize the protocol and evaluate it empirically through 3,604 controlled experiments. These experiments employ five models—three locally deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B) and two API-served (GPT-4, Claude Sonnet 4.5)—on four NLP tasks. All five models are evaluated on single-turn extraction and summarization under five conditions varying seed, temperature, and decoding strategy (10–30 abstracts per model; GPT-4 primarily under variable-seed greedy decoding due to quota constraints). Multi-turn refinement and RAG extraction are evaluated for the three local models and Claude Sonnet 4.5 under greedy decoding (10 abstracts each). We measure output variability using Exact Match Rate, Normalized Edit Distance, ROUGE-L, and BERTScore, and quantify the protocol’s own overhead in terms of time and storage.

Results: Under greedy decoding ($t=0$), local models achieve near-perfect reproducibility: Gemma 2 9B reaches $\text{EMR} = 1.000$ across all tasks, LLaMA 3 attains $\text{EMR} = 0.987$ for extraction, and Mistral 7B achieves $\text{EMR} = 0.960$. By contrast, API-served models exhibit substantial hidden non-determinism: GPT-4 achieves only $\text{EMR} = 0.443$ for extraction, while Claude Sonnet 4.5 achieves $\text{EMR} = 0.190$ for extraction and $\text{EMR} = 0.020$ for summarization—the lowest observed in our study. This local-vs-API reproducibility gap (average single-turn EMR : 0.956 vs. 0.221, a more than 4-fold difference) is confirmed across two independent API providers. The gap extends to complex interaction regimes: under multi-turn refinement and RAG extraction, local models maintain high reproducibility ($\text{EMR} \geq 0.880$), while Claude Sonnet 4.5—the only API model tested on these tasks—achieves $\text{EMR} = 0.040$ for multi-turn and $\text{EMR} = 0.000$ for RAG. The protocol adds less than 1% overhead across all five models.

Conclusions: Our results provide evidence that (1) API-served models exhibit substantial non-determinism under greedy decoding that is not attributable to user-controllable parameters, a pattern observed independently for both GPT-4 and Claude; (2) locally deployed models achieve near-perfect to perfect bitwise reproducibility under greedy decoding; (3) the local-vs-API gap extends to multi-turn refinement and RAG extraction,

*Corresponding Author.

Authors' Contact Information: Lucas Rover, ORCID: [0000-0001-6641-9224](#), lucasrover@utfpr.edu.br, UTFPR – Universidade Tecnológica Federal do Paraná, Programa de Pós-Graduação em Engenharia Mecânica, Ponta Grossa, Paraná, Brazil; Yara de Souza Tadano, ORCID: [0000-0002-3975-3419](#), yaratadano@utfpr.edu.br, UTFPR – Universidade Tecnológica Federal do Paraná, Programa de Pós-Graduação em Engenharia Mecânica, Ponta Grossa, Paraná, Brazil.



This work is licensed under a Creative Commons Attribution International 4.0 License.

1 where Claude Sonnet 4.5 (the only API model tested on these tasks) achieves near-zero EMR while local models maintain EMR ≥ 0.880 ; (4) temperature is the dominant user-controllable factor affecting variability; and (5) comprehensive provenance logging adds negligible overhead (<1%). The protocol, reference implementation, and all experimental data are publicly available.

5 CCS Concepts: • Software and its engineering → Software testing and debugging; Documentation; • Computing methodologies → Machine learning.

7 Additional Key Words and Phrases: reproducibility, large language models, non-determinism, provenance, generative AI, experiment tracking, W3C PROV, scientific methodology

9 **JAIR Associate Editor:** Insert JAIR AE Name

10 **JAIR Reference Format:**

12 Lucas Rover and Yara de Souza Tadano. 2026. Hidden Non-Determinism in Large Language Model APIs: A
13 Lightweight Provenance Protocol for Reproducible Generative AI Research. *Journal of Artificial Intelligence
14 Research* (2026), 34 pages. DOI: [10.1613/jair.1.xxxxx](https://doi.org/10.1613/jair.1.xxxxx)

15 1 Introduction

17 When a researcher queries a cloud-hosted LLM with the same prompt and temperature zero, one would
18 reasonably expect identical outputs. Our experiments show otherwise: across five controlled seeds under
19 greedy decoding, GPT-4 produces the same extraction result only 44% of the time, and Claude Sonnet
20 4.5 achieves only 19%. Meanwhile, locally deployed models such as Gemma 2 9B produce *perfectly identical*
21 outputs every time. This hidden, provider-dependent non-determinism exemplifies a fundamental
22 challenge introduced by the rapid adoption of large language models (LLMs) in scientific research: how
23 to ensure that studies relying on generative AI outputs are reproducible, auditable, and scientifically
24 rigorous. Unlike traditional computational experiments, in which deterministic algorithms produce iden-
25 tical results given identical inputs, LLMs exhibit inherent variability in their outputs due to stochastic
26 sampling, floating-point non-determinism, and opaque model-versioning practices (Y. Chen et al. 2023;
27 Zhu et al. 2023).

28 Importantly, “non-reproducible” does not necessarily mean “unreliable”: our experiments also show
29 that semantic similarity (measured by BERTScore F1) remains above 0.94 across all conditions, even
30 when exact textual match drops to zero. In other words, API outputs typically convey the same *meaning*
31 despite differing in *phrasing*—but this distinction is invisible without systematic measurement, and many
32 downstream analyses (meta-analyses, comparative studies, regulatory audits) require exact reproducibil-
33 ity.

34 A related subtlety concerns the `seed` parameter offered by some APIs. For API-served models, the
35 seed parameter is advisory, not a guarantee of determinism: OpenAI explicitly documents this caveat
36 for GPT-4, and Anthropic’s Claude API does not support a seed parameter at all. Our experimental
37 design accounts for this by treating seed variation as a control condition and measuring actual output
38 reproducibility directly, rather than relying on API-side determinism guarantees.

39 This reproducibility challenge is not merely theoretical. Baker (2016) reported that over 70% of re-
40 searchers have failed to reproduce another scientist’s experiment, a crisis that extends to AI research
41 (Gundersen and Kjensmo 2018; Hutson 2018; Kapoor and A. Narayanan 2023; Stodden et al. 2016).
42 For generative AI specifically, the problem is compounded by several factors unique to text-generation
43 workflows: (1) the same prompt can yield semantically similar yet textually distinct outputs across runs;
44 (2) API-based models may undergo silent updates that alter behavior; (3) temperature and sampling
45 parameters create a high-dimensional space of possible outputs; and (4) no established standard exists
46 for documenting the full context needed to understand, audit, or reproduce a generative output.

48 Existing experiment-tracking tools such as MLflow (Zaharia et al. 2018), Weights & Biases (Biewald
 49 2020), and DVC (Kuprieiev et al. 2024) were designed primarily for training pipelines and numerical
 50 metrics. Although valuable for their intended purposes, these tools lack features critical for generative
 51 AI studies: structured prompt versioning, cryptographic output hashing for tamper detection, prove-
 52 nance graphs linking outputs to their full generation context, and environment fingerprinting specific to
 53 inference-time conditions.

54 In this paper, we make three contributions:

- 55 (1) **A lightweight protocol** for logging, versioning, and provenance tracking of generative AI exper-
 56 iments. The protocol introduces *Prompt Cards* and *Run Cards* as structured documentation arti-
 57 facts, and adopts the W3C PROV data model (Moreau and Missier 2013) for machine-readable
 58 provenance graphs.
- 59 (2) **An empirical evaluation** of both the protocol’s effectiveness and the reproducibility charac-
 60 teristics of LLM outputs. Through 3,604 controlled experiments with five models—three locally
 61 deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B) and two API-served (GPT-4, Claude Son-
 62 net 4.5)—across four tasks (extraction, summarization, multi-turn refinement, RAG extraction),
 63 30 abstracts, and five conditions, we quantify output variability using four complementary met-
 64 rrics and measure the protocol’s overhead. Our results document a striking, provider-independent
 65 reproducibility gap between local and API-based inference that is invisible without systematic
 66 logging.
- 67 (3) **A reference implementation** in Python that demonstrates the protocol’s practical applicabil-
 68 ity, together with all experimental data, to facilitate adoption and independent verification.

70 The remainder of this paper is organized as follows. Section 2 reviews related work on reproducibility in
 71 AI and experiment tracking. Section 3 formalizes the protocol design. Section 4 describes the experimental
 72 methodology. Section 5 presents the empirical results. Section 6 discusses findings, limitations, and
 73 practical implications. Section 7 concludes with directions for future work.

74 2 Related Work

75 2.1 Reproducibility in AI Research

76 The reproducibility crisis in AI has been documented extensively. Gundersen and Kjensmo (2018) sur-
 77veyed 400 AI papers and found that only 6% provided sufficient information for full reproducibility.
 78 Pineau et al. (2021) reported on the NeurIPS 2019 Reproducibility Program, which introduced re-
 79 producibility checklists and found significant gaps between reported and actual reproducibility. More
 80 recently, Gundersen, Helmert, et al. (2024) described four institutional mechanisms adopted by JAIR—
 81 reproducibility checklists, structured abstracts, badges, and reproducibility reports—establishing a com-
 82 munity standard for what should be documented in AI research. Gundersen, Gil, et al. (2018) identified
 83 three levels of reproducibility in AI—method, data, and experiment—and argued that all three are nec-
 84 essary for scientific progress. Belz et al. (2021) conducted a systematic review of 601 NLP papers and
 85 confirmed pervasive under-reporting of experimental details, while Dodge et al. (2019) proposed improved
 86 reporting standards for ML experiments, including confidence intervals and significance tests across mul-
 87 tiple runs. More broadly, Kapoor and A. Narayanan (2023) identified data leakage as a widespread driver
 88 of irreproducible results across 17 scientific fields that use ML-based methods.

89 For generative AI specifically, Y. Chen et al. (2023) demonstrated that ChatGPT’s outputs on NLP
 90 benchmarks exhibit non-trivial variability across identical queries, even with temperature set to zero.
 91 Zhu et al. (2023) showed that reproducibility degrades further when tasks involve subjective judgment,

such as social computing annotations. Most recently, Atil et al. (2024) systematically measured the non-determinism of five LLMs under supposedly deterministic settings across eight tasks, finding accuracy variations up to 15% across runs and introducing the Total Agreement Rate (TAR) metric. Ouyang et al. (2024) confirmed that temperature zero does not guarantee determinism in ChatGPT code generation. Most recently, Yuan et al. (2025) traced such non-determinism to numerical precision issues in GPU kernels and proposed LayerCast as a mitigation strategy—a hardware-level fix that reduces but does not eliminate non-determinism, and that is not available to researchers using closed API services. Our Exact Match Rate (EMR) metric is closely related to Atil et al.’s Total Agreement Rate (TAR), which measures the fraction of runs producing the modal output; EMR instead measures the fraction of *all output pairs* that match exactly, providing a more sensitive measure when agreement is low and no clear modal output exists. Our work complements these studies in four specific ways. First, whereas prior studies (including Atil et al.’s five-model, eight-task study) measure variability post hoc, we provide a structured provenance protocol that enables *prospective* documentation and audit—answering not only “how much variability?” but also “why did these outputs differ?” through cryptographic hashing and W3C PROV graphs. Second, we directly compare local and API-based inference on identical tasks with identical prompts across *five* models and *two* independent API providers (OpenAI and Anthropic), isolating the deployment paradigm as a variable and suggesting that API non-determinism is a consistent pattern across providers rather than a provider-specific artifact. Third, we extend beyond single-turn evaluation to include multi-turn refinement and retrieval-augmented generation, demonstrating that reproducibility characteristics generalize across interaction regimes. Fourth, we quantify the overhead of systematic logging, demonstrating that the “cost of knowing” is negligible.

2.2 Experiment Tracking Tools

Several tools exist for tracking machine learning experiments, although none was designed specifically for generative AI text-output workflows:

MLflow (Zaharia et al. 2018) provides experiment tracking, model packaging, and deployment. It logs parameters, metrics, and artifacts, but focuses on training pipelines and numerical outcomes rather than text-generation provenance.

Weights & Biases (Biewald 2020) offers experiment tracking with visualization dashboards. It supports prompt logging but lacks structured prompt versioning, cryptographic output hashing, and provenance graph generation.

DVC (Kuprieiev et al. 2024) provides data versioning through git-like operations. While effective for dataset management, it does not address run-level provenance or prompt documentation.

OpenAI Eval (OpenAI 2023) is a framework for evaluating LLM outputs against benchmarks. It provides structured evaluation but is tightly coupled to OpenAI’s ecosystem and does not generate interoperable provenance records.

LangSmith (LangChain 2023) offers tracing and evaluation for LLM applications. It captures detailed execution traces but uses a proprietary format and requires cloud connectivity.

More broadly, Bommasani et al. (2022) identified reproducibility as a key risk for foundation models, and Liang et al. (2023) proposed the HELM benchmark for holistic evaluation of language models, including robustness and fairness dimensions that complement our reproducibility focus. In the provenance space, Padovani et al. (2025) recently introduced yProv4ML, a framework that captures ML provenance in PROV-JSON format with minimal code modifications; our protocol shares the commitment to W3C PROV and SHA-256 hashing but differs in three key respects: (i) we target inference-time stochastic text generation rather than training pipelines; (ii) our Run Cards capture prompt-level metadata (prompt

142 Table 1. Comparison of our protocol with existing reproducibility tools and frameworks for GenAI experiments. Check-
 143 marks (✓) indicate full support; tildes (~) indicate partial support; dashes (–) indicate no support.

144

145 Feature	Ours	MLflow	W&B	DVC	OpenAI Eval	LangSmith
146 Prompt versioning (Prompt Card)	✓	–	~	–	~	~
147 Run-level provenance (W3C PROV)	✓	–	–	–	–	–
148 Cryptographic output hashing	✓	–	–	✓	–	–
149 Seed & param logging	✓	✓	✓	–	✓	✓
150 Environment fingerprinting	✓	~	~	~	–	–
151 Model weights hashing	✓	–	~	✓	–	–
152 Overhead <1% of inference	✓	~	~	N/A	N/A	~
153 Designed for GenAI text output	✓	–	–	–	✓	✓
154 Open standard (PROV-JSON)	✓	–	–	–	–	–
155 Local-first (no cloud dependency)	✓	✓	–	✓	–	–

156

157

158 hash, seed status, interaction regime) not present in training-oriented schemas; and (iii) we provide
 159 empirical evidence quantifying why such logging is necessary for API-served models.

160 Table 1 provides a systematic feature-by-feature comparison of our protocol with these tools. The key
 161 distinction is not merely one of tooling but of *scientific capability*: existing tools log what happened
 162 during training (parameters, metrics, artifacts), whereas our protocol enables answering questions that
 163 these tools cannot—specifically, whether two generative outputs are provably derived from identical
 164 configurations, which exact factor caused a divergence between non-identical outputs, and whether an
 165 output has been tampered with post-generation. These capabilities require the combination of crypto-
 166 graphic hashing, structured prompt documentation, and W3C PROV provenance graphs that no existing
 167 tool provides. In short, our contribution is not an alternative experiment tracker but a *reproducibility*
 168 *assessment framework* designed for the unique challenges of stochastic text generation.
 169

170 2.3 Provenance in Scientific Computing

171 Data provenance—the lineage of data through transformations—has a rich history in database systems
 172 and scientific workflows (Herschel et al. 2017). The W3C PROV family of specifications (Moreau and
 173 Missier 2013) provides a standardized data model for representing provenance as directed acyclic graphs
 174 of *entities*, *activities*, and *agents*. Samuel and König-Ries (2022) applied provenance tracking to compu-
 175 tational biology workflows, demonstrating its value for reproducibility. However, to our knowledge, no
 176 prior work has applied W3C PROV specifically to generative AI experiment workflows, in which the
 177 challenge involves not only tracking data lineage but also capturing the stochastic generation context
 178 that determines output variability.
 179

180 Taken together, these gaps point to a clear need: a lightweight, standards-based protocol that bridges
 181 generative AI inference with the provenance infrastructure already established in scientific computing.
 182 The next section presents our design for such a protocol.

183

184

3 Protocol Design

185 Our protocol addresses the question: *What is the minimum set of metadata that must be captured for each*
 186 *generative AI run to enable auditing, reproducibility assessment, and provenance tracking?* We address
 187 this question through four complementary components.

188

189 **3.1 Scope and Design Principles**

190 The protocol is designed around three principles:

- 191 (1) **Completeness**: Every factor that can influence a generative output must be captured—prompt
192 text, model identity and version, inference parameters, environment state, and timestamps.
193 (2) **Negligible overhead**: The logging process must not materially affect the experiment. We target
194 <1% overhead relative to inference time.
195 (3) **Interoperability**: All artifacts are stored in open, machine-readable formats (JSON, PROV-
196 JSON), aligned with the FAIR (Findable, Accessible, Interoperable, Reusable) principles ([Wilkin-
197 son et al. 2016](#)), to enable tool integration and long-term preservation.

199 **3.2 Prompt Cards**

200 A *Prompt Card* is a versioned documentation artifact that captures the design rationale and metadata
201 for a prompt template used in experiments. Each Prompt Card contains:

- 203 • `prompt_id`: Unique identifier
204 • `prompt_hash`: SHA-256 hash of the prompt text, enabling tamper detection
205 • `version`: Semantic version number
206 • `task_category`: Classification of the task (e.g., summarization, extraction)
207 • `objective`: Natural-language description of what the prompt is designed to achieve
208 • `assumptions`: Explicit assumptions about inputs and expected behavior
209 • `limitations`: Known limitations or failure modes
210 • `target_models`: Models for which the prompt was designed and tested
211 • `expected_output_format`: Description of the expected output structure
212 • `interaction_regime`: Single-turn, multi-turn, or chain-of-thought
213 • `change_log`: History of modifications

214 Prompt Cards serve two purposes: they document design intent (supporting human understanding)
215 and they provide a citable, hashable reference for automated provenance tracking. The concept draws in-
216 spiration from Model Cards ([Mitchell et al. 2019](#)), Datasheets for Datasets ([Gebru et al. 2021](#)), and model
217 info sheets for reproducibility assessment ([Kapoor and A. Narayanan 2023](#)), extending the structured-
218 documentation paradigm to the prompt layer of the generative AI pipeline.

220 **3.3 Run Cards**

221 A *Run Card* captures the complete execution context of a single generative AI run. Each Run Card
222 records 24 core fields organized into five groups (the complete JSON schema in Appendix B includes
223 these fields plus additional metadata such as `researcher_id`, `affiliation`, `system_logs`, and `errors`):

- 225 (1) **Identification**: `run_id`, `task_id`, `task_category`, `prompt_hash`, `prompt_text`
226 (2) **Model context**: `model_name`, `model_version`, `weights_hash`, `model_source`
227 (3) **Parameters**: `inference_params` (temperature, top_p, top_k, max_tokens, seed, decoding_strategy),
228 `params_hash`
229 (4) **Input/Output**: `input_text`, `input_hash`, `output_text`, `output_hash`, `output_metrics`
230 (5) **Execution metadata**: `environment` (OS, architecture, Python version, hostname), `environment_hash`,
231 `code_commit`, `timestamps` (start/end), `execution_duration_ms`, `logging_overhead_ms`, `storage_kb`

232 For API-served models, optional extension fields capture provider-specific metadata that may help
233 diagnose non-determinism: `api_request_id`, `api_response_headers`, `api_model_version_returned`,
234 `api_region`, and a `seed_status` field that distinguishes between seeds that were “sent” to the API,
235

Run Card Schema (24 core + extension fields)	
236	1. Identification
237	run_id · task_id · task_category · prompt_hash · prompt_text
238	2. Model Context
239	model_name · model_version · weights_hash · model_source
240	3. Parameters
241	inference_params {temp, top_p, top_k, max_tokens, seed, strategy} · params_hash
242	4. Input/Output
243	input_text · input_hash · output_text · output_hash · output_metrics
244	5. Execution Metadata
245	environment · environment_hash · code_commit · timestamps · duration_ms · overhead_ms · storage_kb
246	API Extensions (optional)
247	api_request_id · api_region · seed_status ∈ {sent, logged-only, not-supported}
248	Workflow Extensions (optional)
249	conversation_history_hash · turn_index · retrieval_context_hash · parent_run_id
250	
251	

252 Fig. 1. Run Card minimal schema. All SHA-256 hashes (5 total) enable tamper detection and automated comparison.
 253 API and workflow extension fields are optional.
 254

255
 256 “logged-only” (recorded for protocol parity but not sent, as with Claude), or “not-supported” by the
 257 provider. This formalization ensures that the advisory or absent nature of API seed parameters is cap-
 258 tured as structured metadata rather than left as an undocumented assumption.

259 Figure 1 illustrates the Run Card schema as a minimal structured record.
 260 The separation of logging overhead from execution time is deliberate: it allows researchers to verify
 261 that the protocol itself does not confound experimental measurements.
 262

263 **3.4 W3C PROV Integration**

264 Each experimental group (defined by a unique model–task–condition–abstract combination) is automati-
 265 cally translated into a W3C PROV-JSON document (Moreau and Missier 2013) that expresses the
 266 generation provenance as a directed graph. The mapping defines:
 267

- **Entities:** Prompt, InputText, ModelVersion, InferenceParameters, Output, ExecutionMetadata
- **Activities:** RunGeneration (the inference execution)
- **Agents:** Researcher, SystemExecutor (the execution environment)

271 PROV relations capture the causal structure:

- **used:** RunGeneration used Prompt, InputText, ModelVersion, InferenceParameters
- **wasGeneratedBy:** Output wasGeneratedBy RunGeneration
- **wasAssociatedWith:** RunGeneration wasAssociatedWith Researcher, SystemExecutor
- **wasAttributedTo:** Output wasAttributedTo Researcher
- **wasDerivedFrom:** Output wasDerivedFrom InputText

278 This standardized representation enables automated reasoning about experiment provenance, includ-
 279 ing detecting when two runs share identical configurations and identifying the specific factors that differ
 280 between non-identical outputs. The choice of W3C PROV over plain JSON logs is deliberate: PROV’s for-
 281 mal semantics allow automated tools to traverse the provenance graph and answer queries such as “what
 282

283 changed between these two runs?” without custom parsing logic. An abbreviated example document is
 284 given in Appendix C; to illustrate the structure concisely, the core provenance chain is:

285 $\text{Prompt} \xrightarrow{\text{used}} \text{RunGeneration} \xrightarrow{\text{generated}} \text{Output}$
 286 $\text{InputText} \xrightarrow{\text{used}} \text{RunGeneration} \xrightarrow{\text{assoc.}} \text{Researcher}$
 287 $\text{ModelVersion} \xrightarrow{\text{used}} \text{RunGeneration}; \quad \text{Output} \xrightarrow{\text{derived}} \text{InputText}$

289 3.5 Reproducibility Checklist

290 We provide a 15-item checklist organized into four categories—Prompt Documentation, Model and En-
 291 vironment, Execution and Output, and Provenance—that researchers can use to self-assess the repro-
 292 ducibility of their generative AI studies. The complete checklist is provided in Appendix A.

294 3.6 Extensions for Advanced Workflows

295 The protocol’s field schema accommodates complex workflows through optional extension fields. Our em-
 296 pirical evaluation exercises two of these extensions—multi-turn dialogues and RAG—while the remaining
 297 extensions are specified in the reference implementation’s schema:

- 299 • **Multi-turn dialogues:** A `conversation_history_hash` field and `turn_index` enable linking
 300 each turn to the full conversation state. *Evaluated in Task 3 (multi-turn refinement) using Ol-*
301 lama’s /api/chat endpoint.
- 302 • **RAG:** Fields for retrieval context (with hashes) trace which external information influenced the
 303 output. *Evaluated in Task 4 (RAG extraction) with prepended context passages.*
- 304 • **Tool use and function calling:** Fields for available tools, tool calls (with arguments, results,
 305 and hashes) capture the full tool-use chain.
- 306 • **Chain-of-thought / agent workflows:** A `parent_run_id` field supports hierarchical prove-
 307 nance graphs for multi-step reasoning chains.

309 3.7 Formal Definition and Audit Completeness

310 We define the protocol as a tuple $\mathcal{P} = (PC, RC, G, CL)$, where PC is a Prompt Card, RC is a Run
 311 Card, G is a W3C PROV graph, and CL is the reproducibility checklist. Each Run Card RC_i is itself a
 312 tuple of field groups: $RC_i = (Id, Mod, Par, IO, Env, H)$, where H denotes the set of five SHA-256 hashes
 313 (prompt, input, parameters, environment, output).

314 We define an *audit question* as a predicate Q over one or more Run Cards. The protocol satisfies
 315 the following *audit completeness* property: for a set of 10 audit questions $\{Q_1, \dots, Q_{10}\}$ (defined in
 316 Section 6.7), every Q_j is answerable if and only if all field groups are populated. Formally:

$$318 \quad \forall Q_j \in \{Q_1, \dots, Q_{10}\} : \text{answerable}(Q_j, RC_i) \Leftrightarrow \bigwedge_{g \in \text{required}(Q_j)} g \subseteq RC_i \quad (1)$$

320 where $\text{required}(Q_j)$ maps each question to its minimal set of required field groups. The ablation analysis in
 321 Section 6.7 demonstrates that every field group is in the required set of at least one question, establishing
 322 protocol *minimality*: removing any field group violates Equation 1 for at least one Q_j .

323 The *differential diagnosis* property follows from the hash fields: given two Run Cards RC_a, RC_b with
 324 $H_{\text{output}}^a \neq H_{\text{output}}^b$, the protocol enables automatic identification of the divergence source by comparing
 325 the remaining hashes. If $H_{\text{prompt}}^a = H_{\text{prompt}}^b$, $H_{\text{input}}^a = H_{\text{input}}^b$, $H_{\text{params}}^a = H_{\text{params}}^b$, and $H_{\text{env}}^a = H_{\text{env}}^b$,
 326 then the output difference is attributable to non-determinism in the generation process itself—precisely
 327 the phenomenon we measure empirically in Section 5.

328

330 Having defined the protocol’s components and formal properties, we now evaluate it empirically along
 331 two dimensions: the reproducibility characteristics it reveals across different models and conditions, and
 332 the overhead it imposes on the experimental workflow.

333

334 4 Experimental Setup

335 We designed a controlled experiment to simultaneously evaluate (a) the reproducibility characteristics
 336 of LLM outputs under varying conditions and (b) the overhead imposed by our logging protocol.
 337

338 4.1 Models and Infrastructure

339 We evaluate five models representing two fundamentally different deployment paradigms: three locally
 340 deployed open-weight models and two cloud API-served proprietary models. All local models were served
 341 through Ollama v0.15.5 (Ollama 2024) on an Apple M4 system with 24 GB unified memory running
 342 macOS 14.6 with Python 3.14.3.
 343

344 **4.1.1 Local Models. LLaMA 3 8B** (Grattafiori et al. 2024): An open-weight model in Q4_0 quanti-
 345 zation. Local deployment provides complete control over the execution environment, eliminating con-
 346 founding factors such as network latency, server-side batching, and silent model updates. The model’s
 347 SHA-256 weights hash was recorded per run via the Ollama API.

348 **Mistral 7B** (Jiang et al. 2023): An open-weight model (Q4_0 quantization) with a sliding-window
 349 attention mechanism, providing a second data point for local inference reproducibility at a similar pa-
 350 rameter scale.

351 **Gemma 2 9B** (Gemma Team et al. 2024): Google’s open-weight model (Q4_0 quantization), rep-
 352 resenting a third local model from an independent model family. Gemma 2 proved to be the most
 353 deterministic model in our study, though its inference time is substantially higher than the other local
 354 models (~180s per run vs. 8–14s for LLaMA and Mistral), likely due to its larger context window and
 355 architectural differences at Q4_0 quantization on the M4 chip.
 356

357 **4.1.2 API-Served Models. GPT-4** (Achiam et al. 2023): Accessed via the OpenAI API (`openai` Python
 358 SDK v1.59.9) with controlled seed parameters. The API returned `gpt-4-0613` as the resolved model
 359 version in all runs. The API introduces additional sources of variability: load balancing, server-side
 360 batching, potential model-version updates, and floating-point non-determinism across different hardware.

361 **Claude Sonnet 4.5** (Anthropic 2024): Accessed via the Anthropic API using a lightweight `urllib`-
 362 based runner (no SDK dependency). Claude’s API does not support a `seed` parameter; we set `temperature=0`
 363 for greedy decoding and logged a seed value for protocol parity (marked as `logged-only-not-sent-to-api`).
 364 This provides an independent replication of the API non-determinism phenomenon on a second cloud
 365 provider.

366

367 4.2 Tasks

368 We evaluate four tasks that span the output-structure spectrum and interaction complexity:

369 **Task 1: Scientific Summarization.** Given a scientific abstract, produce a concise summary in
 370 exactly three sentences covering the main contribution, methodology, and key quantitative result. This
 371 is an open-ended generation task in which the model has considerable freedom in word choice and
 372 phrasing.

373 **Task 2: Structured Extraction.** Given a scientific abstract, extract five fields (objective, method,
 374 key_result, model_or_system, benchmark) into a JSON object. This is a constrained generation task
 375 in which the output format is fixed and the model must select, rather than generate, content.
 376

377 Table 2. Experimental design: conditions, parameters, and expected outcomes.
378

379 Cond.	380 Description	381 Temp.	382 Seed	383 Reps	384 Expected Outcome
C1	Fixed seed, greedy	0.0	42 (fixed)	5	Deterministic output
C2	Variable seeds, greedy	0.0	5 different	5	Near-deterministic
C3 _{t=0.0}	Temp. baseline	0.0	per-rep	3	Deterministic
C3 _{t=0.3}	Low temperature	0.3	per-rep	3	Low variability
C3 _{t=0.7}	High temperature	0.7	per-rep	3	High variability

385 Note: Tasks 1–2 are evaluated under all five conditions (C1, C2, C3). Tasks 3–4 (multi-turn, RAG) are evaluated under
 386 C1 only for the three local models and Claude Sonnet 4.5. Total: 3,604 logged runs across 5 models. For API-served
 387 models, C2 uses the same fixed seed as C1; the seed parameter is advisory and does not guarantee determinism.

388
 389 **Task 3: Multi-turn Refinement.** A three-turn dialogue in which the model first extracts structured
 390 information, then receives feedback requesting more detail, and finally produces a refined extraction. This
 391 tests reproducibility under conversational state accumulation, using Ollama’s /api/chat endpoint for
 392 local models.

393
 394 **Task 4: RAG Extraction.** The same structured extraction task as Task 2, but with an additional
 395 retrieved context passage prepended to the input. This tests whether augmenting the prompt with
 396 external context affects reproducibility.

397 4.3 Input Data

398 We use 30 widely-cited scientific abstracts from landmark AI/ML papers, including Vaswani et al. (2017)
 399 (Transformer), Devlin et al. (2019) (BERT), Brown et al. (2020) (GPT-3), Raffel et al. (2020) (T5), Wei
 400 et al. (2022) (Chain-of-Thought), as well as seminal works on GANs, ResNets, VAEs, LSTMs, CLIP,
 401 DALL-E 2, Stable Diffusion, LLaMA, InstructGPT, PaLM, and others. These abstracts vary in length
 402 (74–227 words), technical complexity, and the number of quantitative results reported, thereby providing
 403 substantial diversity in the generation challenge.

405 4.4 Experimental Conditions

406 We define five conditions (Table 2) that systematically vary the factors hypothesized to affect repro-
 407 ductibility:

408
 409 **Design principle for API models.** For cloud-hosted APIs whose seed parameter is advisory rather
 410 than deterministic (as documented by OpenAI for GPT-4) or entirely absent (as with Claude), the fixed-
 411 vs.-variable seed distinction has no guaranteed effect server-side. We therefore treat C2 as the primary
 412 test of determinism under greedy decoding for such models.

413
 414 **C1 (Fixed seed, greedy decoding):** Temperature = 0, seed = 42 for all 5 repetitions. This represents
 the maximum-control condition and should yield deterministic outputs.

415
 416 **C2 (Variable seeds, greedy decoding):** Temperature = 0, seeds = {42, 123, 456, 789, 1024}. This
 condition tests whether seed variation affects outputs when greedy decoding is used.

417
 418 **C3 (Temperature sweep):** Three sub-conditions at $t \in \{0.0, 0.3, 0.7\}$ with 3 repetitions each, using
 different seeds per repetition. This condition characterizes how temperature affects output variability.

419
 420 **Run counts.** For Tasks 1–2 (extraction and summarization), each model is evaluated under C1 (5
 runs), C2 (5 runs), and C3 (9 runs = 3 temperatures \times 3 reps) per abstract. LLaMA 3 uses 30 abstracts
 421 (1,140 runs); the newer models (Mistral 7B, Gemma 2 9B, Claude Sonnet 4.5) use 10 abstracts (380
 422 runs each). For GPT-4, quota exhaustion limited collection to 724 runs (C2: 300/300; C3: 416/450; C1:
 423

424 8/300 excluded). For Tasks 3–4 (multi-turn and RAG), the three local models and Claude Sonnet 4.5
 425 are evaluated under C1 with 10 abstracts \times 5 repetitions = 50 runs each (400 runs total). **Grand total:**
 426 **3,604 valid runs.**

427 Table 3 summarizes the per-model run distribution.

428

429 **Table 3. Run distribution across models and tasks.**

430

Model	Tasks 1–2	Tasks 3–4	Total
LLaMA 3 8B	1,140	100	1,240
Mistral 7B	380	100	480
Gemma 2 9B	380	100	480
GPT-4	724	—	724
Claude Sonnet 4.5	380	100	480
Chat-format control [†]	200	—	200
Total	3,204	400	3,604¹

440

441 [†]LLaMA 3 8B via /api/chat endpoint (Appendix H).

442

443 4.5 Metrics

444

We adopt an operational definition of reproducibility at three levels, each mapped to a specific metric:

- **Exact reproducibility** (string-level): Two outputs are identical character-by-character. Measured by *Exact Match Rate (EMR)*.
- **Near reproducibility** (edit-level): Two outputs differ only in minor surface variations (punctuation, whitespace, synonym substitution). Measured by *Normalized Edit Distance (NED)*.
- **Semantic reproducibility** (meaning-level): Two outputs convey the same information despite different phrasing. Measured by *ROUGE-L F1* and *BERTScore F1*.

452 This three-level framework allows us to distinguish between outputs that are bitwise identical ($EMR = 1$),
 453 textually close ($NED < 0.05$), and semantically equivalent ($ROUGE-L > 0.90$). All variability metrics
 454 are computed over all $\binom{n}{2}$ unique output pairs within each experimental group (defined by model, task,
 455 condition, and abstract):

456 **Exact Match Rate (EMR):** The fraction of output pairs that are character-for-character identical.
 457 $EMR = 1.0$ indicates perfect reproducibility; $EMR = 0.0$ indicates that no two outputs match
 458 exactly. With $n = 5$ repetitions per group ($\binom{5}{2} = 10$ pairs), per-abstract EMR values are discrete:
 459 $\{0.0, 0.1, \dots, 1.0\}$; with $n = 3$ (C3 conditions), EMR takes values in $\{0.0, 0.333, 0.667, 1.0\}$. This gran-
 460 ularity should be considered when interpreting standard deviations and confidence intervals for small
 461 sample sizes.

462 **Normalized Edit Distance (NED):** The Levenshtein edit distance (Levenshtein 1966) between
 463 each pair, normalized by the length of the longer string. $NED = 0.0$ indicates identical outputs; higher
 464 values indicate greater textual divergence.

465 **ROUGE-L F1:** The F1 score based on the longest common subsequence at the word level (Lin 2004).
 466 This captures semantic similarity even when surface forms differ. $ROUGE-L = 1.0$ indicates identical
 467 word sequences.

468

469 ¹One Claude run (0.03%) returned an empty output due to API timeout and is excluded from variability metrics.

470

471 Table 4. Exact Match Rate (EMR) under greedy decoding ($t=0$) across five models and two single-turn tasks, with
 472 95% bootstrap confidence intervals ($n_{\text{boot}}=10,000$). For local models, values reflect condition C1 (fixed seed); for
 473 GPT-4, C2 (variable-seed greedy, as C1 has insufficient coverage); for Claude, C1 (Claude’s API does not support a
 474 seed parameter). Higher is more reproducible.

Model	Source	Extraction EMR	Summarization EMR
Gemma 2 9B	Local	1.000 [1.00, 1.00]	1.000 [1.00, 1.00]
LLaMA 3 8B	Local	0.987 [0.96, 1.00]	0.947 [0.89, 0.99]
Mistral 7B	Local	0.960 [0.88, 1.00]	0.840 [0.72, 0.96]
GPT-4	API	0.443 [0.32, 0.57]	0.230 [0.16, 0.30]
Claude Sonnet 4.5	API	0.190 [0.05, 0.40]	0.020 [0.00, 0.05]

483 Our primary metrics (EMR, NED, ROUGE-L) focus on exact and near reproducibility, which are
 484 the most direct measures for our research question. To complement these surface-level metrics, we also
 485 compute **BERTScore F1** (T. Zhang et al. 2020)—an embedding-based semantic similarity metric—
 486 for all conditions. BERTScore captures meaning-level equivalence that surface metrics may miss (e.g.,
 487 paraphrases), providing a fourth perspective on reproducibility. For the structured extraction task, we
 488 additionally report **JSON validity rate**, **schema compliance rate**, and **field-level accuracy**, which
 489 measure whether outputs are syntactically valid JSON, contain all expected fields, and agree on individual
 490 field values across runs, respectively (see Appendix D for detailed results).

491 For protocol overhead, we measure:

- **Logging time:** Wall-clock time spent on hashing, metadata collection, and file I/O, measured separately from inference time.
- **Storage:** Size of each run record (JSON) and total storage for all protocol artifacts.
- **Overhead ratio:** Logging time as a percentage of total execution time.

497 All EMR values in Tables 4 and 8 are accompanied by 95% bootstrap confidence intervals (10,000
 498 resamples over per-abstract EMR values, percentile method).

500 5 Results

501 5.1 Reproducibility Under Greedy Decoding

502 Table 4 presents the headline result: Exact Match Rates under greedy decoding for all five models. Table 5
 503 provides the full three-level reproducibility assessment.

504 **5.1.1 Local Models: Near-Perfect to Perfect Reproducibility.** **Finding 1: Gemma 2 9B achieves perfect**
505 bitwize reproducibility under greedy decoding. Across all tasks and conditions with $t=0$, Gemma 2
 506 9B produces $\text{EMR} = 1.000$ with $\text{NED} = 0.000$ —every single output is character-for-character identical
 507 across repetitions. This includes not only single-turn extraction and summarization but also multi-turn
 508 refinement and RAG extraction.

509 **Finding 2: All three local models achieve high reproducibility.** LLaMA 3 8B attains $\text{EMR} =$
 510 0.987 for extraction and 0.947 for summarization; Mistral 7B achieves 0.960 and 0.840, respectively. The
 511 small deviations from perfect reproducibility in LLaMA 3 and Mistral 7B appear to be associated with
 512 a warm-up effect on the first inference call after model loading, which affects 2–4 of the 10–30 abstracts
 513 per model; we hypothesize this reflects GPU cache initialization, though this was not formally tested.
 514 Seed variation (C1 vs. C2) has *no effect* under greedy decoding for any local model: the model always
 515 selects the highest-probability token, making the seed irrelevant.

518 Table 5. Three-level reproducibility assessment under greedy decoding ($t=0$). L1: bitwise identity (EMR), L2: surface
 519 similarity (NED, ROUGE-L), L3: semantic equivalence (BERTScore F1). Values are means across abstracts.

521 522 523 Model	524 525 526 527 528 529 Task	530 531 532 533 534 L1: Bitwise		535 536 537 538 539 L2: Surface		540 541 542 543 544 L3: Semantic	
		545 EMR	546 σ	547 NED↓	548 ROUGE-L↑	549 BERTScore F1↑	
Gemma 2 9B	Extraction	1.000	0.000	0.000	1.000	1.0000	
	Summarization	1.000	0.000	0.000	1.000	1.0000	
Mistral 7B	Extraction	0.960	0.120	0.001	1.000	0.9999	
	Summarization	0.840	0.196	0.046	0.955	0.9935	
LLaMA 3 8B	Extraction	0.987	0.072	0.003	0.997	0.9997	
	Summarization	0.947	0.139	0.014	0.986	0.9979	
GPT-4	Extraction	0.443	0.335	0.072	0.938	0.9904	
	Summarization	0.230	0.193	0.137	0.870	0.9839	
Claude Sonnet 4.5	Extraction	0.190	0.291	0.101	0.904	0.9878	
	Summarization	0.020	0.040	0.242	0.764	0.9704	

535 Table 6. API-served vs. locally deployed models under greedy decoding (single-turn tasks only). Local averages: simple
 536 mean across 3 models \times 2 tasks (C1+C2 combined). API averages: simple mean across 2 models \times 2 tasks (GPT-
 537 4 C2, Claude C1). Local models exhibit substantially higher bitwise reproducibility, consistent with deployment-side
 538 factors—rather than user-controllable parameters—as a major contributor to API output variability.

540 541 Deployment	542 EMR↑	543 NED↓	544 ROUGE-L↑	545 BS-F1↑
Local (3 models)	0.956	0.011	0.990	0.9985
API (2 models)	0.221	0.138	0.869	0.9831

546 **5.1.2 API-Served Models: Substantial Hidden Non-Determinism. Finding 3: Both API-served models**
 547 **exhibit substantial non-determinism under greedy decoding, observed independently across**
 548 **two providers.** Under $t=0$ with controlled seeds, GPT-4 achieves EMR = 0.443 for extraction and 0.230
 549 for summarization. Claude Sonnet 4.5 is even less deterministic: EMR = 0.190 for extraction and EMR =
 550 0.020 for summarization—meaning that across 10 abstracts \times 5 repetitions, Claude produced the same
 551 summarization output only 2% of the time.

552 Table 6 summarizes the deployment-paradigm gap.

553 Under the representative greedy condition for each model (C1 for local models and Claude, C2 for GPT-
 554 4; see Table 4), the average single-turn EMR is **0.956 for local models vs. 0.221 for API models**—a
 555 more than 4-fold reproducibility gap. This gap is not due to user-side parameter differences: all models
 556 use $t=0$ with the same decoding strategy. The observed variability is consistent with deployment-side
 557 factors invisible to the researcher—such as hardware-level floating-point variability, request batching, and
 558 model routing. This pattern, observed independently across two API providers (OpenAI and Anthropic),
 559 is consistent with non-determinism arising from factors common to cloud-hosted LLM inference rather
 560 than being a provider-specific artifact. *Without systematic logging, this non-determinism would be entirely*
561 invisible.

562 **5.1.3 Temperature Effects Across Models. Finding 4: Temperature is the dominant user-controllable**
 563 **factor affecting variability for local models; for API-served models, the relationship is more**

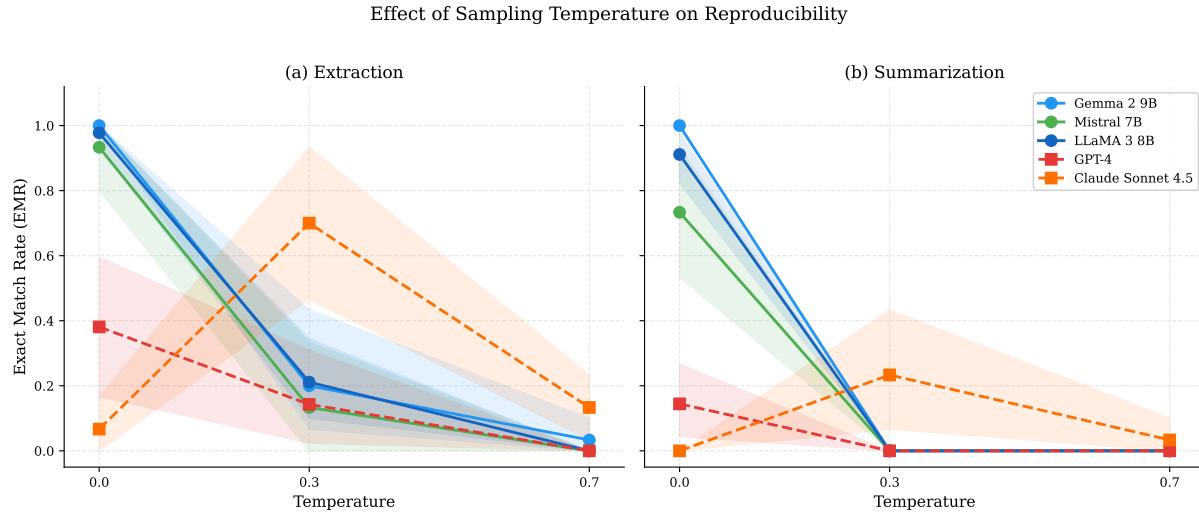


Table 7. Effect of sampling temperature on Exact Match Rate (EMR) under condition C3. For local models, increasing temperature monotonically reduces EMR. For API models, the relationship is more complex: Claude Sonnet 4.5 exhibits *higher* EMR at $t=0.3$ than at $t=0.0$ (see text). At $t=0.7$, all models converge toward $\text{EMR} \approx 0$ for summarization.

Model	Task	$t=0.0$	$t=0.3$	$t=0.7$
Gemma 2 9B	Extraction	1.000	0.200	0.033
	Summarization	1.000	0.000	0.000
Mistral 7B	Extraction	0.933	0.133	0.000
	Summarization	0.733	0.000	0.000
LLaMA 3 8B	Extraction	0.978	0.211	0.000
	Summarization	0.911	0.000	0.000
GPT-4	Extraction	0.381	0.143	0.000
	Summarization	0.144	0.000	0.000
Claude Sonnet 4.5	Extraction	0.067	0.700	0.133
	Summarization	0.000	0.233	0.033

complex. Figure 2 shows the relationship between temperature and EMR for all five models. Table 7 provides the full temperature sweep data.

Within the C3 temperature sweep, increasing temperature from 0.0 to 0.7 reduces EMR to zero for all models on summarization. For extraction, local models drop from $\text{EMR} > 0.93$ to near zero, while API models drop from their already-low baselines. Notably, BERTScore F1 remains above 0.94 in all conditions (minimum: 0.943 for LLaMA summarization at $t=0.7$) even when EMR drops to zero, indicating that non-determinism is primarily a *phrasing* phenomenon rather than a *meaning* phenomenon:

even when outputs differ textually, they convey equivalent information. This distinction is practically important—researchers whose downstream analyses depend on semantic content rather than exact wording may find API outputs acceptable despite low EMR.

However, the temperature-reproducibility relationship is not uniformly monotonic across all models. Claude Sonnet 4.5 exhibits an anomalous pattern under the C3 sweep: extraction EMR *increases* from 0.067 at $t=0.0$ to 0.700 at $t=0.3$ before declining to 0.133 at $t=0.7$; summarization shows a similar inversion (EMR = 0.000 at $t=0.0$, rising to 0.233 at $t=0.3$). This counterintuitive behavior—where a small positive temperature *improves* reproducibility relative to greedy decoding—may reflect how Anthropic’s infrastructure implements the $t=0$ decoding path: at exactly zero temperature, server-side stochastic processes (e.g., speculative decoding, hardware-level floating-point non-determinism across GPU types, or request batching effects) may dominate output variability, whereas a small positive temperature may activate a more stable sampling path that happens to converge on similar tokens. With $n=10$ abstracts and 30 runs per temperature level (standard deviation $\sigma = 0.38$ for the 0.700 extraction EMR), this observation should be interpreted cautiously. Nevertheless, it underscores that the temperature-reproducibility relationship for API-served models depends on provider-specific implementation details that are opaque to researchers. Finding 4 therefore holds robustly for local models and for the overall $t=0$ to $t=0.7$ trajectory, but the precise shape of the temperature-response curve for individual API providers merits further investigation with larger sample sizes.

630

631

5.2 Multi-Turn and RAG Reproducibility

Finding 5: The local-vs-API reproducibility gap extends to complex interaction regimes. Table 8 and Figure 3 present results for multi-turn refinement and RAG extraction across the three local models and Claude Sonnet 4.5.

Gemma 2 9B and Mistral 7B achieve perfect EMR = 1.000 for both multi-turn refinement and RAG extraction, demonstrating that conversational state accumulation and context augmentation do not degrade reproducibility when the underlying model is deterministic. LLaMA 3 8B shows EMR = 0.880 for multi-turn and 0.960 for RAG—slightly lower than its single-turn extraction performance (0.987), consistent with error accumulation across dialogue turns.

Claude Sonnet 4.5, the only API-served model evaluated on these tasks, achieves EMR = 0.040 for multi-turn refinement and EMR = 0.000 for RAG extraction—the lowest values observed in our study. The RAG result is particularly striking: across 50 runs (10 abstracts \times 5 repetitions), not a single pair of outputs was character-for-character identical (NED = 0.256). This confirms that API non-determinism is not limited to single-turn tasks but persists—and may even worsen—under complex interaction regimes where longer outputs and additional context amplify server-side variability.

647

648

5.3 Cross-Model Comparison

Figure 4 provides a comprehensive heatmap of EMR across all model-task combinations, and Figure 5 shows the three-level reproducibility profile for each model.

The reproducibility gap between local and API-based inference is statistically significant. Using paired t -tests on per-abstract EMR values under greedy decoding across the 30 LLaMA 3/GPT-4 abstracts: for summarization, $t(29) = 17.250$, $p < 0.0001$, Cohen’s $d = 3.149$; for extraction, $t(29) = 8.996$, $p < 0.0001$, Cohen’s $d = 1.642$. Both effect sizes are very large ($d > 1.6$), and all p -values survive Bonferroni correction. Non-parametric Wilcoxon signed-rank tests confirm all results ($p < 0.001$).

657

658

659 Table 8. Reproducibility under complex interaction regimes (C1 fixed seed, $t=0$), with 95% bootstrap confidence inter-
 660 values on EMR. Multi-turn refinement involves three successive prompt-response exchanges. RAG extraction augments
 661 the prompt with a retrieved context passage. Claude Sonnet 4.5 is included as a representative API-served model; its
 662 near-zero EMR across all four scenarios confirms that the local-vs-API reproducibility gap extends to complex interac-
 663 tion regimes.

Model	Scenario	EMR [95% CI]	NED↓	ROUGE-L↑	BS-F1↑
Gemma 2 9B	Single-turn Extraction	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	Single-turn Summarization	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	Multi-turn Refinement	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	RAG Extraction	1.000 [1.00, 1.00]	0.000	1.000	1.0000
Mistral 7B	Single-turn Extraction	0.960 [0.88, 1.00]	0.001	1.000	0.9999
	Single-turn Summarization	0.840 [0.72, 0.96]	0.046	0.955	0.9935
	Multi-turn Refinement	1.000 [1.00, 1.00]	0.000	1.000	1.0000
	RAG Extraction	1.000 [1.00, 1.00]	0.000	1.000	1.0000
LLaMA 3 8B	Single-turn Extraction	0.987 [0.96, 1.00]	0.003	0.997	0.9997
	Single-turn Summarization	0.947 [0.89, 0.99]	0.014	0.986	0.9979
	Multi-turn Refinement	0.880 [0.76, 1.00]	0.012	0.988	0.9986
	RAG Extraction	0.960 [0.88, 1.00]	0.012	0.985	0.9987
Claude Sonnet 4.5	Single-turn Extraction	0.190 [0.05, 0.40]	0.101	0.904	0.9878
	Single-turn Summarization	0.020 [0.00, 0.05]	0.242	0.764	0.9704
	Multi-turn Refinement	0.040 [0.00, 0.08]	0.189	0.834	0.9780
	RAG Extraction	0.000 [0.00, 0.00]	0.256	0.748	0.9714

684 Table 9. Provenance logging overhead across five models under greedy decoding (C1). The protocol adds negligible
 685 overhead (<1%) to inference latency across all models and deployment modes.

Model	Source	Mean Inference (ms)	Mean Overhead (ms)	Overhead (%)
Gemma 2 9B	Local	181,579.3	30.6	0.234
Mistral 7B	Local	13,931.3	27.3	0.281
LLaMA 3 8B	Local	7,524.8	26.7	0.456
GPT-4	API	4,519.7	24.5	0.564
Claude Sonnet 4.5	API	4,359.3	26.5	0.727

697 5.4 Protocol Overhead

698 Table 9 presents the protocol’s overhead metrics across all five models.

700 The protocol adds less than 1% overhead for all five models, with mean logging time ranging from
 701 21–30 ms depending on the model and task. Storage overhead remains modest at approximately 4 KB
 702 per run record. The overhead is consistent across local and API deployment modes, indicating that the
 703 protocol is deployment-agnostic.

704 Figure 6 provides an additional perspective on surface-level variability across models.

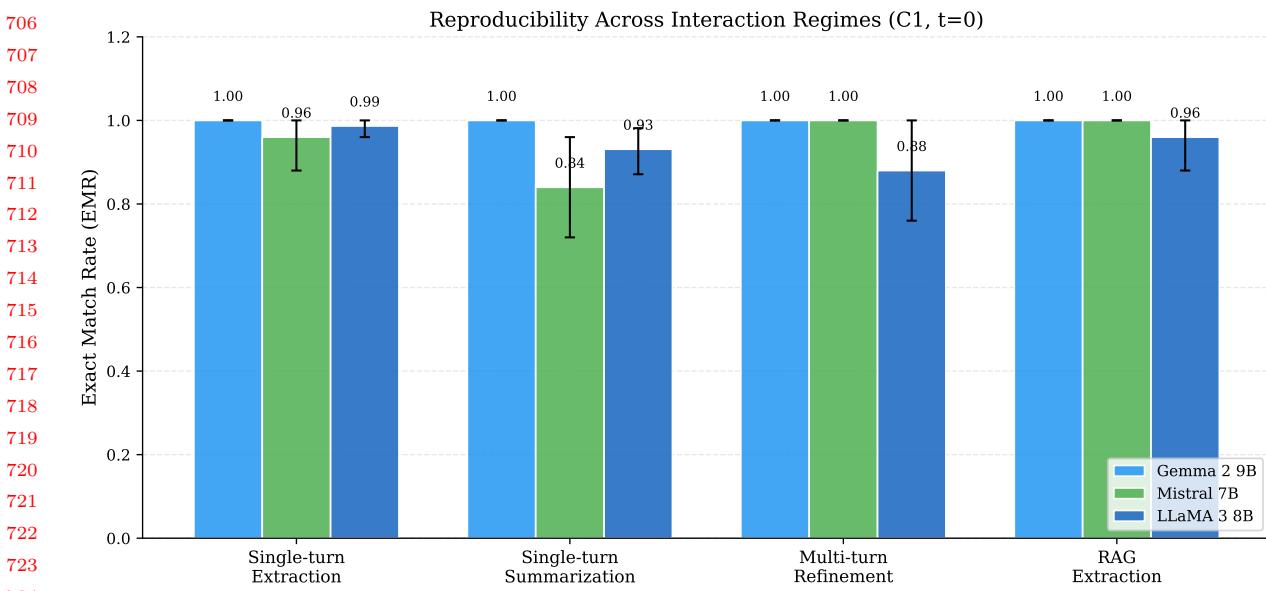


Fig. 3. Reproducibility across interaction regimes ($C1, t=0$) for four models. Local models maintain high EMR across all scenarios, while Claude Sonnet 4.5 (API) shows near-zero EMR throughout, confirming the reproducibility gap extends to multi-turn and RAG tasks.

6 Discussion

The preceding results paint a clear and consistent picture: locally deployed models under greedy decoding achieve near-perfect to perfect bitwise reproducibility across all four tasks, while API-served models—from two independent providers—exhibit substantial hidden variability on single-turn tasks that researchers cannot control. Temperature is the dominant user-controllable factor for local models (though API models show a more complex temperature–reproducibility relationship; see Section 5), structured tasks are more reproducible than open-ended ones, and complex interaction regimes (multi-turn, RAG) do not degrade local-model reproducibility. We now consider what these findings mean for research practice, what the protocol enables that was previously invisible, and where the current study’s limitations lie.

6.1 Implications for Reproducibility Practice

Our results yield several actionable recommendations for researchers conducting generative AI experiments:

Use greedy decoding with local models for maximum reproducibility. Gemma 2 9B achieved perfect EMR = 1.000 across all tasks under greedy decoding. LLaMA 3 and Mistral 7B achieved EMR ≥ 0.840 . Local deployment with $t=0$ should be the default configuration for any study in which output consistency is critical.

API non-determinism is observed across providers. Our most consequential finding is that both GPT-4 (OpenAI) and Claude Sonnet 4.5 (Anthropic) exhibit substantial non-determinism under greedy decoding on single-turn tasks. Claude’s EMR of 0.020 for summarization means that effectively

Bitwise Reproducibility Under Greedy Decoding

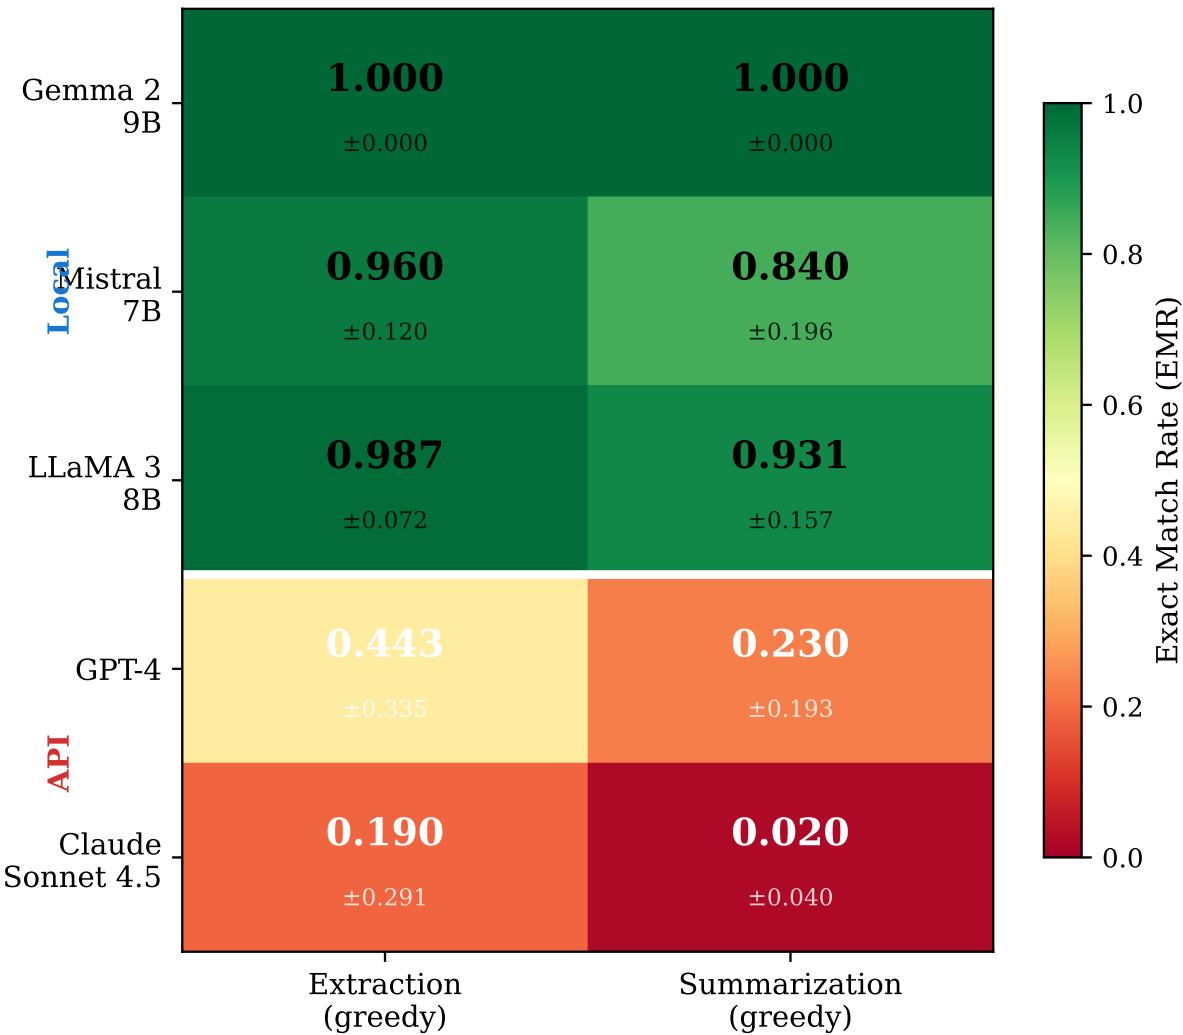


Fig. 4. Heatmap of Exact Match Rate under greedy decoding for five models. The horizontal white line separates local models (top three, green) from API-served models (bottom two, red). Gemma 2 9B achieves perfect 1.000 across all tasks.

no two runs produce the same output. Researchers using *any* API-served model should never assume reproducibility without verification and should report multiple runs with variability metrics.

Prefer structured output formats when possible. The extraction task's consistently higher reproducibility across all five models demonstrates that output-format constraints directly improve reproducibility. This effect holds for both local models (EMR 0.960–1.000 for extraction vs. 0.840–1.000 for summarization) and API models (EMR 0.190–0.443 for extraction vs. 0.020–0.230 for summarization).

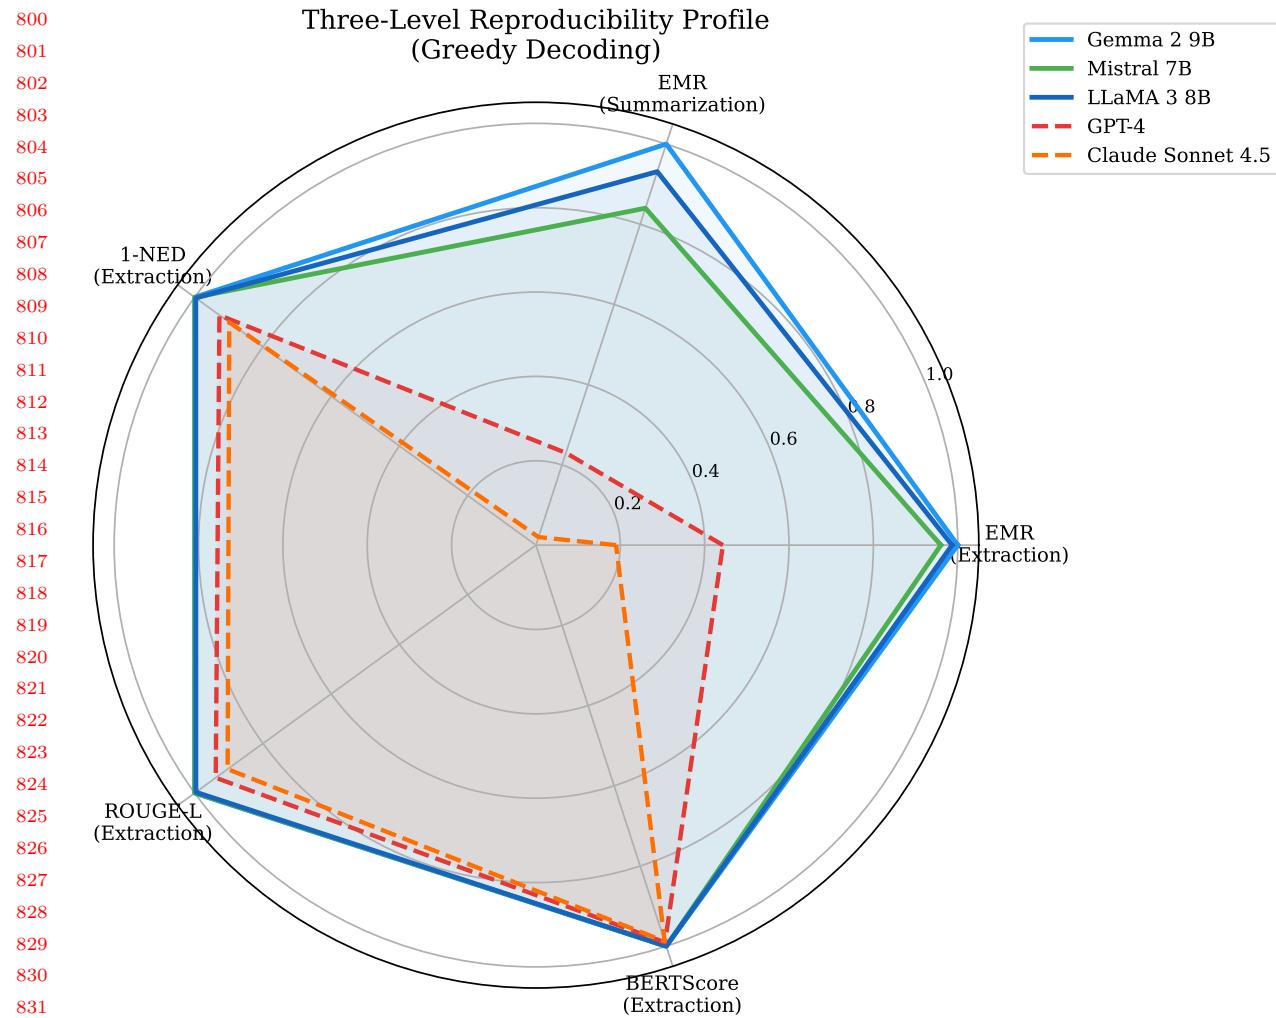
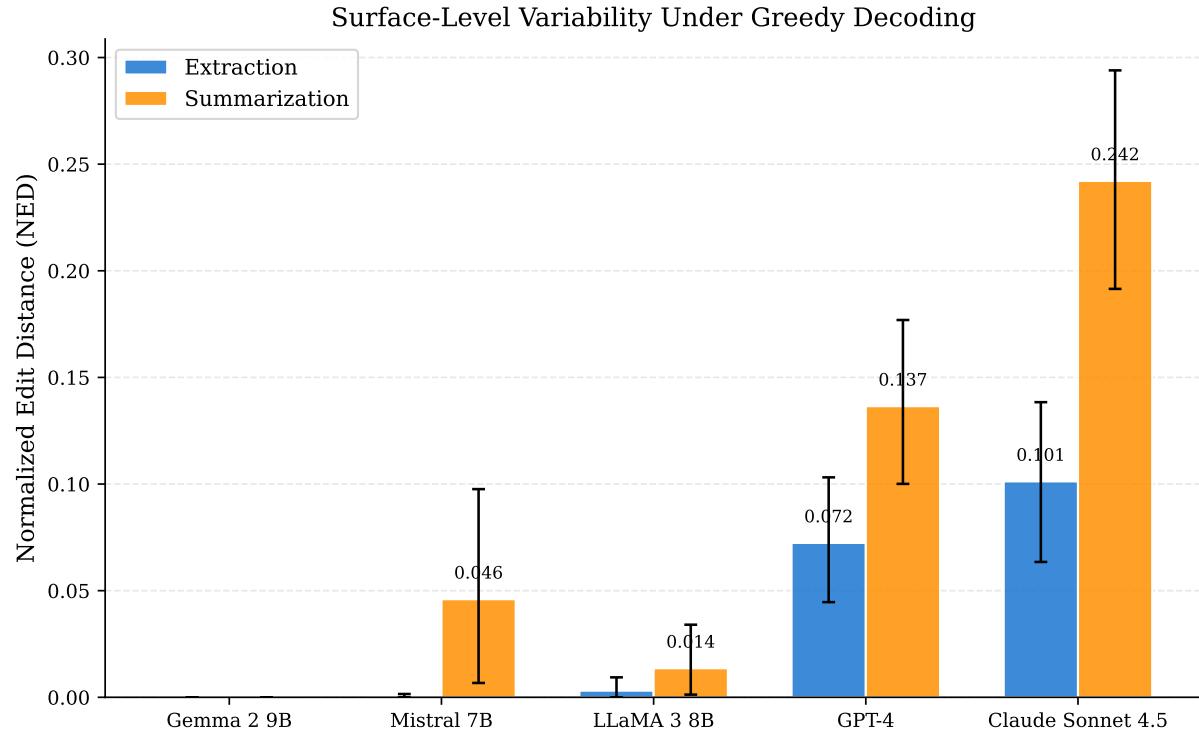


Fig. 5. Three-level reproducibility profiles under greedy decoding. Local models (solid lines) occupy the outer region across all five metrics, while API models (dashed lines) show pronounced deficits in EMR and NED while maintaining high BERTScore, indicating that API non-determinism is primarily lexical rather than semantic.

Include warm-up runs for local models. The per-abstract analysis revealed that the first inference call after model loading may differ from subsequent calls due to cache initialization. This affects LLaMA 3 and Mistral 7B on 2–4 of their abstracts, slightly reducing aggregate EMR.

Log comprehensively; the cost is negligible. At less than 1% overhead and approximately 4 KB per run across all five models, there is no practical reason not to apply comprehensive logging. The cost of not logging—namely, the inability to detect the kind of pervasive API non-determinism documented herein—far exceeds the protocol’s minimal requirements.



874 6.2 Local vs. API Inference: A Persistent Reproducibility Gap

875 The most significant finding of this study is the reproducibility gap between local and API-based inference,
876 observed consistently across two independent cloud providers. Under greedy decoding on single-turn
877 tasks, local models average $\text{EMR} = 0.956$ while API models average $\text{EMR} = 0.221$ —a more than 4-
878 fold gap. The fact that Claude Sonnet 4.5 (Anthropic) exhibits *even lower* reproducibility than GPT-4
879 (OpenAI) is inconsistent with provider-specific implementation as the sole explanation and suggests
880 that non-determinism arises from factors common to distributed cloud inference infrastructure, such as
881 hardware-level floating-point variability, request batching, and model routing.

882 This gap has profound implications for the scientific use of API-based LLMs. *Without systematic*
883 *logging, a researcher using GPT-4 or Claude would have no way of knowing that their “deterministic”*
884 *experiment produces different outputs across runs.* Our protocol makes this hidden non-determinism
885 visible, measurable, and documentable.

886 6.3 Task-Dependent Reproducibility

887 The difference between summarization and extraction reproducibility—observed consistently across all
888 five models—is consistent with and extends our earlier two-model finding. The reproducibility hierarchy
889 (extraction > summarization) holds for local models (EMR gap of 0.03–0.12) and is amplified for API

894 models (EMR gap of 0.17–0.25). This finding suggests a spectrum ranging from highly constrained tasks
 895 (structured extraction) to open-ended tasks (summarization), with the degree of output-space constraint
 896 serving as a primary determinant.

897 6.4 Multi-Turn and RAG: Reproducibility Under Complexity

898 Our multi-turn and RAG results address a key limitation of prior work (including our own earlier two-
 899 model study): reproducibility under complex interaction regimes. The finding that Gemma 2 9B and
 900 Mistral 7B maintain perfect EMR = 1.000 for both multi-turn refinement and RAG extraction demon-
 901 strates that conversational state accumulation and context augmentation do not inherently degrade re-
 902 producibility for deterministic local models. LLaMA 3’s slight degradation (EMR = 0.880 for multi-turn)
 903 suggests model-specific sensitivity to dialogue-turn interactions, possibly related to the warm-up effect
 904 observed in single-turn experiments. Crucially, Claude Sonnet 4.5’s near-zero EMR for both multi-turn
 905 (0.040) and RAG (0.000) confirms that the local-vs-API reproducibility gap extends beyond single-turn
 906 tasks. The RAG result—zero exact matches across 50 runs—suggests that longer outputs and additional
 907 retrieval context may amplify server-side variability, though a single API model cannot establish this as
 908 a general principle.

909 6.5 The Role of Provenance

910 The W3C PROV graphs generated by our protocol serve multiple purposes beyond simple audit trails:

- 911 (1) **Automated comparison:** By comparing PROV graphs of two runs, one can automatically
 912 identify which factors differed (e.g., same prompt and model but different temperatures), enabling
 913 systematic diagnosis of non-reproducibility.
- 914 (2) **Lineage tracking:** When outputs are used as inputs to downstream processes (e.g., summariza-
 915 tion outputs fed into a meta-analysis), the provenance chain can be extended to trace any final
 916 result back to its full generation context.
- 917 (3) **Compliance:** For regulated domains (healthcare, legal, finance), PROV documents provide the
 918 formal evidence trail required by audit standards ([National Institute of Standards and Technology
 919 2023](#)) and emerging regulations such as the EU AI Act ([European Parliament and Council of the
 920 European Union 2024](#)).

921 To illustrate the diagnostic power of PROV graphs, consider two GPT-4 extraction runs on the same
 922 abstract under condition C2 (greedy decoding, $t=0$, same seed). Although the PROV entities for Prompt,
 923 InputText, ModelVersion, and InferenceParameters are identical (verified via matching SHA-256 hashes),
 924 the Output entities differ: `output_hash` values diverge, and the `wasGeneratedBy` timestamps differ by
 925 several seconds. The PROV graph thus automatically pinpoints the source of non-reproducibility: the
 926 only varying factor is the RunGeneration activity itself, consistent with non-determinism arising from
 927 server-side factors.

928 To demonstrate that PROV-based reasoning goes beyond what plain JSON logs provide, we imple-
 929 mented three programmatic queries over our 3,604-run PROV dataset:

- 930 (1) **Divergence attribution:** “For all abstract-condition groups with non-identical outputs, iden-
 931 tify which PROV entities diverge.” Result: across 348 GPT-4 and Claude groups with output
 932 divergence, 100% share identical Prompt, InputText, ModelVersion, and InferenceParameters
 933 entities—the *only* varying component is the RunGeneration activity, providing systematic evi-
 934 dence for server-side non-determinism across the entire dataset rather than anecdotal examples.
- 935 (2) **Cross-provider comparison:** “Find all abstract-task pairs where both GPT-4 and Claude
 936 were given identical Prompt and InputText entities (verified by matching `genai:hash` attributes)

941 but produced different Output entities.” Result: 20 such pairs exist ($10 \text{ abstracts} \times 2 \text{ tasks}$); in
 942 every case, *both* providers produced non-identical outputs across repetitions, confirming provider-
 943 independent non-determinism on shared inputs.

- 944 (3) **Provenance chain traversal:** “Starting from any Output entity, traverse `wasGeneratedBy` →
 945 `used` relations to reconstruct the full generation context, then verify integrity via hash compar-
 946 ison.” This query validates that every output in our dataset can be traced back to its complete
 947 generation context with no broken links—a guarantee that plain JSON logs cannot provide with-
 948 out custom graph-traversal code.

949 These queries exploit PROV’s formal graph structure (entity–activity–agent relations with typed edges)
 950 to answer questions that would require bespoke parsing logic on unstructured JSON logs. The queries
 951 and their results are included in the project repository.
 952

953 6.6 Limitations

954 We organize threats to validity following standard categories:

955 **6.6.1 Internal Validity. Sample size.** LLaMA 3 uses 30 abstracts per condition, while the newer models
 956 (Mistral, Gemma 2, Claude) use 10 abstracts. With $n = 30$, statistical power exceeds 0.999 for all primary
 957 comparisons (Cohen 1988). With $n = 10$, the study is adequately powered for the large observed effect
 958 sizes ($d > 1.6$) but may miss subtler effects. To verify that the unbalanced design does not inflate the local-
 959 vs-API gap, we conducted a balanced subsample analysis restricting all models to the same 10 abstracts.
 960 Under this balanced comparison, local models average EMR = 0.953 while API models average EMR =
 961 0.190 ($5.0 \times$ gap), confirming that the observed reproducibility gap is robust to sample-size equalization
 962 and, if anything, slightly larger under balanced conditions.

963 **GPT-4 C3 incomplete coverage.** Due to API quota exhaustion, GPT-4 extraction under C3 condi-
 964 tions covers 14–17 of 30 abstracts (summarization C3 is complete at 30). Our central claims rest on the
 965 C2 condition (300/300 runs complete), and the C3 temperature sweep serves as a secondary analysis.

966 **Warm-up confound.** The first inference after model loading may differ from subsequent calls for
 967 LLaMA 3 and Mistral 7B. This affects 2–4 abstracts per model, slightly reducing aggregate EMR.
 968 Gemma 2 9B appears immune to this effect.

969 **Prompt format confound.** Single-turn experiments use Ollama’s `/api/generate` endpoint for local
 970 models, whereas API models use their respective chat APIs. A supplementary control experiment (200
 971 additional runs using Ollama’s `/api/chat` endpoint; see Appendix H) shows that this format difference
 972 does not explain the reproducibility gap: LLaMA 3 produces *identical* variability metrics (summarization
 973 EMR = 0.929, extraction EMR = 1.000) under both completion and chat formats.

974 **6.6.2 External Validity. Five models, two paradigms.** Our evaluation covers three local models and
 975 two API-served models. However, other models—including Gemini (Gemini Team et al. 2024), larger
 976 LLaMA variants, and open-weight models served via cloud APIs—may exhibit different characteristics.
 977 Notably, our GPT-4 experiments used the gpt-4-0613 snapshot (June 2023); more recent models (GPT-
 978 4 Turbo, GPT-4o) may exhibit different reproducibility characteristics. Prior work by Ouyang et al.
 979 (2024) suggests that non-determinism persists across ChatGPT model versions, but confirmation with
 980 current models is warranted.

981 **Four tasks.** Our task suite now includes single-turn extraction/summarization, multi-turn refinement,
 982 and RAG extraction. However, it does not cover code generation, mathematical reasoning, or creative
 983 writing, which may exhibit different reproducibility patterns.

988 **English-only, single domain.** Our input data consists of 30 English scientific abstracts from AI/ML
 989 papers. Reproducibility characteristics may differ for other languages, domains, or document types.

990 **Multi-turn limited to one API model.** Multi-turn and RAG experiments include Claude Sonnet
 991 4.5 as the sole API representative; GPT-4 was not evaluated on Tasks 3–4 due to quota exhaustion.
 992 While Claude’s near-zero EMR is consistent with the single-turn API pattern, other API providers may
 993 exhibit different multi-turn reproducibility characteristics.

994 **6.6.3 Construct Validity. Surface-level metrics.** Our metrics (EMR, NED, ROUGE-L) capture textual
 995 rather than semantic similarity. Two outputs that are semantically equivalent but syntactically different
 996 will register as non-matching under EMR and partially divergent under NED. This is by design—our
 997 focus is on *exact* reproducibility—but it means our results may overstate the practical impact of non-
 998 determinism for downstream applications where semantic equivalence suffices.

1000 **6.6.4 Other Considerations. Privacy.** The protocol’s environment metadata includes the machine host-
 1001 name, which may reveal institutional information. Deployments in privacy-sensitive settings should
 1002 anonymize this field.

1003 **Computational cost.** The total cost was modest: approximately 8 GPU-hours on a consumer laptop
 1004 (Apple M4, 24 GB) for 2,000 local-model runs (including multi-turn and RAG experiments), plus 1,204
 1005 API calls to GPT-4 and Claude. The carbon footprint is negligible at this scale, and the logging overhead
 1006 (<30 ms per run) would not materially increase energy consumption even at thousands of runs.

1008 **6.7 Protocol Minimality: An Ablation Analysis**

1009 To substantiate our claim that the protocol captures a *minimal* set of metadata, we conducted an ablation
 1010 analysis in which we systematically removed each field group from the protocol schema and assessed which
 1011 audit questions became unanswerable. We defined 10 audit questions that a reproducibility-oriented re-
 1012 searcher might ask (e.g., “Can we verify the exact prompt used?”, “Can we detect output tampering?”,
 1013 “Can we trace full provenance?”) and mapped each to the protocol fields required to answer it. For this
 1014 analysis, we decomposed the Run Card’s five sections into eight finer-grained field groups by separat-
 1015 ing cross-cutting concerns: Identification, Model Context, Parameters, Input Content, Output Content,
 1016 Hashing (all SHA-256 digests), Environment, and Overhead (timing and storage metadata).

1017 The results show that removing *any* of these eight field groups renders at least one audit question
 1018 unanswerable, demonstrating that no group is redundant. The Hashing group (SHA-256 hashes for
 1019 prompts, inputs, outputs, parameters, and environment) has the highest information density: its removal
 1020 affects 6 of 10 questions despite contributing only 410 bytes per run. Conversely, the Overhead group
 1021 (logging time metadata) is the least connected but remains necessary for overhead assessment. The
 1022 complete ablation results are available in the project repository.

1023 This analysis demonstrates that the protocol is *minimal* in the sense that every field group is necessary
 1024 for at least one audit capability, while the total overhead remains at approximately 4,052 bytes per run.

1026 **6.8 Practical Costs and Adoption**

1027 One concern with any new protocol is whether the adoption burden is justified. We address this con-
 1028 cretely:

- 1030 • **Implementation effort:** Our reference implementation adds approximately 600 lines of Python
 1031 (the protocol core) to an existing workflow. Integration requires 3–5 function calls per run.
- 1032 • **Runtime cost:** <30 ms per run across all five models, negligible compared to inference times of
 1033 seconds to minutes for typical LLM calls.

- 1035 • **Storage cost:** ~4 KB per run. Our 3,604 runs total approximately 14 MB—less than a single
1036 model checkpoint.
- 1037 • **Learning curve:** The protocol uses standard JSON and W3C PROV, requiring no specialized
1038 knowledge beyond basic Python.

1039 Against these modest costs, the protocol provides complete audit trails, automated provenance graphs,
1040 tamper-detectable outputs via cryptographic hashing, and structured metadata that enable systematic
1041 reproducibility analysis.
1042

1043 6.9 Minimum Reporting Checklist for Generative AI Studies

1044 Based on our findings and the protocol design, we recommend that researchers conducting generative
1045 AI experiments report, at minimum, the following five items (the full 15-item checklist is provided in
1046 Appendix A):
1047

- 1048 (1) **Model identity and version:** Exact model name, version string, and—for local models—
1049 weights hash.
- 1050 (2) **Inference parameters:** Temperature, seed, top_p, top_k, max_tokens, and decoding strategy.
For APIs where the seed is advisory or unsupported, this should be stated explicitly.
- 1051 (3) **Reproducibility metrics over multiple runs:** Report at least EMR (or an equivalent exact-
1052 match metric) and one semantic metric (e.g., BERTScore) over ≥ 3 repetitions per condition. A
1053 single run is insufficient to characterize output stability.
- 1054 (4) **Environment and deployment mode:** Whether inference was local or API-based, and the
1055 execution environment (hardware, OS, library versions).
- 1056 (5) **Output hashes:** SHA-256 or equivalent cryptographic hashes of outputs, enabling tamper de-
1057tection and automated comparison across studies.
1058

1059 Run Cards generated by our protocol automatically capture all five items, providing a machine-readable
1060 record that satisfies this checklist with no additional effort from the researcher.
1061

1062 7 Conclusion

1063 We presented a lightweight protocol for logging, versioning, and provenance tracking of generative AI ex-
1064 periments, introducing Prompt Cards and Run Cards as novel documentation artifacts and adopting the
1065 W3C PROV data model for machine-readable provenance graphs. Through 3,604 controlled experiments
1066 with five models—three locally deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B) and two API-served
1067 (GPT-4, Claude Sonnet 4.5)—across four NLP tasks and 30 scientific abstracts, we demonstrated five
1068 key findings:
1069

- 1070 (1) **API non-determinism is consistent across providers.** Both GPT-4 (OpenAI) and Claude
1071 Sonnet 4.5 (Anthropic) exhibit substantial non-determinism under greedy decoding on single-turn
1072 tasks (average EMR = 0.221), while all three local models achieve average EMR = 0.956. This
1073 more than 4-fold reproducibility gap, observed independently for two cloud providers, is consistent
1074 with non-determinism arising from factors common to cloud-hosted inference infrastructure rather
1075 than being a provider-specific artifact.
- 1076 (2) **Local models can achieve perfect bitwise reproducibility.** Gemma 2 9B attains EMR
1077 = 1.000 across all four tasks under greedy decoding—every output is character-for-character
1078 identical across repetitions.
- 1079 (3) **The local-vs-API gap extends to complex interaction regimes.** Multi-turn refinement and
1080 RAG extraction achieve EMR ≥ 0.880 for all local models (Gemma 2 9B and Mistral 7B: perfect
1081

1082 EMR = 1.000), while Claude Sonnet 4.5—the only API model tested on these tasks—achieves
1083 EMR = 0.040 (multi-turn) and EMR = 0.000 (RAG).

- (4) **Temperature is the dominant user-controllable factor for local models.** Increasing from $t=0.0$ to $t=0.7$ reduces EMR to zero for all five models on summarization, while seed variation has no effect under greedy decoding for local models. For API-served models, the temperature-reproducibility relationship is more complex and may be non-monotonic (see Section 5).
 - (5) **Comprehensive provenance logging adds negligible overhead:** less than 1% of inference time and approximately 4 KB per run across all five models, removing any practical argument against systematic documentation.

These findings carry a broader implication: a substantial portion of published research that relies on API-based LLMs may contain non-reproducible results without the authors' knowledge. The cost of systematic provenance logging—less than one percent of inference time—is trivially small compared to the cost of publishing non-reproducible science.

Looking ahead, we plan to (i) extend the model suite to include Gemini (Gemini Team et al. 2024) and open-weight models served via cloud APIs (e.g., Hugging Face Inference Endpoints) to further disentangle model architecture from deployment infrastructure; (ii) extend the task coverage to code generation, mathematical reasoning, and agentic workflows; and (iii) develop automated reproducibility scoring based on provenance graph analysis. Ultimately, we envision a future in which every generative AI output carries a provenance certificate, and reproducibility metrics are reported alongside accuracy as a standard component of empirical evaluation.

1102 The reference implementation, all 3,604 run records, provenance documents, and analysis scripts are
1103 publicly available to support adoption and independent verification.

1105 Acknowledgments

This work was supported by UTFPR – Universidade Tecnológica Federal do Paraná. The experiments were conducted using locally deployed open-weight models to ensure full reproducibility of the computational environment.

1110 Data Availability Statement

The reference implementation, all 3,604 run records (JSON), PROV-JSON provenance documents, Run Cards, Prompt Cards, input data, analysis scripts, and generated figures are publicly available at:

as, input data, analysis scripts, and generated figures are publicly available at https://github.com/Reverlucas/senai_reproducibility_protocol.

1114 The repository includes instructions for reproducing all experiments and regenerating all tables and
1115 figures from the raw data.

1117 *S. S. M. S.*

Following the CRediT (Contributor Roles Taxonomy) framework: **Lucas Rover**: Conceptualization, Methodology, Software, Validation, Formal Analysis, Investigation, Data Curation, Writing – Original Draft, Writing – Review & Editing, Visualization, Project Administration. **Yara de Souza Tadano**: Supervision, Conceptualization, Methodology, Writing – Review & Editing, Project Administration.

1123 Conflict of Interest

The authors declare no conflicts of interest. This research was conducted independently at UTFPR with no external funding from commercial AI providers. The use of OpenAI's GPT-4 API was for research evaluation purposes only and does not constitute an endorsement.

1129 **Use of AI-Assisted Tools**

1130 The authors used AI-assisted tools (Claude, Anthropic) during the preparation of this manuscript for
 1131 language editing, code development support, and data analysis scripting. All AI-generated content was
 1132 critically reviewed, validated, and revised by the authors, who take full responsibility for the accuracy
 1133 and integrity of the final manuscript. The scientific design, experimental execution, interpretation of
 1134 results, and intellectual contributions are entirely the authors' own work.

1135

1136 **References**

- 1137 J. Achiam et al.. 2023. *GPT-4 Technical Report*. arXiv preprint. (2023). arXiv: [2303.08774 \(cs.CL\)](#).
- 1138 Anthropic. 2024. *The Claude Model Family*. (2024). <https://www.anthropic.com/clause>.
- 1139 B. Atil et al.. 2024. *Non-Determinism of “Deterministic” LLM Settings*. arXiv preprint. (2024). arXiv: [2408.04667 \(cs.CL\)](#).
- 1140 M. Baker. 2016. “1,500 Scientists Lift the Lid on Reproducibility.” *Nature*, 533, 7604, 452–454. doi:[10.1038/533452a](#).
- 1141 A. Belz, S. Agarwal, A. Shimorina, and E. Reiter. 2021. “A Systematic Review of Reproducibility Research in Natural Lan-
 1142 guage Processing.” In: *Proceedings of the 16th Conference of the European Chapter of the Association for Computational
 Linguistics*. Association for Computational Linguistics, 381–393. doi:[10.18653/v1/2021.eacl-main.29](#).
- 1143 L. Biewald. 2020. *Experiment Tracking with Weights and Biases*. (2020). <https://wandb.com/>.
- 1144 R. Bommasani et al.. 2022. *On the Opportunities and Risks of Foundation Models*. arXiv preprint. (2022). arXiv: [2108.07258 \(cs.LG\)](#).
- 1145 T. Brown et al.. 2020. “Language Models are Few-Shot Learners.” In: *Advances in Neural Information Processing Systems*.
 1146 Vol. 33, 1877–1901. arXiv: [2005.14165 \(cs.CL\)](#).
- 1147 Y. Chen, J. Li, X. Liu, and Y. Li. 2023. *On the Reproducibility of ChatGPT in NLP Tasks*. arXiv preprint. (2023). arXiv: [2304.02554 \(cs.CL\)](#).
- 1148 J. Cohen. 1988. *Statistical Power Analysis for the Behavioral Sciences*. (2nd ed.). Lawrence Erlbaum Associates. ISBN:
 1149 978-0-8058-0283-2.
- 1150 J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. 2019. “BERT: Pre-training of Deep Bidirectional Transformers for
 1151 Language Understanding.” In: *Proceedings of the 2019 Conference of the North American Chapter of the Association
 1152 for Computational Linguistics*. Association for Computational Linguistics, 4171–4186. doi:[10.18653/v1/N19-1423](#).
- 1153 J. Dodge, S. Gururangan, D. Card, R. Schwartz, and N. A. Smith. 2019. “Show Your Work: Improved Reporting of
 1154 Experimental Results.” In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*.
 1155 Association for Computational Linguistics, 2185–2194. doi:[10.18653/v1/D19-1224](#).
- 1156 European Parliament and Council of the European Union. 2024. *Regulation (EU) 2024/1689 Laying Down Harmonised
 1157 Rules on Artificial Intelligence (AI Act)*. (2024). <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>.
- 1158 T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford. 2021. “Datasheets
 1159 for Datasets.” *Communications of the ACM*, 64, 12, 86–92. doi:[10.1145/3458723](#).
- 1160 Gemini Team, R. Anil, S. Borgeaud, J.-B. Alayrac, J. Yu, R. Soricut, J. Schalkwyk, A. M. Dai, A. Hauth, et al.. 2024.
 1161 *Gemini: A Family of Highly Capable Multimodal Models*. arXiv preprint. (2024). arXiv: [2312.11805 \(cs.CL\)](#).
- 1162 Gemma Team et al.. 2024. *Gemma 2: Improving Open Language Models at a Practical Size*. arXiv preprint. (2024). arXiv: [2408.00118 \(cs.CL\)](#).
- 1163 A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, et al.. 2024.
 1164 *The LLaMA 3 Herd of Models*. arXiv preprint. (2024). arXiv: [2407.21783 \(cs.AI\)](#).
- 1165 O. E. Gundersen, Y. Gil, and D. W. Aha. 2018. “On Reproducible AI: Towards Reproducible Research, Open Science, and
 1166 Digital Scholarship in AI Publications.” *AI Magazine*, 39, 3, 56–68. doi:[10.1609/aimag.v39i3.2816](#).
- 1167 O. E. Gundersen, M. Helmert, and H. H. Hoos. 2024. “Improving Reproducibility in AI Research: Four Mechanisms Adopted
 1168 by JAIR.” *Journal of Artificial Intelligence Research*, 81, 1019–1041. doi:[10.1613/jair.1.16905](#).
- 1169 O. E. Gundersen and S. Kjensmo. 2018. “State of the Art: Reproducibility in Artificial Intelligence.” *Proceedings of the
 1170 AAAI Conference on Artificial Intelligence*, 32, 1, 1644–1651. doi:[10.1609/aaai.v32i1.11503](#).
- 1171 M. Herschel, R. Diestelkämper, and H. Ben Lahmar. 2017. “A Survey on Provenance: What for? What form? What from?”
 1172 *The VLDB Journal*, 26, 6, 881–906. doi:[10.1007/s00778-017-0486-1](#).
- 1173 M. Hutson. 2018. “Artificial Intelligence Faces Reproducibility Crisis.” *Science*, 359, 6377, 725–726. doi:[10.1126/science.359.6377.725](#).
- 1174 A. Q. Jiang et al.. 2023. *Mistral 7B*. arXiv preprint. (2023). arXiv: [2310.06825 \(cs.CL\)](#).
- 1175 S. Kapoor and A. Narayanan. 2023. “Leakage and the Reproducibility Crisis in Machine-Learning-Based Science.” *Patterns*,
 4, 9, 100804. doi:[10.1016/j.patter.2023.100804](#).

- 1176 R. Kuprieiev, D. Petrov, and Iterative. 2024. *DVC: Data Version Control*. (2024). <https://dvc.org/>.
- 1177 LangChain. 2023. *LangSmith: A Platform for Building Production-Grade LLM Applications*. (2023). <https://smith.langchain.com/>.
- 1178 V. I. Levenshtein. 1966. “Binary Codes Capable of Correcting Deletions, Insertions, and Reversals.” *Soviet Physics Doklady*, 10, 8, 707–710.
- 1180 P. Liang et al.. 2023. “Holistic Evaluation of Language Models.” *Transactions on Machine Learning Research*. <https://openreview.net/forum?id=iO4LZibEqW>.
- 1182 C.-Y. Lin. 2004. “ROUGE: A Package for Automatic Evaluation of Summaries.” In: *Proceedings of the ACL-04 Workshop on Text Summarization Branches Out*. Association for Computational Linguistics, 74–81. <https://aclanthology.org/W04-1013/>.
- 1184 M. Mitchell, S. Wu, A. Zaldivar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, and T. Gebru. 2019. “Model Cards for Model Reporting.” In: *Proceedings of the Conference on Fairness, Accountability, and Transparency*. ACM, 220–229. doi:[10.1145/3287560.3287596](https://doi.org/10.1145/3287560.3287596).
- 1187 L. Moreau and P. Missier. 2013. *PROV-DM: The PROV Data Model*. W3C Recommendation. World Wide Web Consortium. <https://www.w3.org/TR/prov-dm/>.
- 1188 National Institute of Standards and Technology. 2023. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. Tech. rep. U.S. Department of Commerce. doi:[10.6028/NIST.AI.100-1](https://doi.org/10.6028/NIST.AI.100-1).
- 1190 Ollama. 2024. *Ollama: Run Large Language Models Locally*. (2024). <https://ollama.com/>.
- 1191 OpenAI. 2023. *OpenAI Eval: A Framework for Evaluating LLMs*. (2023). <https://github.com/openai/evals>.
- 1192 S. Ouyang, J. M. Zhang, M. Harman, and M. Wang. 2024. “An Empirical Study of the Non-determinism of ChatGPT in Code Generation.” *ACM Transactions on Software Engineering and Methodology*, 34, 2, 1–28. doi:[10.1145/3697010](https://doi.org/10.1145/3697010).
- 1194 G. Padovani, V. Anantharaj, and S. Fiore. 2025. “yProv4ML: Effortless Provenance Tracking for Machine Learning Systems.” *SoftwareX*, 29, 102028. doi:[10.1016/j.softx.2025.102028](https://doi.org/10.1016/j.softx.2025.102028).
- 1195 J. Pineau, P. Vincent-Lamarre, K. Sinha, V. Larivière, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and H. Larochelle. 2021. “Improving Reproducibility in Machine Learning Research: A Report from the NeurIPS 2019 Reproducibility Program.” *Journal of Machine Learning Research*, 22, 164, 1–20. <https://jmlr.org/papers/v22/20-303.html>.
- 1198 C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. 2020. “Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer.” *Journal of Machine Learning Research*, 21, 140, 1–67. <https://jmlr.org/papers/v21/20-074.html>.
- 1200 S. Samuel and B. König-Ries. 2022. “A Provenance-based Semantic Approach to Support Understandability, Reproducibility, and Reuse of Scientific Experiments.” *Journal of Biomedical Semantics*, 13, 1, 1–30. doi:[10.1186/s13326-022-00263-z](https://doi.org/10.1186/s13326-022-00263-z).
- 1202 V. Stodden, M. McNutt, D. H. Bailey, E. Deelman, Y. Gil, B. Hanson, M. A. Heroux, J. P. Ioannidis, and M. Taufer. 2016. “Enhancing Reproducibility for Computational Methods.” *Science*, 354, 6317, 1240–1241. doi:[10.1126/science.aah6168](https://doi.org/10.1126/science.aah6168).
- 1204 A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. 2017. “Attention is All You Need.” In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc. arXiv: [1706.03762 \(cs.CL\)](https://arxiv.org/abs/1706.03762).
- 1206 J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V. Le, and D. Zhou. 2022. “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.” In: *Advances in Neural Information Processing Systems*. Vol. 35, 24824–24837. arXiv: [2201.11903 \(cs.CL\)](https://arxiv.org/abs/2201.11903).
- 1209 M. D. Wilkinson et al.. 2016. “The FAIR Guiding Principles for Scientific Data Management and Stewardship.” *Scientific Data*, 3, 160018. doi:[10.1038/sdata.2016.18](https://doi.org/10.1038/sdata.2016.18).
- 1210 J. Yuan et al.. 2025. “Understanding and Mitigating Numerical Sources of Nondeterminism in LLM Inference.” In: *Advances in Neural Information Processing Systems*. Vol. 38. arXiv preprint. Curran Associates, Inc. arXiv: [2506.09501 \(cs.LG\)](https://arxiv.org/abs/2506.09501).
- 1212 M. Zaharia et al.. 2018. “Accelerating the Machine Learning Lifecycle with MLflow.” *IEEE Data Engineering Bulletin*, 41, 4, 39–45. <http://sites.computer.org/debull/A18dec/p39.pdf>.
- 1214 T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi. 2020. “BERTScore: Evaluating Text Generation with BERT.” In: *Proceedings of the 8th International Conference on Learning Representations*. <https://openreview.net/forum?id=SkeHuCVFDr>.
- 1216 Y. Zhu, P. Zhang, E. Haq, P. Hui, and G. Buchanan. 2023. *Can ChatGPT Reproduce Human-Generated Labels? A Study of Social Computing Tasks*. arXiv preprint. (2023). arXiv: [2304.10145 \(cs.CL\)](https://arxiv.org/abs/2304.10145).
- 1217
- 1218
- 1219
- 1220
- 1221
- 1222

1223 **A Reproducibility Checklist**

1224 The following checklist is designed for self-assessment of reproducibility in generative AI studies. Each
 1225 item maps to a specific field or artifact in our protocol.

1226

1227 **Prompt Documentation**

- 1228 (1) Is the exact prompt text recorded and versioned? [Prompt Card: prompt_text, prompt_hash]
- 1229 (2) Are design assumptions and limitations documented? [Prompt Card: assumptions, limitations]
- 1230 (3) Is the expected output format specified? [Prompt Card: expected_output_format]
- 1231 (4) Is the interaction regime documented (single/multi-turn)? [Prompt Card: interaction_regime]

1234 **Model and Environment**

- 1235 (5) Is the model name and version recorded? [Run Card: model_name, model_version]
- 1236 (6) Are model weights hashed for identity verification? [Run Card: weights_hash]
- 1237 (7) Is the execution environment fingerprinted? [Run Card: environment, environment_hash]
- 1238 (8) Is the source code version recorded? [Run Card: code_commit]

1240 **Execution and Output**

- 1241 (9) Are all inference parameters logged? [Run Card: inference_params]
- 1242 (10) Is the random seed recorded? [Run Card: inference_params.seed]
- 1243 (11) Is the output cryptographically hashed? [Run Card: output_hash]
- 1244 (12) Are execution timestamps recorded? [Run Card: timestamp_start, timestamp_end]
- 1245 (13) Is logging overhead measured separately? [Run Card: logging_overhead_ms]

1247 **Provenance**

- 1249 (14) Is a provenance graph generated per group? [PROV-JSON document]
- 1250 (15) Are provenance documents in an interoperable format? [W3C PROV standard]

1251

1252 **B Run Card Schema**

1253 The complete Run Card schema, with data types and descriptions:

1254

1255 Listing 1. Run Card JSON schema (simplified).

```

1256 1 {
1257 2   "run_id": "string (unique identifier)",
1258 3   "task_id": "string (task identifier)",
1259 4   "task_category": "string (e.g., summarization)",
1260 5   "prompt_hash": "string (SHA-256 of prompt)",
1261 6   "prompt_text": "string (full prompt text)",
1262 7   "input_text": "string (input to the model)",
1263 8   "input_hash": "string (SHA-256 of input)",
1264 9   "model_name": "string (e.g., llama3:8b)",
1265 10  "model_version": "string (e.g., 8.0B)",
1266 11  "weights_hash": "string (SHA-256 of weights)",
1267 12  "model_source": "string (e.g., ollama-local)",
1268 13  "inference_params": {
1269 14    "temperature": "float",

```

```

1270 15   "top_p": "float",
1271 16   "top_k": "integer",
1272 17   "max_tokens": "integer",
1273 18   "seed": "integer|null",
1274 19   "decoding_strategy": "string"
1275 20 },
1276 21   "params_hash": "string (SHA-256 of params)",
1277 22   "environment": {
1278 23     "os": "string",
1279 24     "os_version": "string",
1280 25     "architecture": "string",
1281 26     "python_version": "string",
1282 27     "hostname": "string",
1283 28     "timestamp": "ISO 8601 datetime"
1284 29 },
1285 30   "environment_hash": "string (SHA-256)",
1286 31   "code_commit": "string (git commit hash)",
1287 32   "researcher_id": "string",
1288 33   "affiliation": "string",
1289 34   "timestamp_start": "ISO 8601 datetime",
1290 35   "timestamp_end": "ISO 8601 datetime",
1291 36   "output_text": "string (model output)",
1292 37   "output_hash": "string (SHA-256 of output)",
1293 38   "output_metrics": "object (task-specific)",
1294 39   "execution_duration_ms": "float",
1295 40   "logging_overhead_ms": "float",
1296 41   "storage_kb": "float",
1297 42   "system_logs": "string (raw system info)",
1298 43   "errors": "array of strings",
1299 44
1300 45 // --- API-specific optional fields ---
1301 46   "api_request_id": "string|null (provider request ID)",
1302 47   "api_response_headers": "object|null (selected headers)",
1303 48   "api_model_version_returned": "string|null",
1304 49   "api_region": "string|null (if available)",
1305 50   "seed_status": "string (sent|logged-only|not-supported)",
1306 51
1307 52 // --- Multi-turn extension fields ---
1308 53   "conversation_history_hash": "string|null (SHA-256)",
1309 54   "turn_index": "integer|null",
1310 55   "parent_run_id": "string|null",
1311 56
1312 57 // --- RAG extension fields ---
1313 58   "retrieval_context": "string|null",
1314 59   "retrieval_context_hash": "string|null (SHA-256)"
1315 60 }
1316

```

C Example PROV-JSON Document

An abbreviated example of a PROV-JSON document generated for a single summarization run:

Listing 2. Abbreviated PROV-JSON for a summarization run.

```

1317
1318 1 {
1319 2   "prefix": {
1320 3     "genai": "https://genai-prov.org/ns#",
1321 4     "prov": "http://www.w3.org/ns/prov#"
1322 5   },
1323 6   "entity": {
1324 7     "genai:prompt_c9644358": {
1325 8       "prov:type": "genai:Prompt",
1326 9       "genai:hash": "c9644358805b...",
1327 10      "genai:task_category": "summarization"
1328 11    },
1329 12    "genai:model_llama3_8b": {
1330 13      "prov:type": "genai:ModelVersion",
1331 14      "genai:name": "llama3:8b",
1332 15      "genai:source": "ollama-local"
1333 16    },
1334 17    "genai:output_590d0835": {
1335 18      "prov:type": "genai:Output",
1336 19      "genai:hash": "590d08359e7d..."
1337 20    }
1338 21  },
1339 22   "activity": {
1340 23     "genai:run_llama3_8b_sum_001_C1_rep0": {
1341 24       "prov:type": "genai:RunGeneration",
1342 25       "prov:startTime": "2026-02-07T21:54:34Z",
1343 26       "prov:endTime": "2026-02-07T21:54:40Z"
1344 27     }
1345 28  },
1346 29   "wasGeneratedBy": {
1347 30     "_:wGB1": {
1348 31       "prov:entity": "genai:output_590d0835",
1349 32       "prov:activity": "genai:run_llama3_8b..."
1350 33     }
1351 34  },
1352 35   "used": {
1353 36     "_:u1": {
1354 37       "prov:activity": "genai:run_llama3_...",
1355 38       "prov:entity": "genai:prompt_c9644358"
1356 39     }
1357 40  },
1358 41   "agent": {
1359 42     "genai:researcher_lucas_rover": {
1360 43       "prov:type": "prov:Person",
1361 44       "genai:affiliation": "UTFPR"
1362 45     }
1363 46  },
1364 47   "wasAssociatedWith": {
1365 48     "_:wAW1": {

```

1364 Table 10. JSON extraction quality metrics by model and condition. *Raw Valid* = output parses directly as JSON;
 1365 *Extracted Valid* = JSON extracted via regex from outputs containing preamble text; *Schema* = all five expected
 1366 fields present; *Field EMR* = within-abstract pairwise exact match across runs for each extracted field, averaged over
 1367 abstracts (see Section D for interpretation). LLaMA 3 always prepends introductory text (e.g., “Here is the extracted
 1368 information in JSON format.”), yielding 0% raw validity but near-perfect extracted validity at $t=0$.
 1369

1370 Model	Cond.	Raw	Extr.	Schema	Within-Abstract Field EMR					Overall Field EMR
		Valid	Valid	Compl.	obj	meth	key_r	mod/sys	bench	
1372 LLaMA 3	C1 ($t=0$)	0%	100%	100%	0.987	0.987	0.987	1.000	0.987	0.989
	C2 ($t=0$)	0%	100%	100%	0.987	0.987	0.987	1.000	0.987	0.989
	C3 ($t=0.0$)	0%	100%	100%	0.978	0.978	0.978	1.000	0.978	0.982
	C3 ($t=0.3$)	0%	97.8%	97.8%	0.747	0.460	0.552	0.862	0.805	0.685
	C3 ($t=0.7$)	0%	92.2%	92.2%	0.522	0.167	0.267	0.611	0.711	0.456
1377 GPT-4	C2 ($t=0$)	100%	100%	100%	0.773	0.667	0.637	0.893	0.863	0.767
	C3 ($t=0.0$)	100%	100%	100%	0.833	0.571	0.667	0.905	0.810	0.757
	C3 ($t=0.3$)	100%	100%	100%	0.405	0.262	0.452	0.762	0.690	0.514
	C3 ($t=0.7$)	100%	100%	100%	0.137	0.157	0.255	0.667	0.725	0.388

```

1383 49   "prov:activity": "genai:run_llama3_...",
1384 50   "prov:agent": "genai:researcher_..."
1385 51 }
1386 52 }
1387 53 }
```

D JSON Extraction Quality

1391 Table 10 presents JSON-specific quality metrics for the structured extraction task. Two notable patterns
 1392 emerge.

1393 First, LLaMA 3 never produces raw-valid JSON: all 570 extraction outputs contain preamble text
 1394 (e.g., “Here is the extracted information in JSON format.”) before the JSON object, despite the prompt
 1395 explicitly requesting “JSON only, no explanation.” After extracting the embedded JSON via regex, va-
 1396 lidity rates reach 100% under greedy decoding, degrading slightly at higher temperatures (92.2% at
 1397 $t=0.7$). GPT-4, by contrast, always produces raw-valid JSON with 100% schema compliance across all
 1398 conditions. This instruction-following gap is consistent with the different prompt interfaces: the chat
 1399 completion API’s structured message format may better signal the expected output format.

1400 Second, within-abstract field-level exact match rates—computed by comparing only runs of the *same*
 1401 abstract under the same condition, then averaging across abstracts—confirm the overall reproducibility
 1402 hierarchy. Under greedy decoding, LLaMA 3 achieves near-perfect field EMR (0.982–0.989 overall), with
 1403 all five fields at or above 0.978, consistent with the overall extraction EMR of 0.987 reported in Table 4.
 1404 GPT-4 under greedy shows lower field EMR (0.757–0.767 overall), with open-ended fields (`method`: 0.667,
 1405 `key_result`: 0.637) lagging behind structured fields (`model_or_system`: 0.893, `benchmark`: 0.863). As
 1406 temperature increases, this gap widens: at $t=0.7$, `method` drops to 0.167 (LLaMA) and 0.157 (GPT-4),
 1407 while `benchmark` retains 0.711 and 0.725 respectively—a 4–5× difference. This within-abstract formu-
 1408 lation isolates true reproducibility (same input, same conditions, different runs) from between-abstract
 1409 content variation, providing a methodologically clean measure of field-level consistency.

1411 E Prompt Card Example

1412 The following is a complete, filled-in Prompt Card for the summarization task as used in our experiments:

1413
1414 Listing 3. Prompt Card for the scientific summarization task.

```

1415 1 {
1416 2   "prompt_id": "summarization_v1",
1417 3   "prompt_hash": "c9644358805b4a7e...",
1418 4   "version": "1.0.0",
1419 5   "task_category": "summarization",
1420 6   "objective": "Produce a 3-sentence summary of a
1421 7     scientific abstract covering: (1) main
1422 8       contribution, (2) methodology, (3) key result.",
1423 9   "assumptions": [
1424 10    "Input is a single English scientific abstract",
1425 11    "Abstract contains identifiable methodology
1426 12      and quantitative results",
1427 13    "Model can produce coherent 3-sentence output"
1428 14  ],
1429 15   "limitations": [
1430 16    "Open-ended phrasing allows high output variance",
1431 17    "No explicit output-format constraint (unlike
1432 18      extraction task)"
1433 19  ],
1434 20   "target_models": [
1435 21    "llama3:8b", "mistral:7b", "gemma2:9b",
1436 22    "gpt-4", "claude-sonnet-4-5"
1437 23  ],
1438 24   "expected_output_format": "Three sentences of
1439 25     plain text, no JSON or structured markup",
1440 26   "interaction_regime": "single-turn",
1441 27   "change_log": [
1442 28     {"date": "2026-02-06", "change": "Initial version"}
1443 29  ]
1444 30 }
```

1444 F Representative Prompt Templates

1445 The following are the exact prompt templates used for each of the four experimental tasks. In all tem-
1446 plates, {abstract} is replaced with the scientific abstract text at runtime.

1447 Task 1: Scientific Summarization

1448 Summarize the following scientific abstract in exactly 3 sentences. Cover: (1) the
1449 main contribution, (2) the methodology used, and (3) the key quantitative result.\n\nAbstract:
1450 {abstract}\n\nSummary:

1452 Task 2: Structured Extraction

1453 Extract the following fields from the scientific abstract below. Return JSON only, no
1454 explanation.\n\nFields: objective, method, key_result, model_or_system, benchmark\n\nAbstract:
1455 {abstract}\n\nJSON:

1456 Task 3: Multi-Turn Refinement (3 turns)

1457

1458 Turn 1: [Same as Task 1 prompt]\n Turn 2: Now revise the summary to be more specific
 1459 about the quantitative results mentioned.\n Turn 3: Finally, add one sentence about
 1460 the limitations or future work mentioned in the abstract.

1461 **Task 4: RAG Extraction**

1462
 1463 Using the context passage below and the scientific abstract, extract the following
 1464 fields. Return JSON only.\n\nContext: {retrieved_passage}\nAbstract: {abstract}\n\nFields:
 1465 objective, method, key_result, model_or_system, benchmark\n\nJSON:

1466

1467 **G Experimental Coverage Matrix**

1468 Table 11 provides a complete coverage matrix showing the number of abstracts and runs per model–task–
 1469 condition combination. This matrix enables readers to verify the sample sizes underlying all reported
 1470 metrics.

1471

1472

1473 Table 11. Experimental coverage: number of abstracts (runs) per model–task–condition. Dash (–) indicates the com-
 1474 bination was not evaluated. C1: fixed seed; C2: variable seed (C2_same_params for GPT-4); C3: temperature sweep
 1475 at $t \in \{0.0, 0.3, 0.7\}$.

1476

Model	Task	C1	C2	C3 ($t=0.0$)	C3 ($t=0.3$)	C3 ($t=0.7$)
LLaMA 3 8B	Extraction	30 (150)	30 (150)	30 (90)	30 (90)	30 (90)
	Summarization	30 (150)	30 (150)	30 (90)	30 (90)	30 (90)
	Multi-turn	10 (50)	–	–	–	–
	RAG	10 (50)	–	–	–	–
Mistral 7B	Extraction	10 (50)	10 (50)	10 (30)	10 (30)	10 (30)
	Summarization	10 (50)	10 (50)	10 (30)	10 (30)	10 (30)
	Multi-turn	10 (50)	–	–	–	–
	RAG	10 (50)	–	–	–	–
Gemma 2 9B	Extraction	10 (50)	10 (50)	10 (30)	10 (30)	10 (30)
	Summarization	10 (50)	10 (50)	10 (30)	10 (30)	10 (30)
	Multi-turn	10 (50)	–	–	–	–
	RAG	10 (50)	–	–	–	–
GPT-4	Extraction	–	30 (150)	17 (51)	17 (51)	14 (42)
	Summarization	3 (8) [†]	30 (150)	30 (90)	30 (90)	30 (90)
	Multi-turn	–	–	–	–	–
	RAG	–	–	–	–	–
Claude 4.5	Extraction	10 (49) [‡]	10 (50)	10 (30)	10 (30)	10 (30)
	Summarization	10 (50)	10 (50)	10 (30)	10 (30)	10 (30)
	Multi-turn	10 (50)	–	–	–	–
	RAG	10 (50)	–	–	–	–

1499 [†]GPT-4 C1 summarization: only 3 abstracts completed before quota exhaustion; excluded from primary analysis (C2
 1500 used instead).

1501 [‡]Claude C1 extraction: 49 runs (1 empty output due to API timeout).

1502

1503

1504

1505 H Chat-Format Control Experiment

1506 To assess whether the prompt-format difference between LLaMA 3 (completion-style via `/api/generate`)
 1507 and GPT-4 (chat-style via Chat Completions) contributes to the observed reproducibility gap, we con-
 1508 ducted a supplementary control experiment running LLaMA 3 8B through Ollama’s `/api/chat` endpoint,
 1509 which applies the model’s chat template (including special tokens for system/user/assistant roles) in the
 1510 same message structure used by GPT-4.

1511 **Design:** 10 abstracts \times 2 tasks \times 2 conditions (C1, C2) \times 5 repetitions = 200 runs, all under greedy
 1512 decoding ($t=0$).

1513 **Results:** Table 12 compares the chat-format control with the original completion-format results for
 1514 the same 10 abstracts. The two prompt formats produce *identical* variability metrics across all conditions:
 1515 summarization EMR = 0.929, NED = 0.0066, and ROUGE-L = 0.9922 in both modes; extraction achieves
 1516 perfect reproducibility (EMR = 1.000) regardless of interface. The 0.929 summarization EMR reflects the
 1517 warm-up effect on 2 of 10 abstracts—the same pattern observed in the full 30-abstract experiment. These
 1518 results confirm that prompt format is not a source of variability, and the reproducibility gap between
 1519 LLaMA 3 and GPT-4 is consistent with deployment-side factors (server infrastructure, floating-point
 1520 non-determinism across GPU types, request batching) rather than prompt-format differences.

1521

1522 Table 12. Prompt-format control: LLaMA 3 8B via completion (`/api/generate`) vs. chat (`/api/chat`) for 10 abstracts
 1523 under greedy decoding ($t=0$). EMR computed over conditions C1 and C2 combined.

1524

Task	Metric	Completion	Chat
Summarization	EMR↑	0.929	0.929
	NED↓	0.0066	0.0066
	ROUGE-L↑	0.9922	0.9922
Extraction	EMR↑	1.000	1.000
	NED↓	0.0000	0.0000
	ROUGE-L↑	1.0000	1.0000

1533 Note: Completion and chat formats yield identical metrics for all 10 abstracts under greedy decoding, indicating that
 1534 prompt format is not a source of variability.

1535

1536 Received February 2026

1537

1538

1539

1540

1541

1542

1543

1544

1545

1546

1547

1548

1549

1550

1551