

# Hidden Non-Determinism in Large Language Model APIs: A Lightweight Provenance Protocol for Reproducible Generative AI Research

LUCAS ROVER\*, UTFPR – Universidade Tecnológica Federal do Paraná, Brazil

YARA DE SOUZA TADANO, UTFPR – Universidade Tecnológica Federal do Paraná, Brazil

**Background:** Generative AI models produce non-deterministic outputs that vary across runs, even under nominally identical configurations. This variability threatens the reproducibility of studies that rely on large language model (LLM) outputs, yet most existing experiment-tracking tools were not designed for the specific challenges of text-generation workflows.

**Objectives:** We propose a lightweight, open-standard protocol for logging, versioning, and provenance tracking of generative AI experiments. The protocol introduces two novel documentation artifacts—Prompt Cards and Run Cards—and adopts the W3C PROV data model to create auditable, machine-readable provenance graphs linking every output to its full generation context.

**Methods:** We formalize the protocol and evaluate it empirically through 3,604 controlled experiments. These experiments employ five models—three locally deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B) and two API-served (GPT-4, Claude Sonnet 4.5)—on four NLP tasks (scientific summarization, structured extraction, multi-turn refinement, and retrieval-augmented generation) across 30 scientific abstracts and five experimental conditions that systematically vary the seed, temperature, and decoding strategy. We measure output variability using Exact Match Rate, Normalized Edit Distance, ROUGE-L, and BERTScore, and quantify the protocol’s own overhead in terms of time and storage.

**Results:** Under greedy decoding ( $t=0$ ), local models achieve near-perfect reproducibility: Gemma 2 9B reaches EMR = 1.000 across all tasks, LLaMA 3 attains EMR = 0.987 for extraction, and Mistral 7B achieves EMR = 0.960. By contrast, API-served models exhibit substantial hidden non-determinism: GPT-4 achieves only EMR = 0.443 for extraction, while Claude Sonnet 4.5 achieves EMR = 0.190 for extraction and EMR = 0.020 for summarization—the lowest observed in our study. This local-vs-API reproducibility gap (average single-turn EMR: 0.953 vs. 0.221, a more than 4-fold difference) is confirmed across two independent API providers. The gap extends to complex interaction regimes: under multi-turn refinement and RAG extraction, local models maintain high reproducibility (EMR  $\geq$  0.880), while Claude Sonnet 4.5—the only API model tested on these tasks—achieves EMR = 0.040 for multi-turn and EMR = 0.000 for RAG. The protocol adds less than 1% overhead across all five models.

**Conclusions:** Our results provide evidence that (1) API-served models exhibit substantial non-determinism under greedy decoding that is not attributable to user-controllable parameters, a pattern observed independently for both GPT-4 and Claude; (2) locally deployed models achieve near-perfect to perfect bitwise reproducibility under greedy decoding; (3) the local-vs-API gap extends to multi-turn refinement and RAG extraction,

\*Corresponding Author.

Authors’ Contact Information: Lucas Rover, ORCID: [0000-0001-6641-9224](#), lucasrover@utfpr.edu.br, UTFPR – Universidade Tecnológica Federal do Paraná, Programa de Pós-Graduação em Engenharia Mecânica, Ponta Grossa, Paraná, Brazil; Yara de Souza Tadano, ORCID: [0000-0002-3975-3419](#), yaratadano@utfpr.edu.br, UTFPR – Universidade Tecnológica Federal do Paraná, Programa de Pós-Graduação em Engenharia Mecânica, Ponta Grossa, Paraná, Brazil.



This work is licensed under a Creative Commons Attribution International 4.0 License.

© 2026 Copyright held by the owner/author(s).

DOI: [10.1613/jair.1.xxxxx](#)

1 where Claude Sonnet 4.5 (the only API model tested on these tasks) achieves near-zero EMR while local models maintain EMR  $\geq 0.880$ ; (4) temperature is the dominant user-controllable factor affecting variability; and (5) comprehensive provenance logging adds negligible overhead (<1%). The protocol, reference implementation, and all experimental data are publicly available.

5 CCS Concepts: • Software and its engineering → Software testing and debugging; Documentation; • Computing methodologies → Machine learning.

7 Additional Key Words and Phrases: reproducibility, large language models, non-determinism, provenance, generative AI, experiment tracking, W3C PROV, scientific methodology

9 **JAIR Associate Editor:** Insert JAIR AE Name

10 **JAIR Reference Format:**

12 Lucas Rover and Yara de Souza Tadano. 2026. Hidden Non-Determinism in Large Language Model APIs: A  
13 Lightweight Provenance Protocol for Reproducible Generative AI Research. *Journal of Artificial Intelligence  
14 Research* (2026), 31 pages. DOI: [10.1613/jair.1.xxxxx](https://doi.org/10.1613/jair.1.xxxxx)

## 15 1 Introduction

17 When a researcher queries a cloud-hosted LLM with the same prompt and temperature zero, one would  
18 reasonably expect identical outputs. Our experiments show otherwise: across five controlled seeds under  
19 greedy decoding, GPT-4 produces the same extraction result only 44% of the time, and Claude Sonnet  
20 4.5 achieves only 19%. Meanwhile, locally deployed models such as Gemma 2 9B produce *perfectly identical*  
21 outputs every time. This hidden, provider-dependent non-determinism exemplifies a fundamental  
22 challenge introduced by the rapid adoption of large language models (LLMs) in scientific research: how  
23 to ensure that studies relying on generative AI outputs are reproducible, auditable, and scientifically  
24 rigorous. Unlike traditional computational experiments, in which deterministic algorithms produce iden-  
25 tical results given identical inputs, LLMs exhibit inherent variability in their outputs due to stochastic  
26 sampling, floating-point non-determinism, and opaque model-versioning practices (Y. Chen et al. 2023;  
27 Zhu et al. 2023).

28 Importantly, “non-reproducible” does not necessarily mean “unreliable”: our experiments also show  
29 that semantic similarity (measured by BERTScore F1) remains above 0.94 across all conditions, even  
30 when exact textual match drops to zero. In other words, API outputs typically convey the same *meaning*  
31 despite differing in *phrasing*—but this distinction is invisible without systematic measurement, and many  
32 downstream analyses (meta-analyses, comparative studies, regulatory audits) require exact reproducibil-  
33 ity.

34 A related subtlety concerns the `seed` parameter offered by some APIs. For API-served models, the  
35 seed parameter is advisory, not a guarantee of determinism: OpenAI explicitly documents this caveat  
36 for GPT-4, and Anthropic’s Claude API does not support a seed parameter at all. Our experimental  
37 design accounts for this by treating seed variation as a control condition and measuring actual output  
38 reproducibility directly, rather than relying on API-side determinism guarantees.

39 This reproducibility challenge is not merely theoretical. Baker (2016) reported that over 70% of re-  
40 searchers have failed to reproduce another scientist’s experiment, a crisis that extends to AI research  
41 (Gundersen and Kjensmo 2018; Hutson 2018; Kapoor and A. Narayanan 2023; Stodden et al. 2016).  
42 For generative AI specifically, the problem is compounded by several factors unique to text-generation  
43 workflows: (1) the same prompt can yield semantically similar yet textually distinct outputs across runs;  
44 (2) API-based models may undergo silent updates that alter behavior; (3) temperature and sampling  
45 parameters create a high-dimensional space of possible outputs; and (4) no established standard exists  
46 for documenting the full context needed to understand, audit, or reproduce a generative output.

48 Existing experiment-tracking tools such as MLflow (Zaharia et al. 2018), Weights & Biases (Biewald  
 49 2020), and DVC (Kuprieiev et al. 2024) were designed primarily for training pipelines and numerical  
 50 metrics. Although valuable for their intended purposes, these tools lack features critical for generative  
 51 AI studies: structured prompt versioning, cryptographic output hashing for tamper detection, prove-  
 52 nance graphs linking outputs to their full generation context, and environment fingerprinting specific to  
 53 inference-time conditions.

54 In this paper, we make three contributions:

- 55 (1) **A lightweight protocol** for logging, versioning, and provenance tracking of generative AI exper-  
 56 iments. The protocol introduces *Prompt Cards* and *Run Cards* as structured documentation arti-  
 57 facts, and adopts the W3C PROV data model (Moreau and Missier 2013) for machine-readable  
 58 provenance graphs.
- 59 (2) **An empirical evaluation** of both the protocol’s effectiveness and the reproducibility charac-  
 60 teristics of LLM outputs. Through 3,604 controlled experiments with five models—three locally  
 61 deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B) and two API-served (GPT-4, Claude Son-  
 62 net 4.5)—across four tasks (extraction, summarization, multi-turn refinement, RAG extraction),  
 63 30 abstracts, and five conditions, we quantify output variability using four complementary met-  
 64 rrics and measure the protocol’s overhead. Our results document a striking, provider-independent  
 65 reproducibility gap between local and API-based inference that is invisible without systematic  
 66 logging.
- 67 (3) **A reference implementation** in Python that demonstrates the protocol’s practical applicabil-  
 68 ity, together with all experimental data, to facilitate adoption and independent verification.

69  
 70 The remainder of this paper is organized as follows. Section 2 reviews related work on reproducibility in  
 71 AI and experiment tracking. Section 3 formalizes the protocol design. Section 4 describes the experimental  
 72 methodology. Section 5 presents the empirical results. Section 6 discusses findings, limitations, and  
 73 practical implications. Section 7 concludes with directions for future work.

## 74 2 Related Work

### 75 2.1 Reproducibility in AI Research

76 The reproducibility crisis in AI has been documented extensively. Gundersen and Kjensmo (2018) sur-  
 77veyed 400 AI papers and found that only 6% provided sufficient information for full reproducibility.  
 78 Pineau et al. (2021) reported on the NeurIPS 2019 Reproducibility Program, which introduced re-  
 79 producibility checklists and found significant gaps between reported and actual reproducibility. More  
 80 recently, Gundersen, Helmert, et al. (2024) described four institutional mechanisms adopted by JAIR—  
 81 reproducibility checklists, structured abstracts, badges, and reproducibility reports—establishing a com-  
 82 munity standard for what should be documented in AI research. Gundersen, Gil, et al. (2018) identified  
 83 three levels of reproducibility in AI—method, data, and experiment—and argued that all three are nec-  
 84 essary for scientific progress. Belz et al. (2021) conducted a systematic review of 601 NLP papers and  
 85 confirmed pervasive under-reporting of experimental details, while Dodge et al. (2019) proposed improved  
 86 reporting standards for ML experiments, including confidence intervals and significance tests across mul-  
 87 tiple runs. More broadly, Kapoor and A. Narayanan (2023) identified data leakage as a widespread driver  
 88 of irreproducible results across 17 scientific fields that use ML-based methods.

89 For generative AI specifically, Y. Chen et al. (2023) demonstrated that ChatGPT’s outputs on NLP  
 90 benchmarks exhibit non-trivial variability across identical queries, even with temperature set to zero.  
 91 Zhu et al. (2023) showed that reproducibility degrades further when tasks involve subjective judgment,

95 such as social computing annotations. Most recently, Atil et al. (2024) systematically measured the non-  
 96 determinism of five LLMs under supposedly deterministic settings across eight tasks, finding accuracy  
 97 variations up to 15% across runs and introducing the Total Agreement Rate (TAR) metric. Ouyang et al.  
 98 (2024) confirmed that temperature zero does not guarantee determinism in ChatGPT code generation.  
 99 Most recently, Yuan et al. (2025) traced such non-determinism to numerical precision issues in GPU  
 100 kernels and proposed LayerCast as a mitigation strategy. Our work complements these studies in four  
 101 specific ways. First, whereas prior studies (including Atil et al.’s five-model, eight-task study) measure  
 102 variability post hoc, we provide a structured provenance protocol that enables *prospective* documentation  
 103 and audit—answering not only “how much variability?” but also “why did these outputs differ?” through  
 104 cryptographic hashing and W3C PROV graphs. Second, we directly compare local and API-based inference  
 105 on identical tasks with identical prompts across *five* models and *two* independent API providers  
 106 (OpenAI and Anthropic), isolating the deployment paradigm as a variable and suggesting that API non-  
 107 determinism is a consistent pattern across providers rather than a provider-specific artifact. Third, we  
 108 extend beyond single-turn evaluation to include multi-turn refinement and retrieval-augmented genera-  
 109 tion, demonstrating that reproducibility characteristics generalize across interaction regimes. Fourth, we  
 110 quantify the overhead of systematic logging, demonstrating that the “cost of knowing” is negligible.

111

112

## 113 2.2 Experiment Tracking Tools

114 Several tools exist for tracking machine learning experiments, although none was designed specifically  
 115 for generative AI text-output workflows:

116 **MLflow** (Zaharia et al. 2018) provides experiment tracking, model packaging, and deployment. It logs  
 117 parameters, metrics, and artifacts, but focuses on training pipelines and numerical outcomes rather than  
 118 text-generation provenance.

119 **Weights & Biases** (Biewald 2020) offers experiment tracking with visualization dashboards. It sup-  
 120 ports prompt logging but lacks structured prompt versioning, cryptographic output hashing, and prove-  
 121 nance graph generation.

122 **DVC** (Kuprieiev et al. 2024) provides data versioning through git-like operations. While effective for  
 123 dataset management, it does not address run-level provenance or prompt documentation.

124 **OpenAI Eval** (OpenAI 2023) is a framework for evaluating LLM outputs against benchmarks. It  
 125 provides structured evaluation but is tightly coupled to OpenAI’s ecosystem and does not generate  
 126 interoperable provenance records.

127 **LangSmith** (LangChain 2023) offers tracing and evaluation for LLM applications. It captures detailed  
 128 execution traces but uses a proprietary format and requires cloud connectivity.

129 More broadly, Bommasani et al. (2022) identified reproducibility as a key risk for foundation models,  
 130 and Liang et al. (2023) proposed the HELM benchmark for holistic evaluation of language models, in-  
 131 cluding robustness and fairness dimensions that complement our reproducibility focus. In the provenance  
 132 space, Padovani et al. (2025) recently introduced yProv4ML, a framework that captures ML provenance  
 133 in PROV-JSON format with minimal code modifications; our protocol shares the commitment to W3C  
 134 PROV but targets the specific challenges of stochastic text generation rather than training pipelines.

135 Table 1 provides a systematic feature-by-feature comparison of our protocol with these tools. The key  
 136 distinction is not merely one of tooling but of *scientific capability*: existing tools log what happened  
 137 during training (parameters, metrics, artifacts), whereas our protocol enables answering questions that  
 138 these tools cannot—specifically, whether two generative outputs are provably derived from identical  
 139 configurations, which exact factor caused a divergence between non-identical outputs, and whether an  
 140

141

142 Table 1. Comparison of our protocol with existing reproducibility tools and frameworks for GenAI experiments. Check-  
 143 marks (✓) indicate full support; tildes (~) indicate partial support; dashes (–) indicate no support.

144

145 Feature	Ours	MLflow	W&B	DVC	OpenAI Eval	LangSmith
Prompt versioning (Prompt Card)	✓	–	~	–	~	~
Run-level provenance (W3C PROV)	✓	–	–	–	–	–
Cryptographic output hashing	✓	–	–	✓	–	–
Seed & param logging	✓	✓	✓	–	✓	✓
Environment fingerprinting	✓	~	~	~	–	–
Model weights hashing	✓	–	~	✓	–	–
Overhead <1% of inference	✓	~	~	N/A	N/A	~
Designed for GenAI text output	✓	–	–	–	✓	✓
Open standard (PROV-JSON)	✓	–	–	–	–	–
Local-first (no cloud dependency)	✓	✓	–	✓	–	–

156

157

158 output has been tampered with post-generation. These capabilities require the combination of crypto-  
 159 graphic hashing, structured prompt documentation, and W3C PROV provenance graphs that no existing  
 160 tool provides. In short, our contribution is not an alternative experiment tracker but a *reproducibility*  
 161 *assessment framework* designed for the unique challenges of stochastic text generation.

162

### 163 2.3 Provenance in Scientific Computing

164 Data provenance—the lineage of data through transformations—has a rich history in database systems  
 165 and scientific workflows (Herschel et al. 2017). The W3C PROV family of specifications (Moreau and  
 166 Missier 2013) provides a standardized data model for representing provenance as directed acyclic graphs  
 167 of *entities*, *activities*, and *agents*. Samuel and König-Ries (2022) applied provenance tracking to compu-  
 168 tational biology workflows, demonstrating its value for reproducibility. However, to our knowledge, no  
 169 prior work has applied W3C PROV specifically to generative AI experiment workflows, in which the  
 170 challenge involves not only tracking data lineage but also capturing the stochastic generation context  
 171 that determines output variability.

172 Taken together, these gaps point to a clear need: a lightweight, standards-based protocol that bridges  
 173 generative AI inference with the provenance infrastructure already established in scientific computing.  
 174 The next section presents our design for such a protocol.

175

## 176 3 Protocol Design

177 Our protocol addresses the question: *What is the minimum set of metadata that must be captured for each*  
 178 *generative AI run to enable auditing, reproducibility assessment, and provenance tracking?* We address  
 179 this question through four complementary components.

180

### 181 3.1 Scope and Design Principles

182 The protocol is designed around three principles:

183

- 184 (1) **Completeness:** Every factor that can influence a generative output must be captured—prompt  
 185 text, model identity and version, inference parameters, environment state, and timestamps.
- 186 (2) **Negligible overhead:** The logging process must not materially affect the experiment. We target  
 187 <1% overhead relative to inference time.

188

- 189 (3) **Interoperability:** All artifacts are stored in open, machine-readable formats (JSON, PROV-  
 190 JSON), aligned with the FAIR (Findable, Accessible, Interoperable, Reusable) principles (Wilkin-  
 191 son et al. 2016), to enable tool integration and long-term preservation.

### 193 3.2 Prompt Cards

194 A *Prompt Card* is a versioned documentation artifact that captures the design rationale and metadata  
 195 for a prompt template used in experiments. Each Prompt Card contains:

- 196 • `prompt_id`: Unique identifier
- 197 • `prompt_hash`: SHA-256 hash of the prompt text, enabling tamper detection
- 198 • `version`: Semantic version number
- 199 • `task_category`: Classification of the task (e.g., summarization, extraction)
- 200 • `objective`: Natural-language description of what the prompt is designed to achieve
- 201 • `assumptions`: Explicit assumptions about inputs and expected behavior
- 202 • `limitations`: Known limitations or failure modes
- 203 • `target_models`: Models for which the prompt was designed and tested
- 204 • `expected_output_format`: Description of the expected output structure
- 205 • `interaction_regime`: Single-turn, multi-turn, or chain-of-thought
- 206 • `change_log`: History of modifications

207 Prompt Cards serve two purposes: they document design intent (supporting human understanding)  
 208 and they provide a citable, hashable reference for automated provenance tracking. The concept draws in-  
 209 spiration from Model Cards (Mitchell et al. 2019), Datasheets for Datasets (Gebru et al. 2021), and model  
 210 info sheets for reproducibility assessment (Kapoor and A. Narayanan 2023), extending the structured-  
 211 documentation paradigm to the prompt layer of the generative AI pipeline.

### 214 3.3 Run Cards

215 A *Run Card* captures the complete execution context of a single generative AI run. Each Run Card  
 216 records 24 core fields organized into five groups (the complete JSON schema in Appendix B includes  
 217 these fields plus additional metadata such as `researcher_id`, `affiliation`, `system_logs`, and `errors`):

- 218 (1) **Identification:** `run_id`, `task_id`, `task_category`, `prompt_hash`, `prompt_text`
- 219 (2) **Model context:** `model_name`, `model_version`, `weights_hash`, `model_source`
- 220 (3) **Parameters:** `inference_params` (temperature, top\_p, top\_k, max\_tokens, seed, decoding\_strategy),  
     `params_hash`
- 221 (4) **Input/Output:** `input_text`, `input_hash`, `output_text`, `output_hash`, `output_metrics`
- 222 (5) **Execution metadata:** `environment` (OS, architecture, Python version, hostname), `environment_hash`,  
     `code_commit`, timestamps (start/end), `execution_duration_ms`, `logging_overhead_ms`, `storage_kb`

226 For API-served models, optional extension fields capture provider-specific metadata that may help  
 227 diagnose non-determinism: `api_request_id`, `api_response_headers`, `api_model_version_returned`,  
 228 `api_region`, and a `seed_status` field that distinguishes between seeds that were “sent” to the API,  
 229 “logged-only” (recorded for protocol parity but not sent, as with Claude), or “not-supported” by the  
 230 provider. This formalization ensures that the advisory or absent nature of API seed parameters is cap-  
 231 tured as structured metadata rather than left as an undocumented assumption.

232 Figure 1 illustrates the Run Card schema as a minimal structured record.

233 The separation of logging overhead from execution time is deliberate: it allows researchers to verify  
 234 that the protocol itself does not confound experimental measurements.

<b>Run Card Schema (24 core + extension fields)</b>	
236	<b>1. Identification</b>
237	run_id · task_id · task_category · prompt_hash · prompt_text
238	<b>2. Model Context</b>
239	model_name · model_version · weights_hash · model_source
240	<b>3. Parameters</b>
241	inference_params {temp, top_p, top_k, max_tokens, seed, strategy} · params_hash
242	<b>4. Input/Output</b>
243	input_text · input_hash · output_text · output_hash · output_metrics
244	<b>5. Execution Metadata</b>
245	environment · environment_hash · code_commit · timestamps · duration_ms · overhead_ms · storage_kb
246	<b>API Extensions</b> (optional)
247	api_request_id · api_region · seed_status ∈ {sent, logged-only, not-supported}
248	<b>Workflow Extensions</b> (optional)
249	conversation_history_hash · turn_index · retrieval_context_hash · parent_run_id
250	
251	

252 Fig. 1. Run Card minimal schema. All SHA-256 hashes (5 total) enable tamper detection and automated comparison.  
 253 API and workflow extension fields are optional.

254

255

256

### 3.4 W3C PROV Integration

258 Each experimental group (defined by a unique model–task–condition–abstract combination) is automatically  
 259 translated into a W3C PROV-JSON document (Moreau and Missier 2013) that expresses the  
 260 generation provenance as a directed graph. The mapping defines:

261

- **Entities:** Prompt, InputText, ModelVersion, InferenceParameters, Output, ExecutionMetadata
- **Activities:** RunGeneration (the inference execution)
- **Agents:** Researcher, SystemExecutor (the execution environment)

264

PROV relations capture the causal structure:

266

- **used:** RunGeneration used Prompt, InputText, ModelVersion, InferenceParameters
- **wasGeneratedBy:** Output wasGeneratedBy RunGeneration
- **wasAssociatedWith:** RunGeneration wasAssociatedWith Researcher, SystemExecutor
- **wasAttributedTo:** Output wasAttributedTo Researcher
- **wasDerivedFrom:** Output wasDerivedFrom InputText

271

This standardized representation enables automated reasoning about experiment provenance, including detecting when two runs share identical configurations and identifying the specific factors that differ between non-identical outputs. The choice of W3C PROV over plain JSON logs is deliberate: PROV’s formal semantics allow automated tools to traverse the provenance graph and answer queries such as “what changed between these two runs?” without custom parsing logic. An abbreviated example document is given in Appendix C; to illustrate the structure concisely, the core provenance chain is:

278

```
Prompt →used RunGeneration →generated Output
InputText →used RunGeneration →assoc. Researcher
ModelVersion →used RunGeneration; Output →derived InputText
```

281

282

### 283 3.5 Reproducibility Checklist

284 We provide a 15-item checklist organized into four categories—Prompt Documentation, Model and En-  
 285 vironment, Execution and Output, and Provenance—that researchers can use to self-assess the repro-  
 286 ducibility of their generative AI studies. The complete checklist is provided in Appendix A.  
 287

### 288 3.6 Extensions for Advanced Workflows

289 The protocol’s field schema accommodates complex workflows through optional extension fields. Our em-  
 290 pirical evaluation exercises two of these extensions—multi-turn dialogues and RAG—while the remaining  
 291 extensions are specified in the reference implementation’s schema:  
 292

- 293 • **Multi-turn dialogues:** A `conversation_history_hash` field and `turn_index` enable linking  
   294 each turn to the full conversation state. *Evaluated in Task 3 (multi-turn refinement) using Ollama’s /api/chat endpoint.*
- 296 • **RAG:** Fields for retrieval context (with hashes) trace which external information influenced the  
   297 output. *Evaluated in Task 4 (RAG extraction) with prepended context passages.*
- 298 • **Tool use and function calling:** Fields for available tools, tool calls (with arguments, results,  
   299 and hashes) capture the full tool-use chain.
- 300 • **Chain-of-thought / agent workflows:** A `parent_run_id` field supports hierarchical prove-  
   301 nance graphs for multi-step reasoning chains.

302 Having defined the protocol’s components, we now evaluate it empirically along two dimensions: the  
 303 reproducibility characteristics it reveals across different models and conditions, and the overhead it  
 304 imposes on the experimental workflow.  
 305

## 306 4 Experimental Setup

307 We designed a controlled experiment to simultaneously evaluate (a) the reproducibility characteristics  
 308 of LLM outputs under varying conditions and (b) the overhead imposed by our logging protocol.  
 309

### 310 4.1 Models and Infrastructure

311 We evaluate five models representing two fundamentally different deployment paradigms: three locally  
 312 deployed open-weight models and two cloud API-served proprietary models. All local models were served  
 313 through Ollama v0.15.5 (Ollama 2024) on an Apple M4 system with 24 GB unified memory running  
 314 macOS 14.6 with Python 3.14.3.  
 315

316 **4.1.1 Local Models. LLaMA 3 8B** (Grattafiori et al. 2024): An open-weight model in Q4\_0 quanti-  
 317 zation. Local deployment provides complete control over the execution environment, eliminating con-  
 318 founding factors such as network latency, server-side batching, and silent model updates. The model’s  
 319 SHA-256 weights hash was recorded per run via the Ollama API.

320 **Mistral 7B** (Jiang et al. 2023): An open-weight model (Q4\_0 quantization) with a sliding-window  
 321 attention mechanism, providing a second data point for local inference reproducibility at a similar pa-  
 322 rameter scale.

323 **Gemma 2 9B** (Gemma Team et al. 2024): Google’s open-weight model (Q4\_0 quantization), rep-  
 324 resenting a third local model from an independent model family. Gemma 2 proved to be the most  
 325 deterministic model in our study.  
 326

327 **4.1.2 API-Served Models. GPT-4** (Achiam et al. 2023): Accessed via the OpenAI API (`openai` Python  
 328 SDK v1.59.9) with controlled seed parameters. The API returned `gpt-4-0613` as the resolved model  
 329

330 version in all runs. The API introduces additional sources of variability: load balancing, server-side  
 331 batching, potential model-version updates, and floating-point non-determinism across different hardware.

332 **Claude Sonnet 4.5 (Anthropic 2024)**: Accessed via the Anthropic API using a lightweight `urllib`-  
 333 based runner (no SDK dependency). Claude’s API does not support a `seed` parameter; we set `temperature=0`  
 334 for greedy decoding and logged a seed value for protocol parity (marked as `logged-only-not-sent-to-api`).  
 335 This provides an independent replication of the API non-determinism phenomenon on a second cloud  
 336 provider.

337

## 338 4.2 Tasks

339 We evaluate four tasks that span the output-structure spectrum and interaction complexity:

340 **Task 1: Scientific Summarization.** Given a scientific abstract, produce a concise summary in  
 341 exactly three sentences covering the main contribution, methodology, and key quantitative result. This  
 342 is an open-ended generation task in which the model has considerable freedom in word choice and  
 343 phrasing.

344 **Task 2: Structured Extraction.** Given a scientific abstract, extract five fields (objective, method,  
 345 key\_result, model\_or\_system, benchmark) into a JSON object. This is a constrained generation task  
 346 in which the output format is fixed and the model must select, rather than generate, content.

347 **Task 3: Multi-turn Refinement.** A three-turn dialogue in which the model first extracts structured  
 348 information, then receives feedback requesting more detail, and finally produces a refined extraction. This  
 349 tests reproducibility under conversational state accumulation, using Ollama’s `/api/chat` endpoint for  
 350 local models.

351 **Task 4: RAG Extraction.** The same structured extraction task as Task 2, but with an additional  
 352 retrieved context passage prepended to the input. This tests whether augmenting the prompt with  
 353 external context affects reproducibility.

354

## 355 4.3 Input Data

356 We use 30 widely-cited scientific abstracts from landmark AI/ML papers, including Vaswani et al. (2017)  
 357 (Transformer), Devlin et al. (2019) (BERT), Brown et al. (2020) (GPT-3), Raffel et al. (2020) (T5), Wei  
 358 et al. (2022) (Chain-of-Thought), as well as seminal works on GANs, ResNets, VAEs, LSTMs, CLIP,  
 359 DALL-E 2, Stable Diffusion, LLaMA, InstructGPT, PaLM, and others. These abstracts vary in length  
 360 (74–227 words), technical complexity, and the number of quantitative results reported, thereby providing  
 361 substantial diversity in the generation challenge.

362

## 363 4.4 Experimental Conditions

364 We define five conditions (Table 2) that systematically vary the factors hypothesized to affect repro-  
 365 ducibility:

366 **Design principle for API models.** For cloud-hosted APIs whose `seed` parameter is advisory rather  
 367 than deterministic (as documented by OpenAI for GPT-4) or entirely absent (as with Claude), the fixed-  
 368 vs.-variable seed distinction has no guaranteed effect server-side. We therefore treat C2 as the primary  
 369 test of determinism under greedy decoding for such models.

370 **C1 (Fixed seed, greedy decoding):** Temperature = 0, seed = 42 for all 5 repetitions. This represents  
 371 the maximum-control condition and should yield deterministic outputs.

372 **C2 (Variable seeds, greedy decoding):** Temperature = 0, seeds = {42, 123, 456, 789, 1024}. This  
 373 condition tests whether seed variation affects outputs when greedy decoding is used.

374

377 Table 2. Experimental design: conditions, parameters, and expected outcomes.  
378

Cond.	Description	Temp.	Seed	Reps	Expected Outcome
C1	Fixed seed, greedy	0.0	42 (fixed)	5	Deterministic output
C2	Variable seeds, greedy	0.0	5 different	5	Near-deterministic
C3 <sub>t=0.0</sub>	Temp. baseline	0.0	per-rep	3	Deterministic
C3 <sub>t=0.3</sub>	Low temperature	0.3	per-rep	3	Low variability
C3 <sub>t=0.7</sub>	High temperature	0.7	per-rep	3	High variability

385 Note: Tasks 1–2 are evaluated under all five conditions (C1, C2, C3). Tasks 3–4 (multi-turn, RAG) are evaluated under  
 386 C1 only for the three local models and Claude Sonnet 4.5. Total: 3,604 logged runs across 5 models. For API-served  
 387 models, C2 uses the same fixed seed as C1; the seed parameter is advisory and does not guarantee determinism.

388  
 389 **C3 (Temperature sweep):** Three sub-conditions at  $t \in \{0.0, 0.3, 0.7\}$  with 3 repetitions each, using  
 390 different seeds per repetition. This condition characterizes how temperature affects output variability.

391 **Run counts.** For Tasks 1–2 (extraction and summarization), each model is evaluated under C1 (5  
 392 runs), C2 (5 runs), and C3 (9 runs = 3 temperatures × 3 reps) per abstract. LLaMA 3 uses 30 abstracts  
 393 (1,140 runs); the newer models (Mistral 7B, Gemma 2 9B, Claude Sonnet 4.5) use 10 abstracts (380  
 394 runs each). For GPT-4, quota exhaustion limited collection to 724 runs (C2: 300/300; C3: 416/450; C1:  
 395 8/300 excluded). For Tasks 3–4 (multi-turn and RAG), the three local models and Claude Sonnet 4.5  
 396 are evaluated under C1 with 10 abstracts × 5 repetitions = 50 runs each (400 runs total). **Grand total:**  
 397 **3,604 valid runs.**

398 Table 3 summarizes the per-model run distribution.

400 Table 3. Run distribution across models and tasks.  
401

Model	Tasks 1–2	Tasks 3–4	Total
LLaMA 3 8B	1,140	100	1,240
Mistral 7B	380	100	480
Gemma 2 9B	380	100	480
GPT-4	724	—	724
Claude Sonnet 4.5	380	100	480
Chat-format control <sup>†</sup>	200	—	200
<b>Total</b>	<b>3,204</b>	<b>400</b>	<b>3,604<sup>¶</sup></b>

412  
 413 <sup>†</sup>LLaMA 3 8B via /api/chat endpoint (Appendix E).

## 4.5 Metrics

416 We adopt an operational definition of reproducibility at three levels, each mapped to a specific metric:

- 417 • **Exact reproducibility** (string-level): Two outputs are identical character-by-character. Measured by *Exact Match Rate (EMR)*.
- 418 • **Near reproducibility** (edit-level): Two outputs differ only in minor surface variations (punctuation, whitespace, synonym substitution). Measured by *Normalized Edit Distance (NED)*.

421  
 422 <sup>¶</sup>One Claude run (0.03%) returned an empty output due to API timeout and is excluded from variability metrics.

- 424 • **Semantic reproducibility** (meaning-level): Two outputs convey the same information despite  
 425 different phrasing. Measured by *ROUGE-L F1* and *BERTScore F1*.

426 This three-level framework allows us to distinguish between outputs that are bitwise identical (EMR = 1),  
 427 textually close ( $\text{NED} < 0.05$ ), and semantically equivalent ( $\text{ROUGE-L} > 0.90$ ). All variability metrics  
 428 are computed over all  $\binom{n}{2}$  unique output pairs within each experimental group (defined by model, task,  
 429 condition, and abstract):

430   **Exact Match Rate (EMR):** The fraction of output pairs that are character-for-character identical.  
 431 EMR = 1.0 indicates perfect reproducibility; EMR = 0.0 indicates that no two outputs match exactly.

432   **Normalized Edit Distance (NED):** The Levenshtein edit distance (Levenshtein 1966) between  
 433 each pair, normalized by the length of the longer string. NED = 0.0 indicates identical outputs; higher  
 434 values indicate greater textual divergence.

435   **ROUGE-L F1:** The F1 score based on the longest common subsequence at the word level (Lin 2004).  
 436 This captures semantic similarity even when surface forms differ. ROUGE-L = 1.0 indicates identical  
 437 word sequences.

438 Our primary metrics (EMR, NED, ROUGE-L) focus on exact and near reproducibility, which are  
 439 the most direct measures for our research question. To complement these surface-level metrics, we also  
 440 compute **BERTScore F1** (T. Zhang et al. 2020)—an embedding-based semantic similarity metric—  
 441 for all conditions. BERTScore captures meaning-level equivalence that surface metrics may miss (e.g.,  
 442 paraphrases), providing a fourth perspective on reproducibility. For the structured extraction task, we  
 443 additionally report **JSON validity rate**, **schema compliance rate**, and **field-level accuracy**, which  
 444 measure whether outputs are syntactically valid JSON, contain all expected fields, and agree on individual  
 445 field values across runs, respectively (see Appendix D for detailed results).

446 For protocol overhead, we measure:

- 447 • **Logging time:** Wall-clock time spent on hashing, metadata collection, and file I/O, measured  
 448 separately from inference time.  
 449 • **Storage:** Size of each run record (JSON) and total storage for all protocol artifacts.  
 450 • **Overhead ratio:** Logging time as a percentage of total execution time.

## 453 5 Results

### 454 5.1 Reproducibility Under Greedy Decoding

455 Table 4 presents the headline result: Exact Match Rates under greedy decoding for all five models. Table 5  
 456 provides the full three-level reproducibility assessment.

457   **5.1.1 Local Models: Near-Perfect to Perfect Reproducibility.** **Finding 1: Gemma 2 9B achieves perfect**  
 458 **bitwise reproducibility under greedy decoding.** Across all tasks and conditions with  $t=0$ , Gemma 2  
 459 9B produces EMR = 1.000 with NED = 0.000—every single output is character-for-character identical  
 460 across repetitions. This includes not only single-turn extraction and summarization but also multi-turn  
 461 refinement and RAG extraction.

462   **Finding 2: All three local models achieve high reproducibility.** LLaMA 3 8B attains EMR =  
 463 0.987 for extraction and 0.931 for summarization; Mistral 7B achieves 0.960 and 0.840, respectively. The  
 464 small deviations from perfect reproducibility in LLaMA 3 and Mistral 7B are attributable to a warm-up  
 465 effect on the first inference call after model loading, which affects 2–4 of the 10–30 abstracts per model.  
 466 Seed variation (C1 vs. C2) has *no effect* under greedy decoding for any local model: the model always  
 467 selects the highest-probability token, making the seed irrelevant.

468

471 Table 4. Exact Match Rate (EMR) under greedy decoding ( $t=0$ ) across five models and two single-turn tasks. Values  
 472 for local models aggregate conditions C1 and C2 (both greedy,  $t=0$ ); for GPT-4, values reflect C2 only (C1 has  
 473 insufficient coverage: 3/30 abstracts for summarization); for Claude, values reflect C1 only (Claude’s API does not  
 474 support the seed parameter; C1 and C2 differ slightly due to sampling variability, with C2 extraction EMR = 0.080 vs.  
 475 C1 = 0.190). Higher is more reproducible.

Model	Source	Extraction	Summarization	N Runs	N Abstracts
Gemma 2 9B	Local	1.000	1.000	100	10
Mistral 7B	Local	0.960	0.840	100	10
LLaMA 3 8B	Local	0.987	0.931	400	30
GPT-4	API	0.443	0.230	300	30
Claude Sonnet 4.5	API	0.190	0.020	99	10

484 Table 5. Three-level reproducibility assessment under greedy decoding ( $t=0$ ). L1: bitwise identity (EMR), L2: surface  
 485 similarity (NED, ROUGE-L), L3: semantic equivalence (BERTScore F1). Values are means across abstracts.

Model	Task	L1: Bitwise		L2: Surface		L3: Semantic	
		EMR	$\sigma$	NED $\downarrow$	ROUGE-L $\uparrow$	BERTScore F1 $\uparrow$	
Gemma 2 9B	Extraction	1.000	0.000	0.000	1.000	1.0000	
	Summarization	1.000	0.000	0.000	1.000	1.0000	
Mistral 7B	Extraction	0.960	0.120	0.001	1.000	0.9999	
	Summarization	0.840	0.196	0.046	0.955	0.9935	
LLaMA 3 8B	Extraction	0.987	0.072	0.003	0.997	0.9997	
	Summarization	0.931	0.157	0.014	0.986	0.9979	
GPT-4	Extraction	0.443	0.335	0.072	0.938	0.9904	
	Summarization	0.230	0.193	0.137	0.870	0.9839	
Claude Sonnet 4.5	Extraction	0.190	0.291	0.101	0.904	0.9878	
	Summarization	0.020	0.040	0.242	0.764	0.9704	

503 **5.1.2 API-Served Models: Substantial Hidden Non-Determinism.** **Finding 3:** Both API-served models  
 504 exhibit substantial non-determinism under greedy decoding, observed independently across  
 505 two providers. Under  $t=0$  with controlled seeds, GPT-4 achieves EMR = 0.443 for extraction and 0.230  
 506 for summarization. Claude Sonnet 4.5 is even less deterministic: EMR = 0.190 for extraction and EMR =  
 507 0.020 for summarization—meaning that across 10 abstracts  $\times$  5 repetitions, Claude produced the same  
 508 summarization output only 2% of the time.

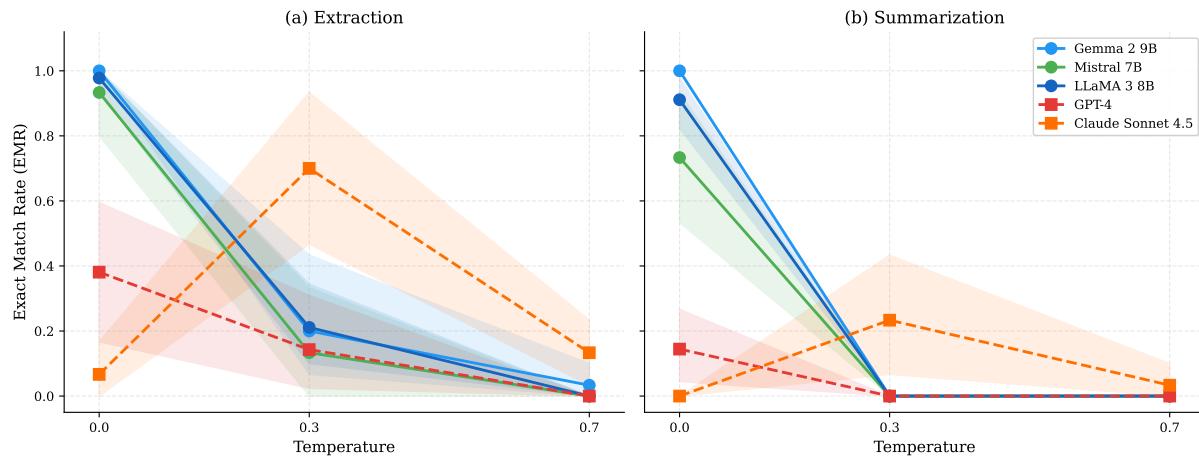
509 Table 6 summarizes the deployment-paradigm gap.

510 Under the representative greedy condition for each model (C1 for local models and Claude, C2 for GPT-  
 511 4; see Table 4), the average single-turn EMR is **0.953 for local models vs. 0.221 for API models**—a  
 512 more than 4-fold reproducibility gap. This gap is not due to user-side parameter differences: all models  
 513 use  $t=0$  with the same decoding strategy. The observed variability is consistent with deployment-side  
 514 factors invisible to the researcher—such as hardware-level floating-point variability, request batching, and  
 515 model routing. This pattern, observed independently across two API providers (OpenAI and Anthropic),  
 516 is consistent with non-determinism arising from factors common to cloud-hosted LLM inference rather  
 517

518 Table 6. API-served vs. locally deployed models under greedy decoding (single-turn tasks only). Local averages: simple  
 519 mean across 3 models  $\times$  2 tasks (C1+C2 combined). API averages: simple mean across 2 models  $\times$  2 tasks (GPT-  
 520 4 C2, Claude C1). Local models exhibit substantially higher bitwise reproducibility, consistent with deployment-side  
 521 factors—rather than user-controllable parameters—as a major contributor to API output variability.  
 522

Deployment	EMR↑	NED↓	ROUGE-L↑	BS-F1↑
Local (3 models)	0.953	0.011	0.990	0.9985
API (2 models)	0.221	0.138	0.869	0.9831

527 Effect of Sampling Temperature on Reproducibility



544 Fig. 2. Effect of temperature on Exact Match Rate across five models. (a) Extraction task. (b) Summarization task.  
 545 Local models (solid lines) start from near-perfect or perfect reproducibility at  $t=0$ , while API models (dashed lines)  
 546 start from a much lower baseline. All models converge toward  $EMR = 0$  at  $t=0.7$ .

548 than being a provider-specific artifact. *Without systematic logging, this non-determinism would be entirely  
 549 invisible.*

551 **5.1.3 Temperature Effects Across Models. Finding 4: Temperature is the dominant user-controllable  
 552 factor affecting variability.** Figure 2 shows the relationship between temperature and EMR for all  
 553 five models. Table 7 provides the full temperature sweep data.

554 Within the C3 temperature sweep, increasing temperature from 0.0 to 0.7 reduces EMR to zero  
 555 for all models on summarization. For extraction, local models drop from  $EMR > 0.93$  to near zero,  
 556 while API models drop from their already-low baselines. Notably, BERTScore F1 remains above 0.94  
 557 across all conditions even when EMR drops to zero, indicating that non-determinism is primarily a  
 558 *phrasing* phenomenon rather than a *meaning* phenomenon: even when outputs differ textually, they  
 559 convey equivalent information. This distinction is practically important—researchers whose downstream  
 560 analyses depend on semantic content rather than exact wording may find API outputs acceptable despite  
 561 low EMR.

562 However, the temperature–reproducibility relationship is not uniformly monotonic across all models.  
 563 Claude Sonnet 4.5 exhibits an anomalous pattern under the C3 sweep: extraction EMR *increases* from  
 564

565 Table 7. Effect of sampling temperature on Exact Match Rate (EMR) under condition C3. For local models, increasing  
 566 temperature monotonically reduces EMR. For API models, the relationship is more complex: Claude Sonnet 4.5 exhibits  
 567 higher EMR at  $t=0.3$  than at  $t=0.0$  (see text). At  $t=0.7$ , all models converge toward  $\text{EMR} \approx 0$  for summarization.

Model	Task	$t=0.0$	$t=0.3$	$t=0.7$
Gemma 2 9B	Extraction	1.000	0.200	0.033
	Summarization	1.000	0.000	0.000
Mistral 7B	Extraction	0.933	0.133	0.000
	Summarization	0.733	0.000	0.000
LLaMA 3 8B	Extraction	0.978	0.211	0.000
	Summarization	0.911	0.000	0.000
GPT-4	Extraction	0.381	0.143	0.000
	Summarization	0.144	0.000	0.000
Claude Sonnet 4.5	Extraction	0.067	0.700	0.133
	Summarization	0.000	0.233	0.033

583 0.067 at  $t=0.0$  to 0.700 at  $t=0.3$  before declining to 0.133 at  $t=0.7$ ; summarization shows a similar  
 584 inversion ( $\text{EMR} = 0.000$  at  $t=0.0$ , rising to 0.233 at  $t=0.3$ ). This counterintuitive behavior—where a small  
 585 positive temperature *improves* reproducibility relative to greedy decoding—may reflect how Anthropic’s  
 586 infrastructure implements the  $t=0$  decoding path: at exactly zero temperature, server-side stochastic  
 587 processes (e.g., speculative decoding, hardware-level floating-point non-determinism across GPU types,  
 588 or request batching effects) may dominate output variability, whereas a small positive temperature  
 589 may activate a more stable sampling path that happens to converge on similar tokens. With  $n=10$   
 590 abstracts and 30 runs per temperature level (standard deviation  $\sigma = 0.38$  for the 0.700 extraction EMR),  
 591 this observation should be interpreted cautiously. Nevertheless, it underscores that the temperature–  
 592 reproducibility relationship for API-served models depends on provider-specific implementation details  
 593 that are opaque to researchers. Finding 4 therefore holds robustly for local models and for the overall  
 594  $t=0$  to  $t=0.7$  trajectory, but the precise shape of the temperature–response curve for individual API  
 595 providers merits further investigation with larger sample sizes.

## 5.2 Multi-Turn and RAG Reproducibility

599 **Finding 5: The local-vs-API reproducibility gap extends to complex interaction regimes.**  
 600 Table 8 and Figure 3 present results for multi-turn refinement and RAG extraction across the three local  
 601 models and Claude Sonnet 4.5.

602 Gemma 2 9B and Mistral 7B achieve perfect  $\text{EMR} = 1.000$  for both multi-turn refinement and RAG  
 603 extraction, demonstrating that conversational state accumulation and context augmentation do not de-  
 604 grade reproducibility when the underlying model is deterministic. LLaMA 3 8B shows  $\text{EMR} = 0.880$   
 605 for multi-turn and 0.960 for RAG—slightly lower than its single-turn extraction performance (0.987),  
 606 consistent with error accumulation across dialogue turns.

607 Claude Sonnet 4.5, the only API-served model evaluated on these tasks, achieves  $\text{EMR} = 0.040$  for  
 608 multi-turn refinement and  $\text{EMR} = 0.000$  for RAG extraction—the lowest values observed in our study.  
 609 The RAG result is particularly striking: across 50 runs (10 abstracts  $\times$  5 repetitions), not a single pair of  
 610 outputs was character-for-character identical ( $\text{NED} = 0.256$ ). This confirms that API non-determinism is

612 Table 8. Reproducibility under complex interaction regimes (C1 fixed seed,  $t=0$ ). Multi-turn refinement involves three  
 613 successive prompt-response exchanges. RAG extraction augments the prompt with a retrieved context passage. Claude  
 614 Sonnet 4.5 is included as a representative API-served model; its near-zero EMR across all four scenarios confirms that  
 615 the local-vs-API reproducibility gap extends to complex interaction regimes.

616	Model	Scenario	EMR	NED↓	ROUGE-L↑	BS-F1↑
617	Gemma 2 9B	Single-turn Extraction	1.000	0.000	1.000	1.0000
618		Single-turn Summarization	1.000	0.000	1.000	1.0000
619		Multi-turn Refinement	1.000	0.000	1.000	1.0000
620		RAG Extraction	1.000	0.000	1.000	1.0000
621	Mistral 7B	Single-turn Extraction	0.960	0.001	1.000	0.9999
622		Single-turn Summarization	0.840	0.046	0.955	0.9935
623		Multi-turn Refinement	1.000	0.000	1.000	1.0000
624		RAG Extraction	1.000	0.000	1.000	1.0000
625	LLaMA 3 8B	Single-turn Extraction	0.987	0.003	0.997	0.9997
626		Single-turn Summarization	0.931	0.014	0.986	0.9979
627		Multi-turn Refinement	0.880	0.012	0.988	0.9986
628		RAG Extraction	0.960	0.012	0.985	0.9987
629	Claude Sonnet 4.5	Single-turn Extraction	0.190	0.101	0.904	0.9878
630		Single-turn Summarization	0.020	0.242	0.764	0.9704
631		Multi-turn Refinement	0.040	0.189	0.834	0.9780
632		RAG Extraction	0.000	0.256	0.748	0.9714

634

635

636

637

not limited to single-turn tasks but persists—and may even worsen—under complex interaction regimes where longer outputs and additional context amplify server-side variability.

638

639

### 5.3 Cross-Model Comparison

Figure 4 provides a comprehensive heatmap of EMR across all model-task combinations, and Figure 5 shows the three-level reproducibility profile for each model.

640

641

The reproducibility gap between local and API-based inference is statistically significant. Using paired  $t$ -tests on per-abstract EMR values under greedy decoding across the 30 LLaMA 3/GPT-4 abstracts: for summarization,  $t(29) = 17.250$ ,  $p < 0.0001$ , Cohen's  $d = 3.149$ ; for extraction,  $t(29) = 8.996$ ,  $p < 0.0001$ , Cohen's  $d = 1.642$ . Both effect sizes are very large ( $d > 1.6$ ), and all  $p$ -values survive Bonferroni correction. Non-parametric Wilcoxon signed-rank tests confirm all results ( $p < 0.001$ ).

642

643

644

645

646

647

648

649

650

651

### 5.4 Protocol Overhead

Table 9 presents the protocol's overhead metrics across all five models.

652

653

654

655

656

657

The protocol adds less than 1% overhead for all five models, with mean logging time ranging from 21–30 ms depending on the model and task. Storage overhead remains modest at approximately 4 KB per run record. The overhead is consistent across local and API deployment modes, indicating that the protocol is deployment-agnostic.

658

Figure 6 provides an additional perspective on surface-level variability across models.

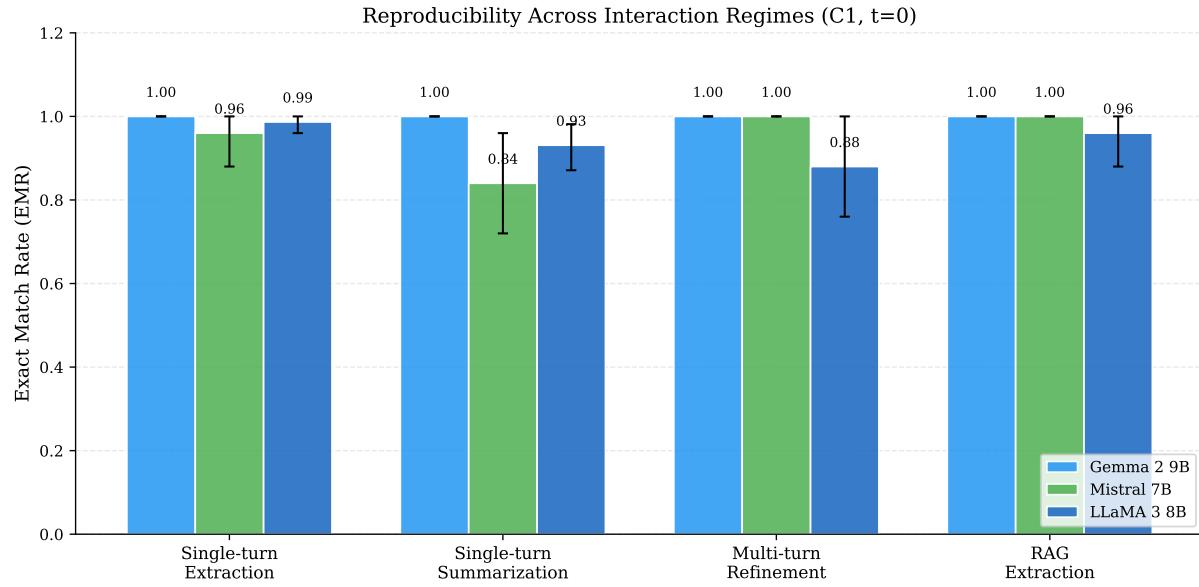


Fig. 3. Reproducibility across interaction regimes ( $C_1, t=0$ ) for four models. Local models maintain high EMR across all scenarios, while Claude Sonnet 4.5 (API) shows near-zero EMR throughout, confirming the reproducibility gap extends to multi-turn and RAG tasks.

Table 9. Provenance logging overhead across five models under greedy decoding ( $C_1$ ). The protocol adds negligible overhead (<1%) to inference latency across all models and deployment modes.

Model	Source	Mean Inference (ms)	Mean Overhead (ms)	Overhead (%)
Gemma 2 9B	Local	181,579.3	30.6	0.234
Mistral 7B	Local	13,931.3	27.3	0.281
LLaMA 3 8B	Local	7,524.8	26.7	0.456
GPT-4	API	4,519.7	24.5	0.564
Claude Sonnet 4.5	API	4,359.3	26.5	0.727

## 6 Discussion

The preceding results paint a clear and consistent picture: locally deployed models under greedy decoding achieve near-perfect to perfect bitwise reproducibility across all four tasks, while API-served models—from two independent providers—exhibit substantial hidden variability on single-turn tasks that researchers cannot control. Temperature is the dominant user-controllable factor for local models (though API models show a more complex temperature–reproducibility relationship; see Section 5), structured tasks are more reproducible than open-ended ones, and complex interaction regimes (multi-turn, RAG) do not degrade local-model reproducibility. We now consider what these findings mean for research practice, what the protocol enables that was previously invisible, and where the current study’s limitations lie.

## Bitwise Reproducibility Under Greedy Decoding

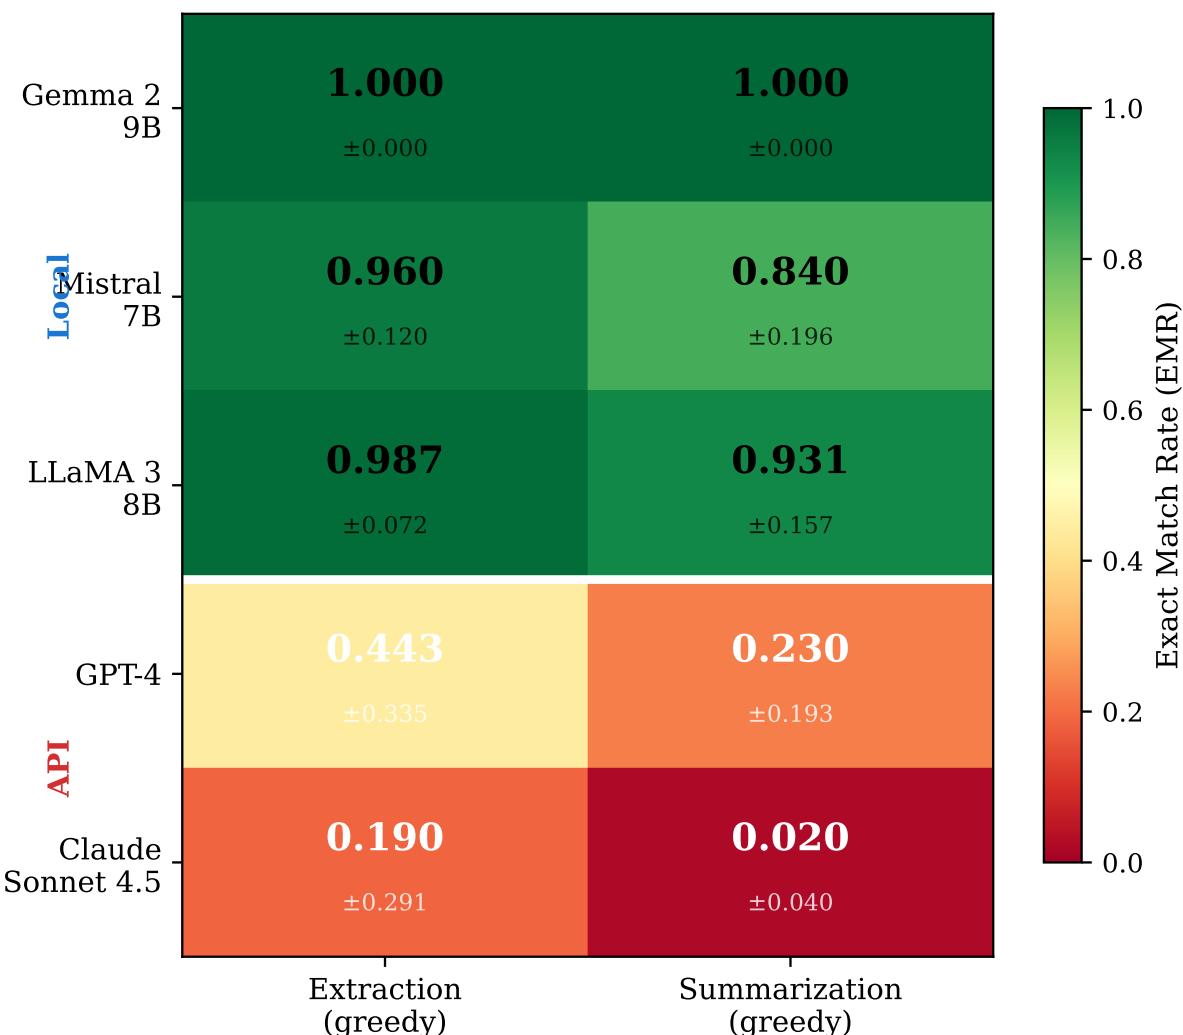
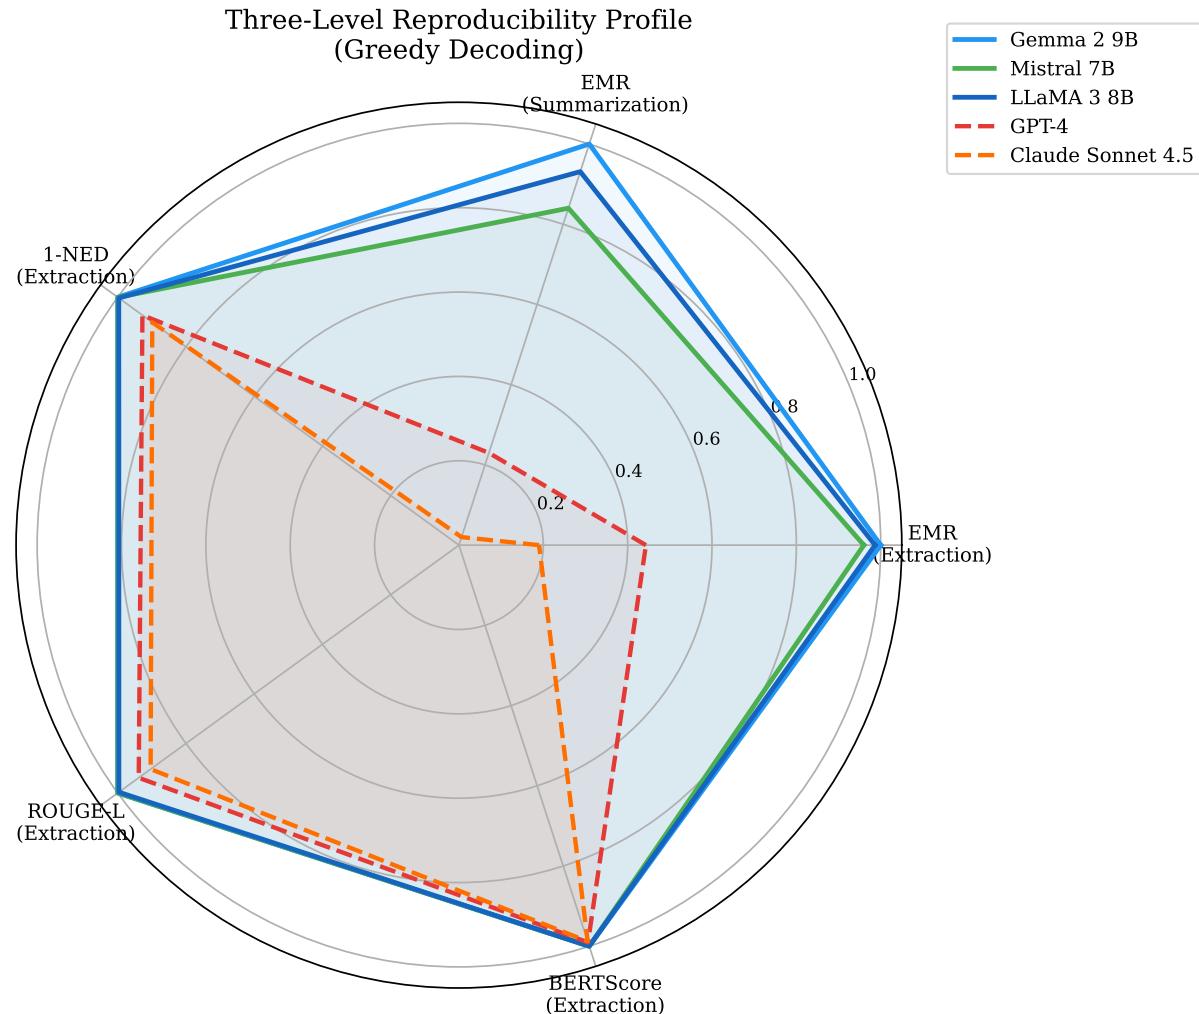


Fig. 4. Heatmap of Exact Match Rate under greedy decoding for five models. The horizontal white line separates local models (top three, green) from API-served models (bottom two, red). Gemma 2 9B achieves perfect 1.000 across all tasks.

### 6.1 Implications for Reproducibility Practice

Our results yield several actionable recommendations for researchers conducting generative AI experiments:

**Use greedy decoding with local models for maximum reproducibility.** Gemma 2 9B achieved perfect EMR = 1.000 across all tasks under greedy decoding. LLaMA 3 and Mistral 7B achieved EMR



$\geq 0.840$ . Local deployment with  $t=0$  should be the default configuration for any study in which output consistency is critical.

**API non-determinism is observed across providers.** Our most consequential finding is that both GPT-4 (OpenAI) and Claude Sonnet 4.5 (Anthropic) exhibit substantial non-determinism under greedy decoding on single-turn tasks. Claude's EMR of 0.020 for summarization means that effectively no two runs produce the same output. Researchers using *any* API-served model should never assume reproducibility without verification and should report multiple runs with variability metrics.

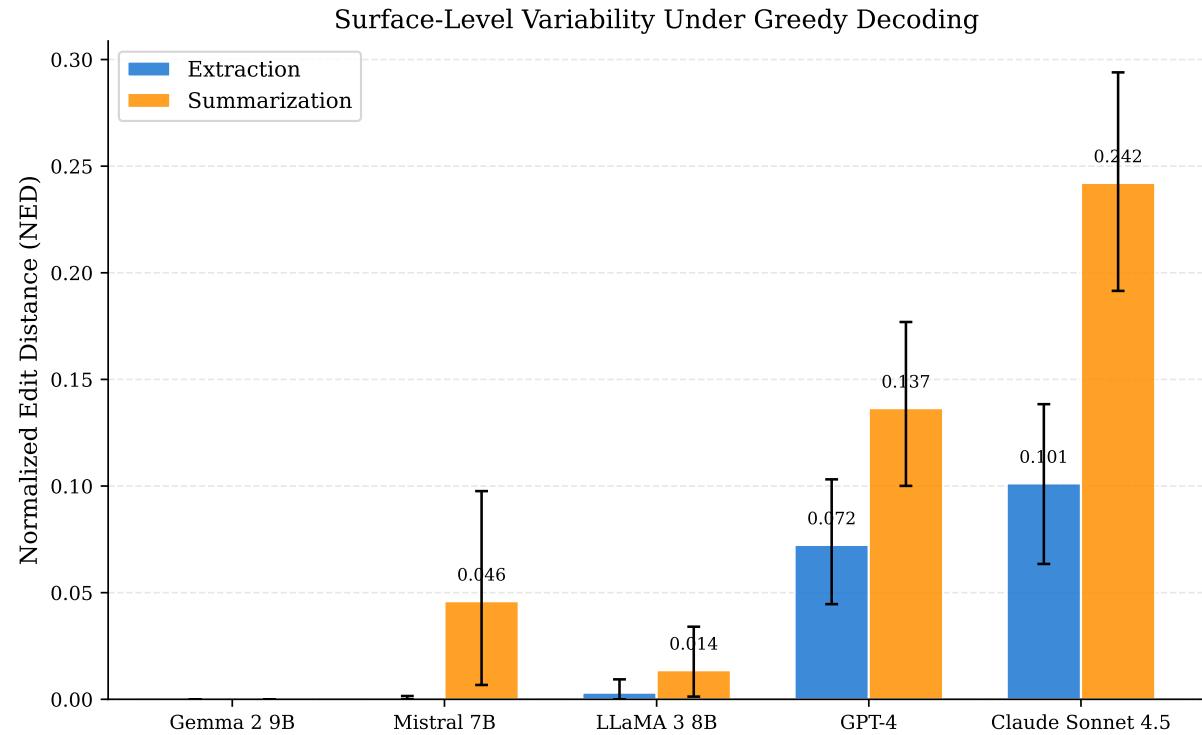


Fig. 6. Normalized Edit Distance (NED) under greedy decoding. Local models show near-zero NED (Gemma 2: 0.000, Mistral: 0.001), while API models exhibit NED 0.07–0.30, quantifying the surface-level divergence that accompanies the EMR gap.

**Prefer structured output formats when possible.** The extraction task's consistently higher reproducibility across all five models demonstrates that output-format constraints directly improve reproducibility. This effect holds for both local models (EMR 0.960–1.000 for extraction vs. 0.840–1.000 for summarization) and API models (EMR 0.190–0.443 for extraction vs. 0.020–0.230 for summarization).

**Include warm-up runs for local models.** The per-abstract analysis revealed that the first inference call after model loading may differ from subsequent calls due to cache initialization. This affects LLaMA 3 and Mistral 7B on 2–4 of their abstracts, slightly reducing aggregate EMR.

**Log comprehensively; the cost is negligible.** At less than 1% overhead and approximately 4 KB per run across all five models, there is no practical reason not to apply comprehensive logging. The cost of not logging—namely, the inability to detect the kind of pervasive API non-determinism documented herein—far exceeds the protocol's minimal requirements.

## 6.2 Local vs. API Inference: A Persistent Reproducibility Gap

The most significant finding of this study is the reproducibility gap between local and API-based inference, observed consistently across two independent cloud providers. Under greedy decoding on single-turn tasks, local models average EMR = 0.953 while API models average EMR = 0.221—a more than 4-fold gap. The fact that Claude Sonnet 4.5 (Anthropic) exhibits *even lower* reproducibility than GPT-4

(OpenAI) is inconsistent with provider-specific implementation as the sole explanation and suggests that non-determinism arises from factors common to distributed cloud inference infrastructure, such as hardware-level floating-point variability, request batching, and model routing.

This gap has profound implications for the scientific use of API-based LLMs. *Without systematic logging, a researcher using GPT-4 or Claude would have no way of knowing that their “deterministic” experiment produces different outputs across runs.* Our protocol makes this hidden non-determinism visible, measurable, and documentable.

### 6.3 Task-Dependent Reproducibility

The difference between summarization and extraction reproducibility—observed consistently across all five models—is consistent with and extends our earlier two-model finding. The reproducibility hierarchy (extraction > summarization) holds for local models (EMR gap of 0.03–0.12) and is amplified for API models (EMR gap of 0.17–0.25). This finding suggests a spectrum ranging from highly constrained tasks (structured extraction) to open-ended tasks (summarization), with the degree of output-space constraint serving as a primary determinant.

### 6.4 Multi-Turn and RAG: Reproducibility Under Complexity

Our multi-turn and RAG results address a key limitation of prior work (including our own earlier two-model study): reproducibility under complex interaction regimes. The finding that Gemma 2 9B and Mistral 7B maintain perfect EMR = 1.000 for both multi-turn refinement and RAG extraction demonstrates that conversational state accumulation and context augmentation do not inherently degrade reproducibility for deterministic local models. LLaMA 3’s slight degradation (EMR = 0.880 for multi-turn) suggests model-specific sensitivity to dialogue-turn interactions, possibly related to the warm-up effect observed in single-turn experiments. Crucially, Claude Sonnet 4.5’s near-zero EMR for both multi-turn (0.040) and RAG (0.000) confirms that the local-vs-API reproducibility gap extends beyond single-turn tasks. The RAG result—zero exact matches across 50 runs—suggests that longer outputs and additional retrieval context may amplify server-side variability, though a single API model cannot establish this as a general principle.

### 6.5 The Role of Provenance

The W3C PROV graphs generated by our protocol serve multiple purposes beyond simple audit trails:

- (1) **Automated comparison:** By comparing PROV graphs of two runs, one can automatically identify which factors differed (e.g., same prompt and model but different temperatures), enabling systematic diagnosis of non-reproducibility.
- (2) **Lineage tracking:** When outputs are used as inputs to downstream processes (e.g., summarization outputs fed into a meta-analysis), the provenance chain can be extended to trace any final result back to its full generation context.
- (3) **Compliance:** For regulated domains (healthcare, legal, finance), PROV documents provide the formal evidence trail required by audit standards ([National Institute of Standards and Technology 2023](#)) and emerging regulations such as the EU AI Act ([European Parliament and Council of the European Union 2024](#)).

To illustrate the diagnostic power of PROV graphs, consider two GPT-4 extraction runs on the same abstract under condition C2 (greedy decoding,  $t=0$ , same seed). Although the PROV entities for Prompt, InputText, ModelVersion, and InferenceParameters are identical (verified via matching SHA-256 hashes), the Output entities differ: `output_hash` values diverge, and the `wasGeneratedBy` timestamps differ by

several seconds. The PROV graph thus automatically pinpoints the source of non-reproducibility: the only varying factor is the RunGeneration activity itself, consistent with non-determinism arising from server-side factors. This kind of automated differential diagnosis is infeasible without structured provenance records.

898

## 899 6.6 Limitations

900 We organize threats to validity following standard categories:

902

903 **6.6.1 Internal Validity. Sample size.** LLaMA 3 uses 30 abstracts per condition, while the newer models  
 904 (Mistral, Gemma 2, Claude) use 10 abstracts. With  $n = 30$ , statistical power exceeds 0.999 for all primary  
 905 comparisons (Cohen 1988). With  $n = 10$ , the study is adequately powered for the large observed effect  
 906 sizes ( $d > 1.6$ ) but may miss subtler effects.

907 **GPT-4 C3 incomplete coverage.** Due to API quota exhaustion, GPT-4 extraction under C3 conditions covers 14–17 of 30 abstracts (summarization C3 is complete at 30). Our central claims rest on the  
 908 C2 condition (300/300 runs complete), and the C3 temperature sweep serves as a secondary analysis.

909 **Warm-up confound.** The first inference after model loading may differ from subsequent calls for  
 910 LLaMA 3 and Mistral 7B. This affects 2–4 abstracts per model, slightly reducing aggregate EMR.  
 912 Gemma 2 9B appears immune to this effect.

913 **Prompt format confound.** Single-turn experiments use Ollama’s /api/generate endpoint for local  
 914 models, whereas API models use their respective chat APIs. A supplementary control experiment (200  
 915 additional runs using Ollama’s /api/chat endpoint; see Appendix E) shows that this format difference  
 916 does not explain the reproducibility gap: LLaMA 3 produces *identical* variability metrics (summarization  
 917 EMR = 0.929, extraction EMR = 1.000) under both completion and chat formats.

918

919 **6.6.2 External Validity. Five models, two paradigms.** Our evaluation now covers three local models  
 920 and two API-served models, substantially strengthening the generalizability of the local-vs-API finding  
 921 compared to single-model-per-paradigm designs. However, other models—including Gemini (Gemini  
 922 Team et al. 2024), larger LLaMA variants, and open-weight models served via cloud APIs—may exhibit  
 923 different characteristics.

924 **Four tasks.** Our task suite now includes single-turn extraction/summarization, multi-turn refinement,  
 925 and RAG extraction. However, it does not cover code generation, mathematical reasoning, or creative  
 926 writing, which may exhibit different reproducibility patterns.

927 **English-only, single domain.** Our input data consists of 30 English scientific abstracts from AI/ML  
 928 papers. Reproducibility characteristics may differ for other languages, domains, or document types.

929 **Multi-turn limited to one API model.** Multi-turn and RAG experiments include Claude Sonnet  
 930 4.5 as the sole API representative; GPT-4 was not evaluated on Tasks 3–4 due to quota exhaustion.  
 931 While Claude’s near-zero EMR is consistent with the single-turn API pattern, other API providers may  
 932 exhibit different multi-turn reproducibility characteristics.

933

934 **6.6.3 Construct Validity. Surface-level metrics.** Our metrics (EMR, NED, ROUGE-L) capture textual  
 935 rather than semantic similarity. Two outputs that are semantically equivalent but syntactically different  
 936 will register as non-matching under EMR and partially divergent under NED. This is by design—our  
 937 focus is on *exact* reproducibility—but it means our results may overstate the practical impact of non-  
 938 determinism for downstream applications where semantic equivalence suffices.

939

940

941 **6.6.4 Other Considerations. Privacy.** The protocol’s environment metadata includes the machine host-  
 942 name, which may reveal institutional information. Deployments in privacy-sensitive settings should  
 943 anonymize this field.

944 **Computational cost.** The total cost was modest: approximately 8 GPU-hours on a consumer laptop  
 945 (Apple M4, 24 GB) for 2,000 local-model runs (including multi-turn and RAG experiments), plus 1,204  
 946 API calls to GPT-4 and Claude. The carbon footprint is negligible at this scale, and the logging overhead  
 947 (<30 ms per run) would not materially increase energy consumption even at thousands of runs.  
 948

## 949 **6.7 Protocol Minimality: An Ablation Analysis**

950 To substantiate our claim that the protocol captures a *minimal* set of metadata, we conducted an ablation  
 951 analysis in which we systematically removed each field group from the protocol schema and assessed which  
 952 audit questions became unanswerable. We defined 10 audit questions that a reproducibility-oriented re-  
 953 searcher might ask (e.g., “Can we verify the exact prompt used?”, “Can we detect output tampering?”,  
 954 “Can we trace full provenance?”) and mapped each to the protocol fields required to answer it. For this  
 955 analysis, we decomposed the Run Card’s five sections into eight finer-grained field groups by separat-  
 956 ing cross-cutting concerns: Identification, Model Context, Parameters, Input Content, Output Content,  
 957 Hashing (all SHA-256 digests), Environment, and Overhead (timing and storage metadata).

958 The results show that removing *any* of these eight field groups renders at least one audit question  
 959 unanswerable, demonstrating that no group is redundant. The Hashing group (SHA-256 hashes for  
 960 prompts, inputs, outputs, parameters, and environment) has the highest information density: its removal  
 961 affects 6 of 10 questions despite contributing only 410 bytes per run. Conversely, the Overhead group  
 962 (logging time metadata) is the least connected but remains necessary for overhead assessment. The  
 963 complete ablation results are available in the project repository.

964 This analysis demonstrates that the protocol is *minimal* in the sense that every field group is necessary  
 965 for at least one audit capability, while the total overhead remains at approximately 4,052 bytes per run.  
 966

## 967 **6.8 Practical Costs and Adoption**

968 One concern with any new protocol is whether the adoption burden is justified. We address this con-  
 969 crectly:

- 970 • **Implementation effort:** Our reference implementation adds approximately 600 lines of Python  
     (the protocol core) to an existing workflow. Integration requires 3–5 function calls per run.
- 971 • **Runtime cost:** <30 ms per run across all five models, negligible compared to inference times of  
     seconds to minutes for typical LLM calls.
- 972 • **Storage cost:** ~4 KB per run. Our 3,604 runs total approximately 14 MB—less than a single  
     model checkpoint.
- 973 • **Learning curve:** The protocol uses standard JSON and W3C PROV, requiring no specialized  
     knowledge beyond basic Python.

974 Against these modest costs, the protocol provides complete audit trails, automated provenance graphs,  
 975 tamper-detectable outputs via cryptographic hashing, and structured metadata that enable systematic  
 976 reproducibility analysis.

## 977 **6.9 Minimum Reporting Checklist for Generative AI Studies**

978 Based on our findings and the protocol design, we recommend that researchers conducting generative  
 979 AI experiments report, at minimum, the following five items (the full 15-item checklist is provided in  
 980 Appendix A):  
 981

- 988 (1) **Model identity and version:** Exact model name, version string, and—for local models—  
 989 weights hash.  
 990 (2) **Inference parameters:** Temperature, seed, top\_p, top\_k, max\_tokens, and decoding strategy.  
 991 For APIs where the seed is advisory or unsupported, this should be stated explicitly.  
 992 (3) **Reproducibility metrics over multiple runs:** Report at least EMR (or an equivalent exact-  
 993 match metric) and one semantic metric (e.g., BERTScore) over  $\geq 3$  repetitions per condition. A  
 994 single run is insufficient to characterize output stability.  
 995 (4) **Environment and deployment mode:** Whether inference was local or API-based, and the  
 996 execution environment (hardware, OS, library versions).  
 997 (5) **Output hashes:** SHA-256 or equivalent cryptographic hashes of outputs, enabling tamper de-  
 998 tection and automated comparison across studies.

999 Run Cards generated by our protocol automatically capture all five items, providing a machine-readable  
 1000 record that satisfies this checklist with no additional effort from the researcher.

## 1002 7 Conclusion

1003 We presented a lightweight protocol for logging, versioning, and provenance tracking of generative AI ex-  
 1004 periments, introducing Prompt Cards and Run Cards as novel documentation artifacts and adopting the  
 1005 W3C PROV data model for machine-readable provenance graphs. Through 3,604 controlled experiments  
 1006 with five models—three locally deployed (LLaMA 3 8B, Mistral 7B, Gemma 2 9B) and two API-served  
 1007 (GPT-4, Claude Sonnet 4.5)—across four NLP tasks and 30 scientific abstracts, we demonstrated five  
 1008 key findings:  
 1009

- 1010 (1) **API non-determinism is consistent across providers.** Both GPT-4 (OpenAI) and Claude  
 1011 Sonnet 4.5 (Anthropic) exhibit substantial non-determinism under greedy decoding on single-turn  
 1012 tasks (average EMR = 0.221), while all three local models achieve average EMR = 0.953. This  
 1013 more than 4-fold reproducibility gap, observed independently for two cloud providers, is consistent  
 1014 with non-determinism arising from factors common to cloud-hosted inference infrastructure rather  
 1015 than being a provider-specific artifact.  
 1016 (2) **Local models can achieve perfect bitwise reproducibility.** Gemma 2 9B attains EMR  
 1017 = 1.000 across all four tasks under greedy decoding—every output is character-for-character  
 1018 identical across repetitions.  
 1019 (3) **The local-vs-API gap extends to complex interaction regimes.** Multi-turn refinement and  
 1020 RAG extraction achieve EMR  $\geq 0.880$  for all local models (Gemma 2 9B and Mistral 7B: perfect  
 1021 EMR = 1.000), while Claude Sonnet 4.5—the only API model tested on these tasks—achieves  
 1022 EMR = 0.040 (multi-turn) and EMR = 0.000 (RAG).  
 1023 (4) **Temperature is the dominant user-controllable factor for local models.** Increasing from  
 1024  $t=0.0$  to  $t=0.7$  reduces EMR to zero for all five models on summarization, while seed variation  
 1025 has no effect under greedy decoding for local models. For API-served models, the temperature—  
 1026 reproducibility relationship is more complex and may be non-monotonic (see Section 5).  
 1027 (5) **Comprehensive provenance logging adds negligible overhead:** less than 1% of inference  
 1028 time and approximately 4 KB per run across all five models, removing any practical argument  
 1029 against systematic documentation.

1030 These findings carry a broader implication: a substantial portion of published research that relies  
 1031 on API-based LLMs may contain non-reproducible results without the authors' knowledge. The cost of  
 1032 systematic provenance logging—less than one percent of inference time—is trivially small compared to  
 1033 the cost of publishing non-reproducible science.

1035 Looking ahead, we plan to (i) extend the model suite to include Gemini (Gemini Team et al. 2024)  
 1036 and open-weight models served via cloud APIs (e.g., Hugging Face Inference Endpoints) to further  
 1037 disentangle model architecture from deployment infrastructure; (ii) extend the task coverage to code  
 1038 generation, mathematical reasoning, and agentic workflows; and (iii) develop automated reproducibility  
 1039 scoring based on provenance graph analysis. Ultimately, we envision a future in which every generative  
 1040 AI output carries a provenance certificate, and reproducibility metrics are reported alongside accuracy  
 1041 as a standard component of empirical evaluation.

1042 The reference implementation, all 3,604 run records, provenance documents, and analysis scripts are  
 1043 publicly available to support adoption and independent verification.

## 1044 Acknowledgments

1045 This work was supported by UTFPR – Universidade Tecnológica Federal do Paraná. The experiments  
 1046 were conducted using locally deployed open-weight models to ensure full reproducibility of the computa-  
 1047 tional environment.

## 1048 Data Availability Statement

1049 The reference implementation, all 3,604 run records (JSON), PROV-JSON provenance documents, Run  
 1050 Cards, Prompt Cards, input data, analysis scripts, and generated figures are publicly available at:

1051 <https://github.com/Roverlucas/genai-reproducibility-protocol>

1052 The repository includes instructions for reproducing all experiments and regenerating all tables and  
 1053 figures from the raw data.

## 1054 Author Contributions

1055 Following the CRediT (Contributor Roles Taxonomy) framework: **Lucas Rover**: Conceptualization,  
 1056 Methodology, Software, Validation, Formal Analysis, Investigation, Data Curation, Writing – Original  
 1057 Draft, Writing – Review & Editing, Visualization, Project Administration. **Yara de Souza Tadano**:  
 1058 Supervision, Conceptualization, Methodology, Writing – Review & Editing, Project Administration.

## 1059 Conflict of Interest

1060 The authors declare no conflicts of interest. This research was conducted independently at UTFPR with  
 1061 no external funding from commercial AI providers. The use of OpenAI’s GPT-4 API was for research  
 1062 evaluation purposes only and does not constitute an endorsement.

## 1063 Use of AI-Assisted Tools

1064 The authors used AI-assisted tools (Claude, Anthropic) during the preparation of this manuscript for  
 1065 language editing, code development support, and data analysis scripting. All AI-generated content was  
 1066 critically reviewed, validated, and revised by the authors, who take full responsibility for the accuracy  
 1067 and integrity of the final manuscript. The scientific design, experimental execution, interpretation of  
 1068 results, and intellectual contributions are entirely the authors’ own work.

## 1069 References

- 1070 J. Achiam et al. 2023. *GPT-4 Technical Report*. arXiv preprint. (2023). arXiv: 2303.08774 (cs.CL).
- 1071 Anthropic. 2024. *The Claude Model Family*. (2024). <https://www.anthropic.com/claude>.
- 1072 B. Atil et al. 2024. *Non-Determinism of “Deterministic” LLM Settings*. arXiv preprint. (2024). arXiv: 2408.04667 (cs.CL).
- 1073 M. Baker. 2016. “1,500 Scientists Lift the Lid on Reproducibility.” *Nature*, 533, 7604, 452–454. doi:10.1038/533452a.

- 1082 A. Belz, S. Agarwal, A. Shimorina, and E. Reiter. 2021. “A Systematic Review of Reproducibility Research in Natural Lan-  
1083 guage Processing.” In: *Proceedings of the 16th Conference of the European Chapter of the Association for Computational*  
1084 *Linguistics*. Association for Computational Linguistics, 381–393. doi:[10.18653/v1/2021.eacl-main.29](https://doi.org/10.18653/v1/2021.eacl-main.29).
- 1085 L. Biewald. 2020. *Experiment Tracking with Weights and Biases*. (2020). <https://wandb.com/>.
- 1086 R. Bommasani et al.. 2022. *On the Opportunities and Risks of Foundation Models*. arXiv preprint. (2022). arXiv: [2108.07258](https://arxiv.org/abs/2108.07258) (cs.LG).
- 1087 T. Brown et al.. 2020. “Language Models are Few-Shot Learners.” In: *Advances in Neural Information Processing Systems*.  
1088 Vol. 33, 1877–1901. arXiv: [2005.14165](https://arxiv.org/abs/2005.14165) (cs.CL).
- 1089 Y. Chen, J. Li, X. Liu, and Y. Li. 2023. *On the Reproducibility of ChatGPT in NLP Tasks*. arXiv preprint. (2023). arXiv:  
1090 [2304.02554](https://arxiv.org/abs/2304.02554) (cs.CL).
- 1091 J. Cohen. 1988. *Statistical Power Analysis for the Behavioral Sciences*. (2nd ed.). Lawrence Erlbaum Associates. ISBN:  
1092 978-0-8058-0283-2.
- 1093 J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. 2019. “BERT: Pre-training of Deep Bidirectional Transformers for  
1094 Language Understanding.” In: *Proceedings of the 2019 Conference of the North American Chapter of the Association*  
1095 *for Computational Linguistics*. Association for Computational Linguistics, 4171–4186. doi:[10.18653/v1/N19-1423](https://doi.org/10.18653/v1/N19-1423).
- 1096 J. Dodge, S. Gururangan, D. Card, R. Schwartz, and N. A. Smith. 2019. “Show Your Work: Improved Reporting of  
1097 Experimental Results.” In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*.  
1098 Association for Computational Linguistics, 2185–2194. doi:[10.18653/v1/D19-1224](https://doi.org/10.18653/v1/D19-1224).
- 1099 European Parliament and Council of the European Union. 2024. *Regulation (EU) 2024/1689 Laying Down Harmonised  
1100 Rules on Artificial Intelligence (AI Act)*. (2024). <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>.
- 1101 T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford. 2021. “Datasheets  
1102 for Datasets.” *Communications of the ACM*, 64, 12, 86–92. doi:[10.1145/3458723](https://doi.org/10.1145/3458723).
- 1103 Gemini Team, R. Anil, S. Borgeaud, J.-B. Alayrac, J. Yu, R. Soricut, J. Schalkwyk, A. M. Dai, A. Hauth, et al.. 2024.  
1104 *Gemini: A Family of Highly Capable Multimodal Models*. arXiv preprint. (2024). arXiv: [2312.11805](https://arxiv.org/abs/2312.11805) (cs.CL).
- 1105 Gemma Team et al.. 2024. *Gemma 2: Improving Open Language Models at a Practical Size*. arXiv preprint. (2024). arXiv:  
1106 [2408.00118](https://arxiv.org/abs/2408.00118) (cs.CL).
- 1107 A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, et al.. 2024.  
1108 *The LLaMA 3 Herd of Models*. arXiv preprint. (2024). arXiv: [2407.21783](https://arxiv.org/abs/2407.21783) (cs.AI).
- 1109 O. E. Gundersen, Y. Gil, and D. W. Aha. 2018. “On Reproducible AI: Towards Reproducible Research, Open Science, and  
1110 Digital Scholarship in AI Publications.” *AI Magazine*, 39, 3, 56–68. doi:[10.1609/aimag.v39i3.2816](https://doi.org/10.1609/aimag.v39i3.2816).
- 1111 O. E. Gundersen, M. Helmert, and H. H. Hoos. 2024. “Improving Reproducibility in AI Research: Four Mechanisms Adopted  
1112 by JAIR.” *Journal of Artificial Intelligence Research*, 81, 1019–1041. doi:[10.1613/jair.1.16905](https://doi.org/10.1613/jair.1.16905).
- 1113 O. E. Gundersen and S. Kjensmo. 2018. “State of the Art: Reproducibility in Artificial Intelligence.” *Proceedings of the  
1114 AAAI Conference on Artificial Intelligence*, 32, 1, 1644–1651. doi:[10.1609/aaai.v32i1.11503](https://doi.org/10.1609/aaai.v32i1.11503).
- 1115 M. Herschel, R. Diestelkämper, and H. Ben Lahmar. 2017. “A Survey on Provenance: What for? What form? What from?”  
1116 *The VLDB Journal*, 26, 6, 881–906. doi:[10.1007/s00778-017-0486-1](https://doi.org/10.1007/s00778-017-0486-1).
- 1117 M. Hutson. 2018. “Artificial Intelligence Faces Reproducibility Crisis.” *Science*, 359, 6377, 725–726. doi:[10.1126/science.359.6377.725](https://doi.org/10.1126/science.359.6377.725).
- 1118 A. Q. Jiang et al.. 2023. *Mistral 7B*. arXiv preprint. (2023). arXiv: [2310.06825](https://arxiv.org/abs/2310.06825) (cs.CL).
- 1119 S. Kapoor and A. Narayanan. 2023. “Leakage and the Reproducibility Crisis in Machine-Learning-Based Science.” *Patterns*,  
1120 4, 9, 100804. doi:[10.1016/j.patter.2023.100804](https://doi.org/10.1016/j.patter.2023.100804).
- 1121 R. Kuprieiev, D. Petrov, and Iterative. 2024. *DVC: Data Version Control*. (2024). <https://dvc.org/>.
- 1122 LangChain. 2023. *LangSmith: A Platform for Building Production-Grade LLM Applications*. (2023). <https://smith.langchain.com/>.
- 1123 V. I. Levenshtein. 1966. “Binary Codes Capable of Correcting Deletions, Insertions, and Reversals.” *Soviet Physics Doklady*,  
1124 10, 8, 707–710.
- 1125 P. Liang et al.. 2023. “Holistic Evaluation of Language Models.” *Transactions on Machine Learning Research*. <https://openreview.net/forum?id=iO4LZibEqW>.
- 1126 C.-Y. Lin. 2004. “ROUGE: A Package for Automatic Evaluation of Summaries.” In: *Proceedings of the ACL-04 Workshop on*  
1127 *Text Summarization Branches Out*. Association for Computational Linguistics, 74–81. <https://aclanthology.org/W04-1013>.
- 1128 M. Mitchell, S. Wu, A. Zaldivar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, and T. Gebru. 2019.  
1129 “Model Cards for Model Reporting.” In: *Proceedings of the Conference on Fairness, Accountability, and Transparency*.  
1130 ACM, 220–229. doi:[10.1145/3287560.3287596](https://doi.org/10.1145/3287560.3287596).

- 1129 L. Moreau and P. Missier. 2013. *PROV-DM: The PROV Data Model*. W3C Recommendation. World Wide Web Consortium.  
 1130 <https://www.w3.org/TR/prov-dm/>.
- 1131 National Institute of Standards and Technology. 2023. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*.  
 Tech. rep. U.S. Department of Commerce. doi:[10.6028/NIST.AI.100-1](https://doi.org/10.6028/NIST.AI.100-1).
- 1132 Ollama. 2024. *Ollama: Run Large Language Models Locally*. (2024). <https://ollama.com/>.
- 1133 OpenAI. 2023. *OpenAI Evals: A Framework for Evaluating LLMs*. (2023). <https://github.com/openai/evals>.
- 1134 S. Ouyang, J. M. Zhang, M. Harman, and M. Wang. 2024. “An Empirical Study of the Non-determinism of ChatGPT in  
 1135 Code Generation.” *ACM Transactions on Software Engineering and Methodology*, 34, 2, 1–28. doi:[10.1145/3697010](https://doi.org/10.1145/3697010).
- 1136 G. Padovani, V. Anantharaj, and S. Fiore. 2025. “yProv4ML: Effortless Provenance Tracking for Machine Learning Sys-  
 1137 tems.” *SoftwareX*, 29, 102028. doi:[10.1016/j.softx.2025.102028](https://doi.org/10.1016/j.softx.2025.102028).
- 1138 J. Pineau, P. Vincent-Lamarre, K. Sinha, V. Larivière, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and H. Larochelle. 2021.  
 1139 “Improving Reproducibility in Machine Learning Research: A Report from the NeurIPS 2019 Reproducibility Program.”  
*Journal of Machine Learning Research*, 22, 164, 1–20. <https://jmlr.org/papers/v22/20-303.html>.
- 1140 C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. 2020. “Exploring the  
 1141 Limits of Transfer Learning with a Unified Text-to-Text Transformer.” *Journal of Machine Learning Research*, 21, 140,  
 1–67. <https://jmlr.org/papers/v21/20-074.html>.
- 1142 S. Samuel and B. König-Ries. 2022. “A Provenance-based Semantic Approach to Support Understandability, Reproducibil-  
 1143 ity, and Reuse of Scientific Experiments.” *Journal of Biomedical Semantics*, 13, 1, 1–30. doi:[10.1186/s13326-022-00263-z](https://doi.org/10.1186/s13326-022-00263-z).
- 1144 V. Stodden, M. McNutt, D. H. Bailey, E. Deelman, Y. Gil, B. Hanson, M. A. Heroux, J. P. Ioannidis, and M. Taufer. 2016.  
 1145 “Enhancing Reproducibility for Computational Methods.” *Science*, 354, 6317, 1240–1241. doi:[10.1126/science.aah6168](https://doi.org/10.1126/science.aah6168).
- 1146 A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. 2017. “Attention  
 1147 is All You Need.” In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc. arXiv:  
[1706.03762 \(cs.CL\)](https://arxiv.org/abs/1706.03762).
- 1148 J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V. Le, and D. Zhou. 2022. “Chain-of-Thought  
 1149 Prompting Elicits Reasoning in Large Language Models.” In: *Advances in Neural Information Processing Systems*.  
 1150 Vol. 35, 24824–24837. arXiv: [2201.11903 \(cs.CL\)](https://arxiv.org/abs/2201.11903).
- 1151 M. D. Wilkinson et al. 2016. “The FAIR Guiding Principles for Scientific Data Management and Stewardship.” *Scientific  
 Data*, 3, 160018. doi:[10.1038/sdata.2016.18](https://doi.org/10.1038/sdata.2016.18).
- 1152 J. Yuan et al.. 2025. “Understanding and Mitigating Numerical Sources of Nondeterminism in LLM Inference.” In: *Advances  
 1153 in Neural Information Processing Systems*. Vol. 38. arXiv preprint. Curran Associates, Inc. arXiv: [2506.09501 \(cs.LG\)](https://arxiv.org/abs/2506.09501).
- 1154 M. Zaharia et al.. 2018. “Accelerating the Machine Learning Lifecycle with MLflow.” *IEEE Data Engineering Bulletin*, 41,  
 1155 4, 39–45. <http://sites.computer.org/debull/A18dec/p39.pdf>.
- 1156 T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi. 2020. “BERTScore: Evaluating Text Generation with  
 1157 BERT.” In: *Proceedings of the 8th International Conference on Learning Representations*. <https://openreview.net/forum?id=SkeHuCVFDr>.
- 1158 Y. Zhu, P. Zhang, E. Haq, P. Hui, and G. Buchanan. 2023. *Can ChatGPT Reproduce Human-Generated Labels? A Study  
 1159 of Social Computing Tasks*. arXiv preprint. (2023). arXiv: [2304.10145 \(cs.CL\)](https://arxiv.org/abs/2304.10145).
- 1160
- 1161
- 1162
- 1163
- 1164
- 1165
- 1166
- 1167
- 1168
- 1169
- 1170
- 1171
- 1172
- 1173
- 1174
- 1175

1176 **A Reproducibility Checklist**

1177 The following checklist is designed for self-assessment of reproducibility in generative AI studies. Each  
 1178 item maps to a specific field or artifact in our protocol.

1179

1180 **Prompt Documentation**

- 1181 (1) Is the exact prompt text recorded and versioned? [Prompt Card: `prompt_text, prompt_hash`]
- 1182 (2) Are design assumptions and limitations documented? [Prompt Card: `assumptions, limitations`]
- 1183 (3) Is the expected output format specified? [Prompt Card: `expected_output_format`]
- 1184 (4) Is the interaction regime documented (single/multi-turn)? [Prompt Card: `interaction_regime`]

1187 **Model and Environment**

- 1188 (5) Is the model name and version recorded? [Run Card: `model_name, model_version`]
- 1189 (6) Are model weights hashed for identity verification? [Run Card: `weights_hash`]
- 1190 (7) Is the execution environment fingerprinted? [Run Card: `environment, environment_hash`]
- 1191 (8) Is the source code version recorded? [Run Card: `code_commit`]

1193 **Execution and Output**

- 1195 (9) Are all inference parameters logged? [Run Card: `inference_params`]
- 1196 (10) Is the random seed recorded? [Run Card: `inference_params.seed`]
- 1197 (11) Is the output cryptographically hashed? [Run Card: `output_hash`]
- 1198 (12) Are execution timestamps recorded? [Run Card: `timestamp_start, timestamp_end`]
- 1199 (13) Is logging overhead measured separately? [Run Card: `logging_overhead_ms`]

1200 **Provenance**

- 1202 (14) Is a provenance graph generated per group? [PROV-JSON document]
- 1203 (15) Are provenance documents in an interoperable format? [W3C PROV standard]

1204

1205 **B Run Card Schema**

1206 The complete Run Card schema, with data types and descriptions:

1207

1208 Listing 1. Run Card JSON schema (simplified).

```

1209 1 {
1210 2   "run_id": "string (unique identifier)",
1211 3   "task_id": "string (task identifier)",
1212 4   "task_category": "string (e.g., summarization)",
1213 5   "prompt_hash": "string (SHA-256 of prompt)",
1214 6   "prompt_text": "string (full prompt text)",
1215 7   "input_text": "string (input to the model)",
1216 8   "input_hash": "string (SHA-256 of input)",
1217 9   "model_name": "string (e.g., llama3:8b)",
1218 10  "model_version": "string (e.g., 8.0B)",
1219 11  "weights_hash": "string (SHA-256 of weights)",
1220 12  "model_source": "string (e.g., ollama-local)",
1221 13  "inference_params": {
1222 14    "temperature": "float",

```

```

1223 15   "top_p": "float",
1224 16   "top_k": "integer",
1225 17   "max_tokens": "integer",
1226 18   "seed": "integer|null",
1227 19   "decoding_strategy": "string"
1228 20 },
1229 21   "params_hash": "string (SHA-256 of params)",
1230 22   "environment": {
1231 23     "os": "string",
1232 24     "os_version": "string",
1233 25     "architecture": "string",
1234 26     "python_version": "string",
1235 27     "hostname": "string",
1236 28     "timestamp": "ISO 8601 datetime"
1237 29 },
1238 30   "environment_hash": "string (SHA-256)",
1239 31   "code_commit": "string (git commit hash)",
1240 32   "researcher_id": "string",
1241 33   "affiliation": "string",
1242 34   "timestamp_start": "ISO 8601 datetime",
1243 35   "timestamp_end": "ISO 8601 datetime",
1244 36   "output_text": "string (model output)",
1245 37   "output_hash": "string (SHA-256 of output)",
1246 38   "output_metrics": "object (task-specific)",
1247 39   "execution_duration_ms": "float",
1248 40   "logging_overhead_ms": "float",
1249 41   "storage_kb": "float",
1250 42   "system_logs": "string (raw system info)",
1251 43   "errors": "array of strings",
1252 44
1253 45 // --- API-specific optional fields ---
1254 46   "api_request_id": "string|null (provider request ID)",
1255 47   "api_response_headers": "object|null (selected headers)",
1256 48   "api_model_version_returned": "string|null",
1257 49   "api_region": "string|null (if available)",
1258 50   "seed_status": "string (sent|logged-only|not-supported)",
1259 51
1260 52 // --- Multi-turn extension fields ---
1261 53   "conversation_history_hash": "string|null (SHA-256)",
1262 54   "turn_index": "integer|null",
1263 55   "parent_run_id": "string|null",
1264 56
1265 57 // --- RAG extension fields ---
1266 58   "retrieval_context": "string|null",
1267 59   "retrieval_context_hash": "string|null (SHA-256)"
1268 60 }

```

## C Example PROV-JSON Document

An abbreviated example of a PROV-JSON document generated for a single summarization run:

1269

Journal of Artificial Intelligence Research, Vol. , Article . Publication date: 2026.

Listing 2. Abbreviated PROV-JSON for a summarization run.

```

1270
1271 {
1272   "prefix": {
1273     "genai": "https://genai-prov.org/ns#",
1274     "prov": "http://www.w3.org/ns/prov#"
1275   },
1276   "entity": {
1277     "genai:prompt_c9644358": {
1278       "prov:type": "genai:Prompt",
1279       "genai:hash": "c9644358805b...",
1280       "genai:task_category": "summarization"
1281     },
1282     "genai:model_llama3_8b": {
1283       "prov:type": "genai:ModelVersion",
1284       "genai:name": "llama3:8b",
1285       "genai:source": "ollama-local"
1286     },
1287     "genai:output_590d0835": {
1288       "prov:type": "genai:Output",
1289       "genai:hash": "590d08359e7d..."
1290     }
1291   },
1292   "activity": {
1293     "genai:run_llama3_8b_sum_001_C1_rep0": {
1294       "prov:type": "genai:RunGeneration",
1295       "prov:startTime": "2026-02-07T21:54:34Z",
1296       "prov:endTime": "2026-02-07T21:54:40Z"
1297     }
1298   },
1299   "wasGeneratedBy": {
1300     "_:wGB1": {
1301       "prov:entity": "genai:output_590d0835",
1302       "prov:activity": "genai:run_llama3_8b..."
1303     }
1304   },
1305   "used": {
1306     "_:u1": {
1307       "prov:activity": "genai:run_llama3_...",
1308       "prov:entity": "genai:prompt_c9644358"
1309     }
1310   },
1311   "agent": {
1312     "genai:researcher_lucas_rover": {
1313       "prov:type": "prov:Person",
1314       "genai:affiliation": "UTFPR"
1315     }
1316   },
1317   "wasAssociatedWith": {
1318     "_:wAW1": {
1319

```

1317 Table 10. JSON extraction quality metrics by model and condition. *Raw Valid* = output parses directly as JSON;  
 1318 *Extracted Valid* = JSON extracted via regex from outputs containing preamble text; *Schema* = all five expected  
 1319 fields present; *Field EMR* = within-abstract pairwise exact match across runs for each extracted field, averaged over  
 1320 abstracts (see Section D for interpretation). LLaMA 3 always prepends introductory text (e.g., “Here is the extracted  
 1321 information in JSON format.”), yielding 0% raw validity but near-perfect extracted validity at  $t=0$ .

1322	1323 Model	Cond.	Raw	Extr.	Schema	Within-Abstract Field EMR					Overall Field EMR
			1324 Valid	Valid	Compl.	obj	meth	key_r	mod/sys	bench	
1325	1326 LLaMA 3	C1 ( $t=0$ )	0%	100%	100%	0.987	0.987	0.987	1.000	0.987	0.989
		C2 ( $t=0$ )	0%	100%	100%	0.987	0.987	0.987	1.000	0.987	0.989
		C3 ( $t=0.0$ )	0%	100%	100%	0.978	0.978	0.978	1.000	0.978	0.982
		C3 ( $t=0.3$ )	0%	97.8%	97.8%	0.747	0.460	0.552	0.862	0.805	0.685
		C3 ( $t=0.7$ )	0%	92.2%	92.2%	0.522	0.167	0.267	0.611	0.711	0.456
1330	1331 GPT-4	C2 ( $t=0$ )	100%	100%	100%	0.773	0.667	0.637	0.893	0.863	0.767
		C3 ( $t=0.0$ )	100%	100%	100%	0.833	0.571	0.667	0.905	0.810	0.757
		C3 ( $t=0.3$ )	100%	100%	100%	0.405	0.262	0.452	0.762	0.690	0.514
		C3 ( $t=0.7$ )	100%	100%	100%	0.137	0.157	0.255	0.667	0.725	0.388
		}									

```

1336 49 "prov:activity": "genai:run_llama3_...",
1337 50 "prov:agent": "genai:researcher_..."
1338 51 }
1339 52 }
1340 53 }
```

## D JSON Extraction Quality

1344 Table 10 presents JSON-specific quality metrics for the structured extraction task. Two notable patterns  
 1345 emerge.

1346 First, LLaMA 3 never produces raw-valid JSON: all 570 extraction outputs contain preamble text  
 1347 (e.g., “Here is the extracted information in JSON format.”) before the JSON object, despite the prompt  
 1348 explicitly requesting “JSON only, no explanation.” After extracting the embedded JSON via regex, va-  
 1349 lidity rates reach 100% under greedy decoding, degrading slightly at higher temperatures (92.2% at  
 1350  $t=0.7$ ). GPT-4, by contrast, always produces raw-valid JSON with 100% schema compliance across all  
 1351 conditions. This instruction-following gap is consistent with the different prompt interfaces: the chat  
 1352 completion API’s structured message format may better signal the expected output format.

1353 Second, within-abstract field-level exact match rates—computed by comparing only runs of the *same*  
 1354 abstract under the same condition, then averaging across abstracts—confirm the overall reproducibility  
 1355 hierarchy. Under greedy decoding, LLaMA 3 achieves near-perfect field EMR (0.982–0.989 overall), with  
 1356 all five fields at or above 0.978, consistent with the overall extraction EMR of 0.987 reported in Table 4.  
 1357 GPT-4 under greedy shows lower field EMR (0.757–0.767 overall), with open-ended fields (`method`: 0.667,  
 1358 `key_result`: 0.637) lagging behind structured fields (`model_or_system`: 0.893, `benchmark`: 0.863). As  
 1359 temperature increases, this gap widens: at  $t=0.7$ , `method` drops to 0.167 (LLaMA) and 0.157 (GPT-4),  
 1360 while `benchmark` retains 0.711 and 0.725 respectively—a 4–5× difference. This within-abstract formu-  
 1361 lation isolates true reproducibility (same input, same conditions, different runs) from between-abstract  
 1362 content variation, providing a methodologically clean measure of field-level consistency.

1364 **E Chat-Format Control Experiment**

1365 To assess whether the prompt-format difference between LLaMA 3 (completion-style via `/api/generate`)  
 1366 and GPT-4 (chat-style via Chat Completions) contributes to the observed reproducibility gap, we con-  
 1367 ducted a supplementary control experiment running LLaMA 3 8B through Ollama’s `/api/chat` endpoint,  
 1368 which applies the model’s chat template (including special tokens for system/user/assistant roles) in the  
 1369 same message structure used by GPT-4.

1370 **Design:** 10 abstracts  $\times$  2 tasks  $\times$  2 conditions (C1, C2)  $\times$  5 repetitions = 200 runs, all under greedy  
 1371 decoding ( $t=0$ ).

1372 **Results:** Table 11 compares the chat-format control with the original completion-format results for  
 1373 the same 10 abstracts. The two prompt formats produce *identical* variability metrics across all conditions:  
 1374 summarization EMR = 0.929, NED = 0.0066, and ROUGE-L = 0.9922 in both modes; extraction achieves  
 1375 perfect reproducibility (EMR = 1.000) regardless of interface. The 0.929 summarization EMR reflects the  
 1376 warm-up effect on 2 of 10 abstracts—the same pattern observed in the full 30-abstract experiment. These  
 1377 results confirm that prompt format is not a source of variability, and the reproducibility gap between  
 1378 LLaMA 3 and GPT-4 is consistent with deployment-side factors (server infrastructure, floating-point  
 1379 non-determinism across GPU types, request batching) rather than prompt-format differences.  
 1380

1381 Table 11. Prompt-format control: LLaMA 3 8B via completion (`/api/generate`) vs. chat (`/api/chat`) for 10 abstracts  
 1382 under greedy decoding ( $t=0$ ). EMR computed over conditions C1 and C2 combined.

Task	Metric	Completion	Chat
Summarization	EMR↑	0.929	0.929
	NED↓	0.0066	0.0066
	ROUGE-L↑	0.9922	0.9922
Extraction	EMR↑	1.000	1.000
	NED↓	0.0000	0.0000
	ROUGE-L↑	1.0000	1.0000

1392 *Note:* Completion and chat formats yield identical metrics for all 10 abstracts under greedy decoding, indicating that  
 1393 prompt format is not a source of variability.

1394  
 1395 Received February 2026  
 1396

1397  
 1398  
 1399  
 1400  
 1401  
 1402  
 1403  
 1404  
 1405  
 1406  
 1407  
 1408  
 1409  
 1410